

ElmerSolver Manual

Juha Ruokolainen, Mika Malinen, Peter Råback,
Thomas Zwinger, Antti Pursula and Mikko Byckling

CSC – IT Center for Science

December 19, 2018

ElmerSolver Manual

About this document

The ElmerSolver Manual is part of the documentation of Elmer finite element software. ElmerSolver Manual describes the Elmer Solver options common for all specific equation solvers. The different equations solver options are described separately in Elmer Models Manual. The ElmerSolver Manual is best used as a reference manual rather than a concise introduction to the matter.

The present manual corresponds to Elmer software version 8.4 for Windows NT and Unix platforms. The latest documentation and program versions of Elmer are available (or links are provided) at

`http://www.csc.fi/elmer`.

Copyright information

The original copyright of this document belongs to CSC – IT Center for Science, Finland, 1995–2018. This document is licensed under the Creative Commons Attribution-No Derivative Works 3.0 License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nd/3.0/>.

Elmer program is free software; you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation; either version 2 of the License, or (at your option) any later version. Elmer software is distributed in the hope that it will be useful, but without any warranty. See the GNU General Public License for more details.

Elmer includes a number of libraries licensed also under free licensing schemes compatible with the GPL license. For their details see the copyright notices in the source files.

All information and specifications given in this document have been carefully prepared by the best efforts of CSC, and are believed to be true and accurate as of time writing. CSC assumes no responsibility or liability on any errors or inaccuracies in Elmer software or documentation. CSC reserves the right to modify Elmer software and documentation without notice.

Contents

Table of Contents	3
1 Solving a multiphysics problem with the solver of Elmer: Fundamentals	6
1.1 Basic concepts	6
1.2 Handling interactions of multiple physical phenomena	7
1.3 The key abilities	8
2 Structure of the Solver Input File	10
2.1 Introduction	10
2.2 The sections of solver input file	10
2.3 Keyword syntax	14
2.4 Running several sequences	17
3 Restart from existing solutions	18
3.1 Restart file	18
3.2 Initialization of dependent variables	19
4 Solution methods for linear systems	21
4.1 Introduction	21
4.2 Direct methods	21
4.3 Preconditioned Krylov methods	22
4.4 Multilevel methods	23
4.5 Preconditioning via inner iterations	25
4.6 Keywords related to linear system solvers	26
4.7 Implementation issues	30
Bibliography	32
5 Nonlinear System Options	33
5.1 Introduction	33
5.2 Keywords related to solution of nonlinear systems	33
6 Integration of time-dependent systems	36
6.1 Introduction	36
6.2 Time discretization strategies	36
6.3 Keywords related to time discretization	37
6.4 On the treatment of time derivatives in Elmer Solver code	39
6.5 Predictor Corrector Scheme	40
Bibliography	40
7 Solving eigenvalue problems	41
7.1 Introduction	41
7.2 Theory	41
7.3 Keywords related to eigenvalue problems	42
7.4 Constructing matrices M and D in Solver code	44

8	Dirichlet boundary conditions and limiters	45
8.1	Dirichlet constraints	45
8.2	Soft limiters	47
9	Enforcing shared conditions between boundaries	48
9.1	Periodic conditions	48
9.2	Mortar conditions	49
10	Linear systems with constraints	54
10.1	Introduction	54
10.2	Keywords related to treatment of linear constraints	55
11	Solving constraint modes	56
11.1	Introduction	56
11.2	Keywords related to constraint modes analysis	56
12	Working with nodal loads	57
12.1	Setting nodal loads	57
12.2	Computing nodal loads	57
12.3	Computing nodal weights	58
12.4	Energy norm	58
13	Generic solver utilities	59
13.1	Solver activation	59
13.2	Variable names	59
13.3	Exported and derived variables	60
13.4	Active and passive elements	62
13.5	Timing the solvers	62
13.6	Consistency testing of solvers	62
14	Meshing Utilities	63
14.1	Coordinate transformation	63
14.2	Mesh multiplication	64
14.3	Mesh extrusion	64
14.4	Discontinuous mesh for internal jumps	66
14.5	Discontinuous Galerkin for internal jumps	67
15	Adaptive Solution	68
15.1	Introduction	68
15.2	Theory	68
15.3	Keywords related to the adaptive solution	70
15.4	Implementing own error estimators	70
16	Parallel runs	73
16.1	Introduction	73
16.2	Preprocessing of Parallel Runs	74
16.3	Parallel Computations in Elmer	77
16.4	Post-processing of Parallel Runs	82
17	Compilation and Linking	84
17.1	Compiling with CMake	84
17.2	Compiling a user-defined subroutine	85

18 Basic Programming	86
18.1 Introduction	86
18.2 Basic Elmer Functions and Structures	86
18.3 Writing a User Function	94
18.4 Writing a Solver	106
18.5 Compilation and Linking of User Defined Routines/Functions	118
A Format of mesh files	119
A.1 The format of header file	119
A.2 The format of node file	119
A.3 The format of element file	120
A.4 The format of boundary element file	120
A.5 The format of shared nodes file	120
A.6 Exceptions on parallel mesh format	121
B Format of result output files	122
B.1 Format versions	122
B.2 General structure	122
B.3 The positions file	125
C Format of ElmerPost Input File	126
D Basic element types	128
E Higher-order finite elements	131
E.1 Higher-order elements in Elmer	131
E.2 ElmerSolver services for higher-order elements	134
E.3 Basis functions	136
Bibliography	145
F Face and edge elements	146
F.1 The construction of face element interpolation	146
F.2 The construction of edge element interpolation	150
Bibliography	154
G Advanced finite element definitions	156
G.1 Specification of the degrees of freedom	156
Index	157

Chapter 1

Solving a multiphysics problem with the solver of Elmer: Fundamentals

Elmer software has been developed multiphysics simulations in mind. Thus, in addition to offering ways to produce computational solutions to single-physics models (the available collection of which is described in the Elmer Models Manual), Elmer provides procedures for creating computational models which describe interactions of multiple physical phenomena. Our intention here is to give an overview how this functionality is built into the solver of Elmer associated with the independent program executable ElmerSolver.

1.1 Basic concepts

The models handled by Elmer may generally be stationary or evolutionary, with nonlinearities possible in both the cases. Starting from the weak formulation of the problem, finite element approximation and advancing in time with implicit methods are typically applied in order to obtain the computational version of the model. In the simplest case of single-physics models we are then lead to solving equations

$$F(u) = 0, \quad (1.1)$$

where u represents either the vector of coefficients in the finite element expansion of the stationary solution or the coefficient vector to describe the evolutionary finite element solution at a given time $t = t_k$. Thus, in the case of evolution, the problems of the type (1.1) are solved repeatedly when advancing in time.

For linear models the problem (1.1) reduces to solving a linear system via defining

$$F(u) = b - Ku$$

where the coefficient matrix K is referred to as the stiffness matrix and b corresponds to the right-hand side vector in the linear system. Otherwise F is a nonlinear mapping and an iteration is needed to handle the solution of the problem (1.1). Available nonlinear iteration methods generally depend on the model, as the definition of the linearization strategy is a part of the computational description of each physical model.

We note that many single-physics models offer the possibility of using the Newton iteration where the current nonlinear iterate $u^{(m)}$ to approximate u is updated at each iteration step as

$$\begin{aligned} DF(u^{(m)})[\delta^{(m)}] &= -F(u^{(m)}), \\ u^{(m+1)} &= u^{(m)} + \delta^{(m)}, \end{aligned} \quad (1.2)$$

where $DF(u^{(m)})$ is the derivative of F at $u^{(m)}$. Thus, performing the nonlinear solution update again entails the solution of the linear system at each iteration step. As an alternate to the Newton method, linearization strategies based on lagged-value approximations are also often available. In addition, relaxation is conventionally offered as a way to enable convergence in cases where the basic nonlinear iteration fails

to produce convergence. Given the current nonlinear iterate $u^{(m)}$ and a computed correction $\delta u^{(m)}$ to the approximation, the new nonlinear iterate is then defined by

$$u^{(m+1)} = u^{(m)} + \lambda^{(m)} \delta^{(m)},$$

where $\lambda^{(m)}$ is an adjustable parameter referred to as the relaxation parameter.

1.2 Handling interactions of multiple physical phenomena

Having considered the basic concepts in the context of single-physics models, we now proceed to describe how the modularity employed in the design of Elmer allows us to create models which represent interactions of multiple physical phenomena. To this end, we assume that the complete model describes the interaction of N constituent models, the computational versions of which are primarily associated with the coefficient vectors u_i , $i = 1, 2, \dots, N$. As before, the coefficients contained in u_i are associated with the finite element expansion of either the stationary solution or the evolutionary solution at a time level $t = t_k$.

The fully discrete version of the coupled model leads to handling a problem of the form

$$\begin{aligned} F_1(u_1, u_2, \dots, u_N) &= 0, \\ F_2(u_1, u_2, \dots, u_N) &= 0, \\ &\dots \\ F_N(u_1, u_2, \dots, u_N) &= 0. \end{aligned} \tag{1.3}$$

If all the constituent models are linear, the problem (1.3) corresponds to solving the linear system where the coefficient matrix is a $N \times N$ block matrix. Otherwise (1.3) describes a nonlinear problem. Although the solution of (1.3) could in principle be done in the same way as explained in the context of single-physics models in Section 1.1, i.e. by performing either a coupled linear solve or Newton iteration, the coupled problems have usually been handled differently in order to enable the reuse of solvers for single-physics models and the easy extendability of the code to handle new applications.

To this end, the nonlinear Gauss-Seidel iteration is usually applied, so that the coupling of the models is resolved via generating new coupled system iterates $u^{(j)} = (u_1^{(j)}, u_2^{(j)}, \dots, u_N^{(j)})$ as

$$\begin{aligned} F_1(u_1^{(j)}, u_2^{(j-1)}, u_3^{(j-1)}, \dots, u_N^{(j-1)}) &= 0, \\ F_2(u_1^{(j)}, u_2^{(j)}, u_3^{(j-1)}, \dots, u_N^{(j-1)}) &= 0, \\ &\dots \\ F_N(u_1^{(j)}, u_2^{(j)}, \dots, u_N^{(j)}) &= 0. \end{aligned} \tag{1.4}$$

It is noted that the k th discrete model description in (1.4) depends implicitly only on the coupled system iterate to its primary variable u_k , while the dependencies on the other constituent model variables are treated explicitly. This brings us to solving a nonlinear single-field problem

$$F(u_k^{(j)}) = F_k(v_1, \dots, v_{k-1}, u_k^{(j)}, v_{k+1}, \dots, v_N) = 0, \text{ with all } v_j \text{ given}, \tag{1.5}$$

which is handled by using the methods already described in Section 1.1. We also note that if all the constituent models are linear the nonlinear Gauss-Seidel iteration (1.4) reduces to the block Gauss-Seidel iteration for linear systems. Relaxation may again be applied as an attempt to improve the convergence behaviour of the basic iteration (1.4).

It is good to pause here to stress that the main advantage of the adopted nonlinear Gauss-Seidel scheme is its support for the modular software design. Also, it brings us to handling coupled problems via solving linear systems which are smaller than those which would result from treating all constraints in (1.3) simultaneously. Despite these merits, the suitability of the loosely coupled iteration (1.4) generally is case-dependent as convergence problems may occur in cases where a strong physical coupling is involved. Such problems are often best handled by methods which treat all the constituent models in (1.3) simultaneously. Certain physical models available in Elmer indeed employ this alternate tightly coupled solution strategy. However,

these models have initially been developed independently, as common Elmer utilities for creating tightly coupled iteration methods in a general manner are less developed.

To summarize, the following pseudo-code presentation describes the basic loosely coupled iteration scheme employed by the solver of Elmer. This rough description may be helpful in summarizing what needs to be controlled overall to create a working computational solution procedure for a coupled problem.

! The time integration loop

for $k = 1 : M$

Generate an initial guess $u^{(0)} = (u_1^{(0)}, u_2^{(0)}, \dots, u_N^{(0)})$ for the coupled system solution at $t = t_k$

! The nonlinear Gauss-Seidel iteration

for $j = 1, 2, \dots$

! Generate the next coupled system iterate $u^{(j)}$ by performing single-field updates

for $i = 1 : N$

Set $v_l = u_l^{(j)}$ for $l = 1, 2, \dots, i - 1$

Set $v_l = u_l^{(j-1)}$ for $l = i + 1 : N$

Perform nonlinear iteration to solve $F_i(v_1, \dots, v_{i-1}, u_i^{(j)}, v_{i+1}, \dots, v_N) = 0$

Apply relaxation to set $u_i^{(j)} := u_i^{(j-1)} + \alpha_i(u_i^{(j)} - u_i^{(j-1)})$

end

end

end

Here the descriptions of the termination criteria for the iterations have been omitted. It is also noted that, obviously, the time integration loop is not needed in the case of a stationary problem. On the other hand, in the case of stationary simulation it is possible to replace the time integration loop by a pseudo-version of time stepping to enable performing multiple solutions for a range of model parameter values.

1.3 The key abilities

In the following, we give some additional information on the key abilities of the solver of Elmer to create computational solution procedures.

1.3.1 Extendability by modular design

A module of the Elmer software which enables the creation of the discrete model description of the type (1.5) and its solution with respect to the primary variable is generally called a solver. The solvers of Elmer are truly modular in this manner and have a standard interface. Thus, each solver usually contains an implementation of the nonlinear iteration, instructions to assemble the corresponding linear systems from elementwise contributions, and standard subroutine invocations to actually solve the linear systems assembled.

It follows that enabling an interaction with another field, designated by v_j in (1.5), is simply a matter of solver-level implementation. Therefore, interactions which have not been implemented yet can be enabled by making modifications which are localized to the solvers. In addition, a completely new physical model may be added by introducing a new solver which comprises a separate software module and which can be developed independently with respect to the main program. As a result, a loosely coupled solution procedure for a coupled problem based on the new physical model may again be achieved by making only solver-level modifications.

1.3.2 Model-specific finite element approximation

In the most basic setting all constituent model variables $u_i, i = 1, \dots, N$, of a coupled problem are approximated by using the same piecewise polynomial basis functions defined over a single mesh. In addition to this, the solver of Elmer offers a built-in functionality to perform a coupled problem simulation by using solver-specific finite element meshes. The solver description is then augmented by the specification of the

independent mesh which the solver uses. To make this functional in connection with the solution of coupled problems, Elmer has the capability of performing the solution data transfer, which is needed between the solvers in the loosely coupled solution procedure, even when the meshes are non-matching. It must be understood, however, that the loss of high-resolution details is unavoidable when the high-resolution field is represented by using a coarser finite element mesh.

1.3.3 Approximation by various finite element formulations

Elmer has traditionally employed the Galerkin finite element approximation of weak formulation based on the Lagrange interpolation basis functions. In this connection, the piecewise polynomial approximation of degree $1 \leq p \leq 3$ is possible for 2-D models, while 3-D models may be discretized by using the elements of degree $1 \leq p \leq 2$. The isoparametric mapping to describe curved element shapes is also supported with these traditional elements.

Discrete models based on more recent versions of the Galerkin finite element approximation are also possible. As an alternate to using the standard Lagrange interpolation basis functions, the Galerkin approximation based on using high-degree piecewise polynomials can be employed. In this connection, the degree of polynomial approximation can also be defined elementwise, with Elmer providing an in-built mechanism to guarantee the continuity of any solution candidate, so that in effect the use of the hp -version of the finite element method is enabled. However, generic ways to describe curved body surfaces accurately in connection with the high-degree finite elements have not been implemented yet which may limit the utility of these elements.

The way to define the high-degree approximation is based on the idea that a background mesh for representing the standard lowest-degree continuous finite element expansion is first provided so that a specific element type definition in relation to elements present in the background mesh may then be given to enhance the approximation. The same idea has been adapted to create other alternate finite element formulations. For example, finite element formulations which enhance the approximation defined on the background mesh by a subscale approximation spanned by elementwise bubble basis functions can be obtained in this way. We note that this strategy is widely used in Elmer to stabilize otherwise unstable formulations. Another example of the use of the user-supplied element definition relates to creating approximations based on the discontinuous Galerkin method. As a final example we mention that enhancing the approximation on the background mesh by either face or edge degrees of freedom and then omitting the original nodal degrees of freedom is also possible. This leads to a suitable set of unknowns for creating discretizations based on the Raviart-Thomas or edge element interpolation.

1.3.4 Parallel computing

A strength of the solver of Elmer is that it supports the use of parallel computing. This opportunity significantly widens the range of problems which can be considered. Additional details on utilizing parallel computers are found elsewhere in this manual.

1.3.5 Linear algebra abilities

This exposition should have made it clear that having the ability to solve large linear systems efficiently is a central aspect of the simulation process with Elmer. As explained, in the basic setting a linear solve is needed to obtain the solution update at each step of the nonlinear iteration. In practice linear solves are usually done iteratively, revealing one unexposed iteration level in relation to the pseudo-code presentation in the end of Section 1.2.

The solver of Elmer offers a large selection of strategies to construct linear solvers. The majority of them are directly implemented into Elmer software, but interfaces to exploit external linear algebra libraries are also available. Typically the most demanding aspect in the use of linear solvers relates to identifying an effective preconditioning strategy for the problem at hand. Traditionally Elmer has employed generic preconditioning strategies based on the fully algebraic approach, but recently alternate block preconditioning strategies which typically try to utilize physics-based splittings have also been developed.

Chapter 2

Structure of the Solver Input File

2.1 Introduction

Solving partial differential equation (PDE) models with the solver of Elmer requires that a precise description of the problem is given using the so-called solver input file, briefly referred to as the `sif` file. This file contains user-prepared input data which specify the location of mesh files and control the selection of physical models, material parameters, boundary conditions, initial conditions, stopping tolerances for iterative solvers, etc. In this chapter, the general structure of the file is described. We explain how the input data is organized into different sections and describe the general keyword syntax which is used in these sections to define the values of various model parameters and to control the solution procedures.

In the case of simple problem setups the solver input file may be written automatically by the preprocessor of Elmer software, so that knowing the solver input file format may be unnecessary. Creating a more complicated setup, or using keywords introduced by the user, however, requires the knowledge of the file format and keyword syntax.

In the following the general structure of the input file is first illustrated by using simple examples, without trying to explain all possibilities in an exhaustive manner. We then describe the keyword syntax in more detail, showing also how model parameters whose values depend on solution fields can be created. The later chapters of this manual, and Elmer Models Manual, which focuses on describing the PDE models Elmer can handle, provide more detailed material on specific issues. Elmer Tutorials also gives complete examples of solver input files.

2.2 The sections of solver input file

The material of the solver input file is organized into different sections. Each section is generally started with a row containing the name of the section, followed by a number of keyword commands, and ended with a row containing the word `End`. The names for starting new sections are

- `Header`
- `Simulation`
- `Constants`
- `Body n`
- `Material n`
- `Body Force n`
- `Equation n`

- Solver n
- Boundary Condition n
- Initial Condition n
- Component n

Here n associated with the section name represents an integer identifier needed for distinguishing between sections of the same type. A basic keyword command included in a section is nothing more than a statement which sets the value of a keyword with an equal sign.

In the following we describe how the sections are basically arranged without trying to explain all possibilities in an exhaustive manner. The later chapters of this manual and Elmer Models Manual provide more detailed material on specific issues. Elmer Tutorials also gives complete examples of solver input files.

Header section

The location of mesh files is usually given in the header section. Often this is also the only declaration given in the header section. If the Elmer mesh files (see Appendix A) are located in the directory ./mymesh, the header section may simply be

```
Header
  Mesh DB "." "mymesh"
End
```

Note that separate equations can nevertheless be discretized using different meshes if the location of mesh files is given in the solver section described below.

Simulation section

The simulation section is used for giving general information that is not specific to a particular PDE model involved in the simulation. This information describes the coordinate system used, indicates whether the problem is stationary or evolutionary, defines the file names for outputting, etc. Without trying to describe many possibilities and the details of commands, we only give the following simple example:

```
Simulation
  Coordinate System = "Cartesian 2D"
  Coordinate Mapping(3) = 1 2 3
  Coordinate Scaling = 0.001
  Simulation Type = Steady State
  Steady State Max Iterations = 1
  Output Intervals(1) = 1
  Post File = "case.ep"
  Output File = "case.dat"
End
```

Constants section

The constants section is used for defining certain physical constants. For example the gravity vector and the Stefan-Boltzmann constant may be defined using the commands

```
Constants
  Gravity(4) = 0 -1 0 9.82
  Stefan Boltzmann = 5.67e-08
End
```

If the constants are not actually needed in the simulation, this section can also be left empty.

Body, material, body force and initial condition sections

The Elmer mesh files contain information on how the computational region is divided into parts referred to as bodies. A body section associates each body with an equation set, material properties, body forces, and initial conditions by referring to definitions given in a specified equation section, material section, body force section, and initial condition section. To manage to do this, the different sections of the same type are distinguished by integer identifiers that are parts of the section names. Note that the integer in the body section name is an identifier for the body itself.

For example, one may define

```
Body 1
  Material = 1
  Body Force = 1
  Equation = 1
  Initial Condition = 2
End
```

Material properties, body forces, an equation set, and initial conditions are then defined in the material section

```
Material 1
...
End
```

the body force section

```
Body Force 1
...
End
```

the equation section

```
Equation 1
...
End
```

and the initial condition section

```
Initial Condition 2
...
End
```

What material properties and body forces need to be specified depends on the mathematical models involved in the simulation, and the initial condition section used for giving initial values is only relevant in the solution of evolutionary problems. We here omit the discussion of these very model-dependent issues; after reading this introductory chapter the reader should be able to understand the related documentation given in Elmer Models Manual, which focuses on describing the different mathematical models, while the contents of equation section will be described next.

Equation and solver sections

Equation section provides us a way to associate each body with a set of solvers, where each solver is typically associated with the ability to solve a certain physical model; cf. the definition of solver given in the beginning of Section 1.3.1. That is, if the set defined consists of more than one solver, several physical phenomena may be considered to occur simultaneously over the same region of space. The actual definitions of the solvers are given in solver sections, the contents of an equation section being basically a list of integer identifiers for finding the solver sections that define the solvers. The keyword commands given in the solver sections then control the selection of physical models, linearization procedures of nonlinear models, the selection of solution methods for resulting linear equations, convergence tolerances, etc.

For example, if only two solvers are needed, one may simply define a list of two solver identifiers

```
Equation 1
  Active Solvers(2) = 1 2
End
```

Then the solver definitions are read from the solver sections

```
Solver 1
  ...
End
```

and

```
Solver 2
  ...
End
```

Finally, we give an example of solver definitions, without trying to explain the commands at this point:

```
Solver 1
  Equation = "Poisson"
  Variable = "Potential"
  Variable DOFs = 1
  Procedure = "Poisson" "PoissonSolver"
  Linear System Solver = "Direct"
  Steady State Convergence Tolerance = 1e-06
End
```

Boundary condition section

Boundary condition sections define the boundary conditions for the different models. The Elmer mesh files contain information on how the boundaries of the bodies are divided into parts distinguished by their own boundary numbers. The keyword `Target Boundaries` is used to list the boundary numbers that form the domain for imposing the boundary condition. For example the declaration

```
Boundary Condition 1
  Target Boundaries(2) = 1 2
  ...
End
```

means that the boundary condition definitions that follow concern the two parts having the boundary numbers 1 and 2.

Component section

Component is a physical entity that is not associated with the mesh. It could be a representation of a 0D object such as electrical component, equation of state, chemical reaction, rigid object etc. The component could be independent, or it can be target for model reduction in the higher dimensional parts.

The following is an example of a component which is used as a reduction target for two bodies.

```
Component 1
  Name = "MyComponent"
  Master Bodies(2) = 1 2
  ...
End
```

This section that has been added much later than the other sections and it is still quite little used.

Text outside sections

We finally note that some commands, such as comments started with the symbol `!` and MATC expressions described below, may also be placed outside section definitions. An exception of this type is also the command

```
Check Keywords "Warn"
```

usually placed in the beginning of the input file. When this command is given, the solver outputs warning messages if the input file contains keywords that are not listed in the file of known keywords. If new keywords are introduced, misleading warning messages can be avoided by adding the new keywords to the keyword file `SOLVER.KEYWORDS`, located in the directory of the shared library files of ElmerSolver. The other options include `ignore`, `abort`, `silent`.

There is also the commands `echo on` and `echo off` that may be used to control the output of the parser. This is mainly intended for debugging purposes. The default is `off`.

2.3 Keyword syntax

As already illustrated, a basic keyword command used in the solver input file is a statement which sets the value of a solution parameter with the equal sign. Such a command in its full form also contains the data type declaration; for example

```
Density = Real 1000.0
```

Valid data types generally are

- Real
- Integer
- Logical
- String
- File

If the keyword is listed in the keyword file `SOLVER.KEYWORDS`, the data type declaration may be omitted. Therefore, in the case of our example, we may also define

```
Density = 1000.0
```

The value of a keyword may also be an array of elements of specified data type, with the array size definition associated with the keyword. We recall our previous examples of the equation and boundary condition sections, where we defined two lists of integers using the commands

```
Active Solvers(2) = 1 2
```

and

```
Target Boundaries(2) = 1 2
```

Two-dimensional arrays are also possible and may be defined as

```
My Parameter Array(3,3) = Real 1 2 3 \
                                4 5 6 \
                                7 8 9
```

Defining parameters depending on field variables

Most solver parameters may depend on time, or on the field variables defined in the current simulation run. Such dependencies can generally be created by means of tabular data, MATC functions, or Fortran functions. MATC has the benefit of being an interpreted language, making an additional compilation step with a compiler unnecessary.

Simple interpolating functions can be created by means of tabular data. The following example defines the parameter `Density` the value of which depends on the variable `Temperature`:

```
Density = Variable Temperature
Real
    0      900
    273 1000
    300 1020
    400 1000
End
```

This means that the value of `Density` is 900 when `Temperature` is 0, and the following lines are interpreted similarly. Elmer then uses linear interpolation to approximate the parameter for argument values not given in the table. If the value of the independent variable is outside the data set, the first or the last interpolating function which can be created from the tabulated values is used to extrapolate the value of the parameter.

If the field variable has several independent components, such as the components of displacement vector, the independent components may be used as arguments in a function definition. For example, if a three-component field variable is defined in a solver section using the commands

```
Variable = "Displ"
Variable DOFs = 3
```

then the solver of Elmer knows, in addition to the three-component vector `Displ`, three scalar fields `Displ 1`, `Displ 2` and `Displ 3`. These scalar fields may be used as independent variables in parameter definitions and used in the definitions of initial and boundary conditions, etc.

More complicated functions can be defined using MATC language. Here the basic usage of MATC in connection with the solver input file is illustrated; for an additional documentation, see a separate manual for MATC. For example, one may define

```
Density = Variable Temperature
MATC "1000*(1-1.0e-4*(tx-273))"
```

This means that the parameter `Density` depends on the value of `Temperature` as

$$\rho = \rho_0(1 - \beta(T - T_0)), \quad (2.1)$$

with $\rho_0 = 1000$, $\beta = 10^{-4}$ and $T_0 = 273$. Note that the value of the independent variable is known as `tx` in the MATC expression.

If the independent variable has more than one component, the variable `tx` will contain all the components in values `tx(0)`, `tx(1)`, ..., `tx(n-1)`, where `n` is the number of the components of the independent variable. A MATC expression may also take several scalar arguments; one may define, for example,

```
My Parameter = Variable Time, Displ 1
Real MATC "..."
```

The values of the scalar fields `Time` and `Displ 1` can then be referred in the associated MATC expression by the names `tx(0)` and `tx(1)`, respectively.

In addition to using MATC functions, Fortran 90 functions may also be used to create parameter definitions, etc. In the same manner as MATC functions are used, we may define

```
Density = Variable Temperature
Procedure "filename" "proc"
```

In this case the file "filename" should contain a shareable .so (Unix) or .dll (Windows) code for the user function whose name is "proc". The call interface for the Fortran function is as follows

```
FUNCTION proc( Model, n, T ) RESULT(dens)
  USE DefUtils)
  IMPLICIT None
  TYPE(Model_t) :: Model
  INTEGER :: n
  REAL(KIND=dp) :: T, dens

  dens = 1000*(1-1.0d-4*(T-273.0d0))
END FUNCTION proc
```

The Model structure contains pointers to all information about the model, and may be used to obtain field variable values, node coordinates, etc. The argument n is the index of the node to be processed, and T is the value of the independent variable at the node. The function should finally return the value of the dependent variable.

The independent variable can also be composed of several independent components. We may thus define

```
Density = Variable Coordinate
  Procedure "filename" "proc"
```

Now the argument T in the Fortran function interface should be a real array of three values, which give the x,y and z coordinates of the current node.

Parameterized keyword commands

The solver input file also offers possibilities for creating parameterized commands that utilize MATC. In the solver input file an expression following the symbol \$ is generally interpreted to be in MATC language. If the solver input file contains the lines

```
$solvertype = "Iterative"
$tol = 1.0e-6
```

then one may define, e.g.,

```
Solver 1
...
Linear System Solver = $solvertype
Linear System Convergence Tolerance = $tol
...
End

Solver 2
...
Linear System Solver = $solvertype
Linear System Convergence Tolerance = $100*$tol
...
End
```

Alternative keyword syntax

Some keyword commands may be given by using an alternative syntax that may sometimes be needed. The size of an integer or real number arrays may be given in parenthesis in connection with the keyword, but also with the Size declaration. Therefore the following are exactly the same

```
Timestep Intervals(3) = 1 10 100
Timestep Intervals = Size 3; 1 10 100
```


This feature is useful when giving vectors and matrices in ElmerGUI since there the keyword body is fixed and cannot include any size declaration. Note that in the above the semicolon is used as an alternative character for newline.

Another convention is to use two colons to make in-lined definitions in the sif files. The following two expressions are equal

```
Body Force 1
  Heat Source = 1.0
End
```

and

```
Body Force 1 :: Heat Source = 1.0
```

2.4 Running several sequences

Execution within command file

When reading the string RUN in the command file, the solver stops the reading and performs the computation with the instructions so far obtained. After a successful execution the solver continues to interpret the command file. Using this functionality it is therefore possible to create scripts where some parameter value is changed and the problem is recomputed. For example, adding the same sequence to the end of the sif file could be used to test the solution with another linear solver

```
RUN
Solver 1::Linear System Iterative Method = BiCgstabl
RUN
```

It should be noted that not quite all features support this procedure. For example, some preconditioners create static structures that will not be recreated.

Chapter 3

Restart from existing solutions

Often, the user wants to restart a run. This may be needed simply to continue an interrupted simulation — no matter what caused the interruption — but also to read in values needed either as initial conditions or as boundary conditions. These features are controlled by some keywords in the `Simulation` section that will be presented in the following sections.

3.1 Restart file

Any output file obtained by using the keyword command

```
Output File = String ...
```

can be used to define a restart point for a new simulation. By convention the suffix of the file has been `.result` but basically it may be chosen by the user.

By default all distributed fields will be saved to the restart file. Additionally the user may save global variables with the following keyword

```
Output Global Variables = Logical
```

If one wants to select the fields to be saved they may be given individually by

```
Output Variable i = String
```

where $i = 1, 2, 3, \dots$

The format of the output is either `ascii` or `binary`. `Ascii` is the default, and `binary` format is enforced by setting

```
Binary Output = True
```

The limitation of the restart functionality is that the mesh on which the previous case has been run must be identical to that on which the new run is performed. In parallel runs, additionally, also the partitions of the mesh have to coincide. The permutation vectors may vary, and also the field variables do not have to be the same.

The commands for restarting are then given in the `Simulation` section by declaring the restart file name as well as a restart position. For example we may specify

```
Simulation
  Restart File = "previousrun.result"
  Restart Position = 101
End
```

This would perform the current simulation by restarting from the time/iteration level 101 of the previously stored result file `previousrun.result`.

If one wants to select the fields to be restarted they may be given individually by

```
Restart Variable i = String
```

where $i = 1, 2, 3, \dots$

Upon running the new simulation, a message similar to the following example should be seen in the standard output of Elmer:

```
LoadRestartFile:
LoadRestartFile: -----
LoadRestartFile: Reading data from file: previousrun.result
LoadRestartFile: BINARY 3.L
LoadRestartFile:
LoadRestartFile: Total number of dofs to load: 13
LoadRestartFile: Reading time sequence:      2.000E-03
LoadRestartFile: Reading timestep: 10
LoadRestartFile: Time spent for restart (s):      7.9990E-03
LoadRestartFile: All done
LoadRestartFile: -----
```

If the number of stored time/iteration levels is not known a priori, the user can insert the command

```
Restart Position = 0
```

in order to make sure that the results for the lastly stored time/iteration level are reloaded.

Result files arising from steady state simulations often contain results for multiple iteration steps (with the result for the last step containing the converged solution). Nevertheless, these instances of solutions are — if reloaded — interpreted to describe the solution at different time levels. In this case the user might want to redefine the value of time variable for the restarted simulation, especially if continuing with transient runs. The keyword command

```
Restart Time = Real ...
```

may be given in order to manually set the time to correspond the zeroth time level of the new simulation.

3.2 Initialization of dependent variables

The initialization of variables and their boundary conditions is done by default before reading in the previous results. That has two main implications:

1. Values set in the section `Initial Condition` are overwritten by the corresponding values of the variable loaded afterwards from the restart file
2. Variable values given as initial or boundary conditions and specified to depend on other variables are not initiated with those values from the restart file by default.

The latter can be influenced with two keywords, `Restart Before Initial Conditions` (by default `False`) and `Initialize Dirichlet Condition` (by default `True`).

By setting

```
Restart Before Initial Conditions = Logical True
```

Elmer would first load the variables from the restart file and then apply initial conditions to those variables that have not been set by the earlier solution. This is necessary if one of the initial conditions depends on the earlier solution. By default, first the initial conditions from the solver input file are set and thereafter the restart file (if existing) is read.

The value of the keyword `Initialize Dirichlet Condition` is by default set to be true, which means that Dirichlet conditions are set before the simulation and thus also before the particular solver for handling that variable is executed. If a boundary condition for a certain variable now depends on the value of another, the first-time Dirichlet condition is set using the initial values of variables — either set or read in from the restart file. If this is not wanted, the user can switch to using the option

```
Initialize Dirichlet Condition = False
```

which will set the Dirichlet condition on the fly during the execution of the solver associated with the variable.

Chapter 4

Solution methods for linear systems

4.1 Introduction

Discretization and linearization of a system of partial differential equations generally leads to solving linear systems

$$Ax = b, \quad (4.1)$$

where A and b are of orders $n \times n$ and $n \times 1$, respectively. The coefficient matrix A resulting from the finite element discretization has a specific feature that the matrix is sparse, i.e. only a few of the matrix entries in each row differ from zero. In many applications the system can also have a very large order n , so that the chief part of the computation time in performing the simulation is typically spent by solvers for the linear systems.

Solution methods for linear systems fall into two large categories: direct methods and iterative methods. Direct methods determine the solution of the linear system exactly up to a machine precision. They perform in a robust manner leading to the solution after a predetermined number of floating-point operations. Nevertheless, the drawback of direct methods is that they are expensive in computation time and computer memory requirements and therefore cannot be applied to linear systems of very large order. The efficient solution of large systems generally requires the use of iterative methods which work by generating sequences of (hopefully) improving approximate solutions.

ElmerSolver provides access to both direct and iterative methods. The iterative methods available fall into two main categories: preconditioned Krylov subspace methods and multilevel methods. Iteration methods that combine the ideas of these two approaches may also be constructed. Such methods may be very efficient leading to a solution after a nearly optimal number of operation counts.

The development of efficient solution methods for linear systems is still an active area of research, the amount of literature on the topic being nowadays vast. The aim of the following discussion is to provide the user the basic knowledge of the solution methods available in ElmerSolver. The detailed description of methods is omitted. For a more comprehensive treatment the reader is referred to references mentioned.

4.2 Direct methods

A linear system may be solved in a robust way by using direct methods. ElmerSolver offers two main options for using direct methods. The default method utilizes the well-known LAPACK collection of subroutines for band matrices. In practice, this solution method can only be used for the solution of small linear systems as the operation count for this method is of order n^3 .

The other direct solver employs the Umfpack routines to solve sparse linear systems [1]. Umfpack uses the Unsymmetric MultiFrontal method. In practice it may be the most efficient method for solving 2D problems as long as there is enough memory available.

It should be noted that the success of the direct solvers depends very much on the bandwidth of the sparse matrix. In 3D these routines therefore usually fail miserably.

Elmer may be also compiled with Mumps, SuperLU, and Pardiso. The licensing scheme of these software do not allow the distribution of precompiled binaries but everybody may themselves compile a version that includes these solvers. Many times the best linear solver for a particular problem may be found among these.

4.3 Preconditioned Krylov methods

ElmerSolver contains a set of Krylov subspace methods for the iterative solution of linear systems. These methods may be applied to large linear systems but rapid convergence generally requires the use of preconditioning.

4.3.1 Krylov subspace methods

The Krylov subspace methods available in ElmerSolver are

- Conjugate Gradient (CG)
- Conjugate Gradient Squared (CGS)
- Biconjugate Gradient Stabilized (BiCGStab)
- BiCGStab(ℓ)
- Transpose-Free Quasi-Minimal Residual (TFQMR)
- Generalized Minimal Residual (GMRES)
- Generalized Conjugate Residual (GCR)

Both real and complex systems can be solved using these algorithms. For the detailed description of some of these methods see [3] and [4].

A definite answer to the question of what is the best iteration method for a particular case cannot be given. In the following only some remarks on the applicability of the methods are made.

The CG method is an ideal solution algorithm for cases where the coefficient matrix A is symmetric and positive definite. The other methods may also be applied to cases where A is non-symmetric. It is noted that the convergence of the CGS method may be irregular. The BiCGStab and TFQMR methods are expected to give smoother convergence. In cases where BiCGStab does not work well it may be advantageous to use the BiCGStab(ℓ) method, with $\ell \geq 2$ a parameter. Faster convergence in terms of iteration counts may be expected for increasing values of the parameter ℓ . However, since more work is required to obtain the iterate as ℓ increases, optimal performance in terms of computational work may be realized for quite a small value of ℓ . Starting with the value $\ell = 2$ is recommended.

The GMRES and GCR methods generate gradually improving iterates that satisfy an optimality condition. The optimality may however come with a significant cost since the computational work and computer memory requirements of these methods increase as the number of iterations grows. In practice these methods may be restarted after m solution updates have been performed in order to avoid the increasing work and storage requirements. The resulting methods are referred to as the GMRES(m) and GCR(m) methods. Here the choice of m has to be controlled by the user. It should be noted that the convergence of the restarted algorithms may be considerably slower than that of full versions. Unfortunately, general guidelines for determining a reasonable value for m cannot be given as this value is case-dependent.

4.3.2 Basic preconditioning strategies

The performance of iterative Krylov methods depends greatly on the spectrum of the coefficient matrix A . The rate at which an iteration method converges can often be improved by transforming the original system into an equivalent one that has more favorable spectral properties. This transformation is called preconditioning and a matrix which determines the transformation is called a preconditioner.

In ElmerSolver preconditioning is usually done by transforming (4.1) into the system

$$AM^{-1}z = b, \quad (4.2)$$

where the preconditioner M is an approximation to A and z is related to the solution x by $z = Mx$. In practice, the explicit construction of the inverse M^{-1} is not needed, since only a subroutine that for a given v returns a solution u to the system

$$Mu = v \quad (4.3)$$

is required.

The solver of Elmer provides several ways to obtain preconditioners. The basic strategies described in this section may be used to form the Jacobi preconditioner and incomplete factorization preconditioners. In addition to these preconditioners which remain stationary during the linear system iteration, the solver of Elmer allows the use of variable preconditioning where the action of the preconditioner is defined in terms of applying another iteration method to auxiliary preconditioning systems. Such inner-outer iterations are typically applied in connection with block preconditioners and preconditioning by multilevel methods. These advanced preconditioning strategies are considered in Section 4.5 below.

The Jacobi preconditioner is simply based on taking M to be the diagonal of A . More sophisticated preconditioners may be created by computing incomplete LU factorizations of A . The resulting preconditioners are referred to as the ILU preconditioners. This approach gives the preconditioner matrix M in the form $M = LU$ where L and U are lower and upper triangular with certain elements that arise in the factorization process ignored.

There are several ways to choose a set of matrix positions that are allowed to be filled with nonzero elements. ILU preconditioners of fill level N referred to as the ILU(N) preconditioners are built so that ILU(0) accepts nonzero elements in the positions in which A has nonzero elements. ILU(1) allows nonzero elements in the positions that are filled if the first step of Gaussian elimination is performed for A . ILU(2) accepts fill in positions that are needed if the next step of Gaussian elimination is performed with ILU(1) factorization, etc.

Another approximate factorization strategy is based on numerical tolerances. The resulting preconditioner is referred to as the ILUT preconditioner. In the creation of this preconditioner Gaussian elimination is performed so that elements of a given row of the LU factorization are obtained but only elements whose absolute value (scaled by the norm of all values of the row) is over a given threshold value are accepted in the preconditioner matrix.

Obviously, the additional computation time that is spent in creating the preconditioner matrix and solving systems of the type (4.3) should be compensated by faster convergence. Finding an optimal ILU preconditioner for a particular case may require the use of trial and error. Start with ILU(0) and try to increase the fill level N . As N increases, more and more elements in the incomplete LU factorization of the coefficient matrix are computed, so the preconditioner should in principle be better and the number of iterations needed to obtain a solution should decrease. At the same time the memory usage grows rapidly and so does the time spent in building the preconditioner matrix and in applying the preconditioner during iterations. The same applies to the ILUT preconditioner with decreasing threshold values.

4.4 Multilevel methods

A class of iterative methods referred to as multilevel methods provides an efficient way to solve large linear systems. For certain class of problems they perform nearly optimally, the operation count needed to obtain a solution being nearly of order n . Two different multilevel-method approaches are available in ElmerSolver, namely the geometric multigrid (GMG) and algebraic multigrid (AMG).

4.4.1 Geometric multigrid

Given a mesh \mathcal{T}_1 for the finite element discretization of problem the geometric multigrid method utilizes a set of coarser meshes \mathcal{T}_k , $k = 2, \dots, N$, to solve the linear system arising from the discretization. One of the fundamental ideas underlying the method is based on the idea of coarse grid correction. That is, a

coarser grid is utilized to obtain an approximation to the error in the current approximate solution of the linear system. The recursive application of this strategy leads us to multigrid methods.

To utilize different meshes multigrid methods require the development of methods for transferring vectors between fine and coarse meshes. Projection operators are used to transfer vectors from a fine mesh \mathcal{T}_k to a coarse mesh \mathcal{T}_{k+1} and will be denoted by I_k^{k+1} , while interpolation operators I_{k+1}^k transfer vectors from a coarse mesh to a fine mesh.

The multigrid method is defined by the following recursive algorithm: Given A , b and an initial guess y for the solution of the system $Ax = b$ set $i = 1$ and do the following steps:

1. If $i = N$, then solve the system $Ax = b$ by using the direct method and return.
2. Do pre-smoothing by applying some iterative algorithm for a given number of times to obtain a new approximate solution y .
3. Perform coarse grid correction by starting a new application of this algorithm with $A = I_i^{i+1} A I_{i+1}^i$, $b = I_i^{i+1}(Ay - b)$, $i = i + 1$ and the initial guess $e = 0$.
4. Compute a new approximate solution by setting $y = y + I_{i+1}^i e$
5. Do post-smoothing by applying some iterative algorithm for a given number of times to obtain a new approximate solution y .
6. If the solution has not yet converged, go to step 2.

In ElmerSolver one may choose the Jacobi, CG or BiCGStab algorithm as the method for smoothing iterations.

The full success of multigrid methods is based on the favorable combination of the properties of basic iteration methods and methods for transferring vectors between meshes. The smoothing iterations give rapid convergence for oscillatory solution components while coarse grid correction entails an efficient solution method for smooth solution components. For a comprehensive introduction to the geometric multigrid method the reader is referred to [2].

4.4.2 Algebraic multigrid

In many cases the geometric multigrid may not be applied because we do not have the luxury of having a set of appropriate hierarchical meshes. The alternative is the algebraic multigrid (AMG) method which uses only the matrix A to construct the projectors and the coarse level equations. AMG is best suited for symmetric and positive semidefinite problems. For other types of problems the standard algorithm may fail. For more information on AMG see reference [5].

The AMG method has two main phases. The set-up phase includes the recursive selection of the coarser levels and definition of the transfer and coarse-grid operators. The solution phase uses the resulting components to perform a normal multigrid cycling until a desired accuracy is reached. The solution phase is similar to that of the GMG.

Note that the AMG solvers in ElmerSolver are not fully mature. They may provide good solutions for some problems while desperately failing for others.

Classical Ruge-Stüben algorithm

The coarsening is performed using a standard Ruge-Stüben coarsening algorithm. The possible connections are defined by the entries in the matrix A . The variable i is strongly coupled to another variable j if

$$a_{ij} < -c_- \max |a_{ik}| \quad \text{or} \quad a_{ij} > c_+ \max |a_{ik}|, \quad (4.4)$$

where $0 < c_- < 1$ and $0 < c_+ < 1$ are parameters. Typically $c_- \approx 0.2$ and $c_+ \approx 0.5$. Once the negative (P^-) and positive (P^+) strong couplings have been determined the variables are divided into coarse (C) and fine (F) variables using the standard coarsening scheme.

The interpolation matrix may be constructed using the C/F -splitting and the strong couplings of the matrix. The interpolation of coarse nodes is simple as they remain unchanged. The interpolation of fine nodes starts from the fact the smooth error e must roughly satisfy the condition $Ae = 0$ or

$$a_{ii}e_i + \sum_{j \neq i} a_{ij}e_j = 0. \quad (4.5)$$

By manipulation

$$a_{ii}e_i + \alpha_i \sum_{j \in C \cap P_i^-} a_{ij}e_j + \beta_i \sum_{j \in C \cap P_i^+} a_{ij}e_j = 0, \quad (4.6)$$

where

$$\alpha_i = \frac{\sum_{j \in P_i^-} a_{ij}}{\sum_{j \in C \cap P_i^-} a_{ij}} \quad \text{and} \quad \beta_i = \frac{\sum_{j \in P_i^+} a_{ij}}{\sum_{j \in C \cap P_i^+} a_{ij}}. \quad (4.7)$$

The interpolation thus becomes

$$e_i = \sum_{j \in C \cap P_i} w_{ij}e_j \quad \text{with} \quad w_{ij} = \begin{cases} -\alpha_i a_{ij}/a_{ii}, & j \in P_i^-, \\ -\beta_i a_{ij}/a_{ii}, & j \in P_i^+. \end{cases} \quad (4.8)$$

This is known as *direct interpolation*. It may be modified by using also the strong F -nodes in the interpolation. This means that in formula (4.5) the following elimination is made for each $j \in F \cap P_i$

$$e_j \rightarrow - \sum_{k \in C \cap P_j} a_{jk}e_k/a_{jj}. \quad (4.9)$$

This is known as *standard interpolation*. In practice it means that the number of nodes used in the interpolation is increased. This may be important to the quality of the interpolation particularly if the number of direct C -neighbors is small.

After the interpolation weights have been computed the smallest coefficients may be truncated if they are small, i.e., $w_j < c_w \max |w_k|$, where $c_w \approx 0.2$. The other values must accordingly be increased so that the sum of weights remains constant. The truncation is essential in preventing the filling of the coarse level matrices.

Cluster multigrid

There is also an implementation of the agglomeration or cluster multigrid method. It is a variant of the algebraic multigrid method. In this method the components are grouped and the coarse-level matrices are created simply by summing up the corresponding rows and columns. In other words, the projection matrix includes just ones and zeros.

The cluster multigrid method should be more robust for problems where it is difficult to generate an optimal projection matrix. However, for simple problems it is usually beaten by the standard Ruge-Stüben method.

4.5 Preconditioning via inner iterations

In the following, we return to the subject of preconditioning and consider the construction of more advanced preconditioners that utilize the idea of variable preconditioning. We recall that the procedural description of preconditioning allows us to consider schemes where the action of M^{-1} in (4.2) is replaced by performing an inexact solution of the preconditioning system of the type (4.3) with other iterative methods. In this connection, the iteration steps of the method applied to the preconditioning system are referred to as inner iterations, while the iteration steps of the method which is being preconditioned are referred to as outer iterations. In addition, the concept of variable preconditioning is motivated by the fact that the preconditioner can change during the outer iteration process. It should be noted that all the Krylov methods do not necessarily admit using variable preconditioning. Anyhow, GCR algorithm can always be employed as the outer iterative method, and we generally recommend using it in connection with inner-outer iterations.

4.5.1 Preconditioning by multilevel methods

Multilevel methods are iteration methods on their own but they can also be applied to define preconditioners for the Krylov subspace methods. This preconditioning approach corresponds to taking $M = A$ in (4.3) and performing an inaccurate solution of the resulting system using multilevel methods to obtain u . A rather mild stopping criterion may be used in this connection. Preconditioning by multilevel methods may lead to very efficient solution methods for large linear systems.

4.5.2 Block preconditioning

Performing preconditioning by solving auxiliary systems iteratively also arises naturally in connection with a strategy referred to as the block preconditioning. Tacit in the concept of block preconditioning is that the unknowns of the linear system are grouped into sets of variables so that a natural block partitioning of the linear system is induced. Often such partitioning is naturally induced by physics. Block factorizations, block matrix splittings, or even other iteration schemes based on updating different types of variables independently can then be utilized to derive preconditioners. The bottom line is that such block preconditioners typically lead to solving subsidiary problems which can be handled by employing inner iterations.

The designs of efficient block preconditioners tend to involve problem-specific features, especially, if block factorizations or other segregated iteration methods are used to derive preconditioners. Therefore the model solvers in Elmer that offer the possibility of using block preconditioning have originally been developed independently, with the block preconditioning ability being a part of the solver-level implementation. Recently, routines which are suited for constructing block preconditioners on a more general level have also been implemented.

4.6 Keywords related to linear system solvers

The following keywords may be given in Solver section of the solver input file (.sif file).

Linear System Solver String

Using this keyword the type of linear system solver is selected. This keyword may take the following values:

- Direct
- Iterative
- Multigrid

Here Iterative and Multigrid refer to the Krylov and multilevel methods, respectively.

Linear System Direct Method String

If the value of the Linear System Solver keyword is set to be Direct, one may choose a band matrix solver with the value Banded or a sparse matrix solver with the value umfpack, mumps, Pardiso or superlu, . The default is Banded.

Linear System Iterative Method String

If the value of the Linear System Solver keyword is set to be Iterative, one should choose a Krylov method by setting the value of this keyword to be one of the following alternatives:

- CG
- CGS
- BiCGStab
- BiCGStabl
- TFQMR
- GMRES
- GCR

See also the `MG Smoother` keyword.

`Linear System GMRES Restart Integer [10]`

The restart parameter m for the `GMRES(m)` method may be given using this keyword.

`Linear System GCR Restart Integer`

The restart parameter m for the `GCR(m)` method may be given using this keyword. The default option is that restarting is not performed, i.e. the full GCR is used.

`BiCGstabl polynomial degree Integer`

The parameter ℓ for the `BiCGStab(ℓ)` method may be given. By default the minimal applicable value $\ell = 2$ is used.

`Linear System Preconditioning String`

A preconditioner for the Krylov methods may be declared by setting the value of this keyword to be one of the following alternatives:

- None
- Diagonal
- ILUn, where the literal n may take values 0,1,...,9.
- ILUT
- Multigrid

See also the `MG Preconditioning` keyword.

`Linear System ILUT Tolerance Real [0.0]`

This keyword is used to define the value of the numerical tolerance for the ILUT preconditioner.

`Linear System Convergence Tolerance Real [0.0]`

This keyword is used to define a stopping criterion for the Krylov methods. The approximate solution is considered to be accurate enough if the iterate satisfies

$$\frac{\|Ax - b\|}{\|b\|} \leq \epsilon$$

where ϵ is the value of this keyword. See also `MG Tolerance`.

`Linear System Max Iterations Integer [0]`

This keyword is used to define the maximum number of the iterations the Krylov methods are permitted to perform. If this limit is reached and the approximate solution does not satisfy the stopping criterion, `ElmerSolver` either continues the run using the current approximate solution as the solution of the system or aborts the run depending on the value of `Linear System Abort Not Converged` keyword. See also `MG Max Iterations` keyword.

`Linear System Abort Not Converged Logical [True]`

If the value of this keyword is set to be `True`, `ElmerSolver` aborts the run when the maximum number of iterations the algorithm is permitted to perform is reached and the approximate solution does not satisfy the stopping criterion. Otherwise the run will be continued using the current approximate solution as the solution of the system (this may lead to troubles at later steps of computation).

`Linear System Residual Mode Logical`

Usually the solvers of `Elmer` are written so that the linear system is solved for the full amplitude of the solution. However, often it is numerically to solve the residual equation where a change in the system of solved such that the residual would be minimized. This has the advantage that the linear system does not need to be solved as accurately. In practice we take the linear equation $Ax = b$ and make a change of variables such that $A dx = r$, where $r = b - Ax_0$ and $dx = x_1 - x_0$. After solution of the modified equation we revert back to the original system by setting $x_1 = x_0 + dx$.

Linear System Residual Output Integer [1]

By default the iterative algorithms display the value of the (scaled) residual norm after each iteration step. Giving a value $n > 1$ for this keyword may be used to display the residual norm only after every n iterations. If the value 0 is given, the residual output is disabled.

Linear System Precondition Recompute Integer [1]

By default the ElmerSolver computes the preconditioner when a new application of iterative algorithm is started. If the value of this keyword is set to be n , the preconditioner is computed only after n successive subroutine calls for linear systems arising from same source. This may speed up the solution procedure especially in cases where the coefficient matrix does not change much between successive subroutine calls. On the other hand if the coefficient matrix has changed significantly, the preconditioner may not be efficient anymore.

Optimize Bandwidth Logical [True]

If the value of this keyword is set to be `True`, the Cuthill-McKee bandwidth optimization scheme is used to order the unknowns in such a way that band matrices can be handled efficiently. The bandwidth optimization is recommended when the direct solver or incomplete factorization preconditioners are used.

The keywords beginning with `MG` are activated only if either the `Linear System Solver` or `Linear System Preconditioning` keyword has the value `Multigrid`. If a multigrid method is used as the linear system solver, some keywords starting with `MG` may be replaced by corresponding keywords starting with phrase `Linear System`. It should be noted that in the case of a multigrid solver there are some limitations to what values the keywords starting with the phrase `Linear System` may take, see below.

MG Levels Integer [1]

This keyword is used to define the number of levels for the multigrid method.

MG Equal Split Logical [False]

A hierarchy of meshes utilized by the multigrid method may be generated automatically by setting the value of this keyword to be `True`. The coarsest mesh must be supplied by the user and is declared in the usual way in the Header section of the solver input file. The other meshes are obtained using an equal division of the coarse mesh. The solution of the problem will be sought for the finest mesh.

MG Mesh Name File

A hierarchy of meshes utilized by the multigrid method may be supplied by the user. A base name of the mesh directories is declared using this keyword. The names of mesh directories must be composed of the base name appended with a level number such that if the base name is `mgridmesh`, the mesh directories should have names `mgridmesh1`, `mgridmesh2`, etc. The meshes are numbered starting from the coarsest mesh. In addition, the finest mesh must be declared in the Header section of the solver input file. It should be noted that the `MG Equal Split` keyword must be set to be `False` to enable the use of user-supplied meshes.

MG Max Iterations Integer [0]

If a multigrid method is used as a preconditioner for the Krylov methods, the value of this keyword defines the maximum number of iterations the multigrid method is allowed to perform to solve the preconditioning equation (4.3). Usually one or two iterations are sufficient. If a multigrid method is the linear system solver, the use of this keyword is similar to that of the `Linear System Max Iterations` keyword.

MG Convergence Tolerance Real [0.0]

If a multigrid method is used as a preconditioner for the Krylov methods, this keyword defines the solution accuracy for the preconditioning equation (4.3). This keyword is not usually needed if the `MG Max Iterations` keyword has a small value. If a multigrid method is the linear system solver, the use of this keyword is similar to that of the `Linear System Convergence Tolerance` keyword.

MG Smoother `String`

This keyword defines the algorithm for pre- and post-smoothing. It may take one of the following values:

- Jacobi
- CG
- BiCGStab

If the linear system solver is a multigrid method, the `Linear System Iterative Method` keyword may be used instead of this keyword, but only the three algorithms mentioned here can be applied.

MG Pre Smoothing Iterations `Integer [0]`

This keyword defines the number of pre-smoothing iterations.

MG Post Smoothing Iterations `Integer [0]`

This keyword defines the number of post-smoothing iterations.

MG Preconditioning `String`

This keyword declares the preconditioner for the algorithm which is used in smoothing iterations. It may take one of the following values:

- None
- ILUn, where the literal `n` may take values 0,1,...,9.
- ILUT

Note that this keyword is not related to using multigrid method as a preconditioner. It is also noted that preconditioning the smoothing algorithms does not seem to work well if a multigrid method is used as a preconditioner for Krylov methods.

MG ILUT Tolerance `Real [0.0]`

This keyword defines the numerical tolerance for the ILUT preconditioner in connection with smoothing iterations.

The keywords for the algebraic multigrid solver are in a large part the same as for the geometric multigrid. There are however some keywords that are related only to AMG.

MG Lowest Linear Solver Limit `Integer`

This value gives a lower limit for the set of coarse nodes after which the recursive multilevel routine is terminated. A proper value might be around 100.

MG Recompute Projector `Logical`

This flag may be used to enforce recomputation of the projector each time the algebraic multigrid solver is called. The default is `False` as usually the same projector is appropriate for all computations.

MG Eliminate Dirichlet `Logical`

At the highest level the fixed nodes may all be set to be coarse since their value is not affected by the lower levels. The default is `True`

MG Eliminate Dirichlet Limit `Real`

Gives the maximum fraction of non-diagonal entries for a Dirichlet node.

MG Smoother `String`

In addition to the selection for the GMG option `sor` (symmetric over relaxation) is possible.

MG SOR Relax `String`

The relaxation factor for the SOR method. The default is 1.

MG Strong Connection Limit Real

The coefficient c_- in the coarsening scheme. Default is 0.25.

MG Positive Connection Limit Real

The coefficient c_+ in the coarsening scheme. Default is 1.0.

MG Projection Limit Real

The coefficient c_w in the truncation of the small weights. The default is 0.1.

MG Direct Interpolate Logical

Chooses between direct and standard interpolation. The default is `False`.

MG Direct Interpolate Limit Integer

The standard interpolation may also be applied to nodes with only a small number of coarse connection. This gives the smallest number of nodes for which direct interpolation is used.

Finally, there are also some keywords related only to the clustering multigrid.

MG Cluster Size Integer

The desired choice of the cluster. Possible choices are 2,3,4,5,... and zero which corresponds to the maximum cluster.

MG Cluster Alpha Real

In the clustering algorithm the coarse level matrix is not optimal for getting the correct convergence. Tuning this value between 1 and 2 may give better performance.

MG Strong Connection Limit Real

This is used similarly as in the AMG method except it is related to positive and negative connections alike.

MG Strong Connection Minimum Integer

If the number of strong connections with the given limit is smaller than this number then increase the set of strong connection if available connections exist.

4.7 Implementation issues

4.7.1 The sparse matrix storage

To be efficient, iteration methods require that a matrix-vector product for sparse matrices is efficiently implemented. A special storage scheme called the Compressed Row Storage (CRS) [3] is used in ElmerSolver to store only those matrix coefficients that differ from zero.

The matrix structure is defined in module `Types` as:

```
TYPE Matrix_t
...
INTEGER :: NumberOfRows

REAL(KIND=dp), POINTER :: Values(:)
INTEGER, POINTER :: Rows(:), Cols(:), Diag(:)
...
END TYPE Matrix_t
```

The matrix type has several additional fields, but the basic storage scheme can be implemented using the fields shown. The array `Values` is used to store the nonzero elements of the coefficient matrix. The array `Cols` contains the column numbers for the elements stored in the array `Values`, while the array `Rows` contains indices to elements that start new rows. In addition, `Row(n+1)` gives the number of nonzero matrix elements + 1. The array `Diag` is used to store the indices of the diagonal elements.

For example, to go through the matrix row by row the following loop may be used

```

USE Types
TYPE(Matrix_t), POINTER :: A
REAL(KIND=dp) :: val
INTEGER :: i, j, row, col

DO i=1, A % NumberOfRows
  PRINT *, 'Diagonal element for row ', i, ' is ', A % Values( A % Diag(i) )
  DO j=A % Rows(i), A % Rows(i+1)-1
    row = i
    col = A % Cols(j)
    val = A % Values(j)
    PRINT *, 'Matrix element at position: ', row,col, ' is ', val
  END DO
END DO

```

4.7.2 Subroutine calls

Most of the functionality of the sparse linear system solver of the ElmerSolver is available by using the function call

```
Norm = DefaultSolve().
```

The return value `Norm` is a norm of the solution vector.

Sometimes it may be convenient to modify the linear system before solving it. A Fortran function which performs this modification can be written by the user with the name of the function being declared in the solver input file. For example, this technique may be used to define a user-supplied linear system solver.

If the name of the user-supplied Fortran function is `proc` and it can be found in the file having the name `Filename`, the declaration

```
Before Linsolve File Filename proc
```

in the solver input file has the effect that the function will be called just before the default call of linear system solver. The arguments the function can take are fixed and are declared as

```

FUNCTION proc( Model, Solver, A, b, x, n, DOFs, Norm ) RESULT(stat)
  USE SolverUtils
  TYPE(Model_t) :: Model
  TYPE(Solver_t) :: Solver
  TYPE(Matrix_t), POINTER :: A
  REAL(KIND=dp) :: b(:), x(:), Norm
  INTEGER :: n, DOFs, stat
  ...
END FUNCTION proc

```

Here the `Model` structure contains the whole definition of the elmer run. The `Solver` structure contains information for the equation solver from which this linear system originates. The coefficient matrix `A` is in CRS format, `b` is the right-hand side vector, and `x` contains the previous solution. The argument `n` is the number of unknowns, and `DOFs` is the number of unknowns at a single node.

If the return value from this function is zero, the (possibly) modified linear system is solved after the return. If the return value is 1, the linear system is assumed to be already solved and the vector `x` should contain the result. It is noted that the user-supplied Fortran function may also call the default linear equation solver within the function, i.e. one may write the subroutine call

```
CALL SolveLinearSystem( A, b, x, Norm, DOFs, Solver )
```

Here `A` and `b` may be modified so that the linear system which is solved need not be same as the input system.

In a similar way the user may also supply a user-defined Fortran function which will be called just after the solution of linear system. This is done using the declaration

After Linsolve File Filename proc

in the solver input file. The arguments of this function are the same as for a function in the context of Before Linsolve keyword.

Bibliography

- [1] Umfpack home page. <http://www.cise.ufl.edu/research/sparse/umfpack/>.
- [2] W.L. Briggs. *A Multigrid Tutorial*. SIAM, 1987.
- [3] Richard Barrett et al. *Templates for the Solution of Linear Systems: Building Blocks for Iterative Methods*. SIAM, 1993.
- [4] R.W. Freund. A transpose-free quasi-minimal residual algorithm for non-hermitian linear systems. *SIAM J. Sci. Comput.*, 14:470–482, 1993.
- [5] K. Stüben. *Algebraic Multigrid (AMG): An introduction with applications*. GMD – Forschungszentrum Informationstechnik GmbH, 1999.

Chapter 5

Nonlinear System Options

5.1 Introduction

Numerical methods in linear algebra are usually intended for the solution of linear problems. However, there are many problems which are not linear in nature. The nonlinearity may a intrinsic characteristics of the equation, such as is the case with inertial forces in the Navier-Stokes equation. The nonlinearity might also a result of nonlinear material parameters that depend on the solution. What ever the reason for nonlinearity the equations in Elmer are always first linearized to the form

$$A(u_{i-1})u_i = b(u_{i-1}), \quad (5.1)$$

where i refers to the iteration cycle.

How the equations are linearized varies from solver to another. For example, in the Navier-Stokes solver there are tow different methods – the Picard linearization and the Newton linearization that may be used. Also a hybrid scheme where the Picard type of scheme is switched to the Newton kind of scheme when certain criteria are met is available. Therefore this section will not deal with the particular linearization technique of different solver but tries to give some light to the generic keywords that are available. Some keywords may also be defined in the Models Manual related to particular solvers.

In multiphysical simulations there are also a number of keywords related to the solution of coupled systems. Basically one may may define how many times a system of equations is solved repeatedly at maximum and how what are the convergence criteria of the individual solvers that must be met simulataneously.

5.2 Keywords related to solution of nonlinear systems

These keywords are located in the Solver section of each solver, if requited at all.

Nonlinear System Convergence Measure String

The change of solution between two consecutive iterations may be estimated by a number of different measures which are envoked by values `norm`, `solution` and `residual`. The default way of checking for convergence is to test the change of norm

$$\delta = 2 * ||u_i| - |u_{i-1}|| / (|u_i| + |u_{i-1}|). \quad (5.2)$$

This measure is rather liberal since the norm of two solutions may be the same even though the solutions would not. Therefore it is often desirable to look at the norm of change,

$$\delta = 2 * |u_i - u_{i-1}| / (|u_i| + |u_{i-1}|). \quad (5.3)$$

The third choice is to use a backward norm of the residual where the old solution is used with the new matrix.

$$\delta = |Ax_{i-1} - b| / |b|. \quad (5.4)$$

In the current implementation this norm lags one step behind and therefore always performs one extra iteration.

Nonlinear System Norm Degree Integer

The choice of norms used in the evaluation of the convergence measures is not self evident. The default is the L_2 norm. This keyword may be used to replace this by L_n norm where value $n = 0$ corresponds to the infinity (i.e. maximum) norm.

Nonlinear System Norm Dofs Integer

For vector valued field variables by default all components are used in the computation of the norm. However, sometimes it may be desirable only to use some of them. This keyword may be used to give the number of components used in the evaluation. For example, in the Navier-Stokes equations the norm is only taken in respect to the velocity components while pressure is omitted.

Nonlinear System Convergence Absolute Logical

This keyword may be used to enforce absolute convergence measures rather than relative. The default is False.

Nonlinear System Convergence Tolerance Real

This keyword gives a criterion to terminate the nonlinear iteration after the relative change of the norm of the field variable between two consecutive iterations is small enough $\delta < \epsilon$, where ϵ is the value given with this keyword.

Nonlinear System Max Iterations Integer

The maximum number of nonlinear iterations the solver is allowed to do.

Nonlinear System Newton After Iterations Integer

Change the nonlinear solver type to Newton iteration after a number of Picard iterations have been performed. If a given convergence tolerance between two iterations is met before the iteration count is met, it will switch the iteration type instead. This applies only to some few solvers (as the Navier-Stokes) where different linearization strategies are available.

Nonlinear System Newton After Tolerance Real

Change the nonlinear solver type to Newton iteration, if the relative change of the norm of the field variable meets a tolerance criterion:

$$\delta < \epsilon,$$

where ϵ is the value given with this keyword.

Nonlinear System Relaxation Factor Real

Giving this keyword triggers the use of relaxation in the nonlinear equation solver. Using a factor below unity is sometimes required to achieve convergence of the nonlinear system. Typical values range between 0.3 and unity. If one must use smaller values for the relaxation factor some other methods to boost up the convergence might be needed to improve the convergence. A factor above unity might rarely speed up the convergence. Relaxed variable is defined as follows:

$$u'_i = \lambda u_i + (1 - \lambda)u_{i-1},$$

where λ is the factor given with this keyword. The default value for the relaxation factor is unity.

Many of the keywords used to control the Nonlinear System have a corresponding keyword for the Steady State. Basically the operation is similar except the reference value for the current solution u_i is the last converged value of the nonlinear system before starting a new loosely coupled iteration cycle. Otherwise the explanations given above are valid.

Steady State Convergence Measure String

Steady State Norm Degree Integer

Steady State Norm Dofs Integer

Steady State Convergence Tolerance Real

Steady State Relaxation Factor Real

Additionally these keywords are located in the `Simulation` section of the command file.

Steady State Max Iterations Integer

The maximum number of coupled system iterations. For steady state analysis this means it literally, for transient analysis this is the maximum number of iterations within each timestep.

Steady State Min Iterations Integer

Sometimes the coupling is such that nontrivial solutions are obtained only after some basic cycle is repeated. Therefore the user may sometimes need to set also the minimum number of iterations. Sometimes the steady state loop is also used in a dirty way to do some systematic procedures – for example computing the capacitance matrix, or lumped elastic springs. Then this value may be set to an a priori known constant value.

Chapter 6

Integration of time-dependent systems

6.1 Introduction

Solving time-dependent systems is becoming more and more common in various branches of computational science, as the computer resources grow steadily. ElmerSolver may be adapted to solve such systems. The first order time derivatives may be discretized by using the following methods:

- the Crank-Nicolson method
- the Backward Differences Formulae (BDF) of several orders

In the case of the first order BDF scheme adaptive time-stepping strategy may also be used.

The second order time derivatives are approximated by either using the Bossak method or reformulating the second order equations as equivalent systems of first order equations.

In addition, a predictor-corrector scheme is available to automatically control the time steps according to the local truncation error. The analysis and implementation details can be seen in [1].

6.2 Time discretization strategies

Consider the numerical solution of the evolutionary field equation

$$\frac{\partial \phi}{\partial t} + \mathcal{K}\phi = f, \quad (6.1)$$

where the differential operator \mathcal{K} does not involve differentiation with respect to time t and f is a given function of spatial coordinates and time. The spatial discretization of (6.1) leads to the algebraic equations

$$M \frac{\partial \Phi}{\partial t} + K\Phi = F, \quad (6.2)$$

where M , K and F result from the discretization of the identity operator, the operator \mathcal{K} and f , respectively. The vector Φ contains the values of the unknown field ϕ at nodes.

The applications of the first three BDF methods to discretize the time derivative term in (6.2) yield the following systems:

$$\left(\frac{1}{\Delta t} M + K \right) \Phi^{i+1} = F^{i+1} + \frac{1}{\Delta t} M \Phi^i, \quad (6.3)$$

$$\left(\frac{1}{\Delta t} M + \frac{2}{3} K \right) \Phi^{i+1} = \frac{2}{3} F^{i+1} + \frac{1}{\Delta t} M \left(\frac{4}{3} \Phi^i - \frac{1}{3} \Phi^{i-1} \right), \quad (6.4)$$

$$\left(\frac{1}{\Delta t} M + \frac{6}{11} K \right) \Phi^{i+1} = \frac{6}{11} F^{i+1} + \frac{1}{\Delta t} M \left(\frac{18}{11} \Phi^i - \frac{9}{11} \Phi^{i-1} + \frac{2}{11} \Phi^{i-2} \right), \quad (6.5)$$

where Δt is the time step and Φ^i is the solution at time step i . Similarly, F^i is the value of F at time step i .

All the BDF methods are implicit in time and stable. The accuracies of the methods increase along with the increasing order. The starting values for the BDF schemes of order $n > 1$ are computed using the BDF schemes of order 1, ..., $n - 1$ as starting procedures. It should be noted that the BDF discretizations of order $n > 3$ do not allow the use of variable time-step size. Adaptive time-stepping strategy may also be used in the case of the first order BDF scheme.

The adaptive time-stepping is accomplished by first solving the system using a trial time step and then using two time steps the lengths of which equal to the half of that of the trial time step and comparing the results. If the difference between the results is found to be sufficiently small, the use of the trial time step is accepted. Otherwise a new trial time step is defined by dividing the previous trial time step into two steps of equal length and then the procedure is repeated. One may define one's own criterion for determining whether the use of the current time step is accepted. The default criterion is that the norms of the solutions to each system of field equations do not differ more than the given threshold value.

The time discretization of the second order equation

$$M \frac{\partial^2 \Phi}{\partial t^2} + B \frac{\partial \Phi}{\partial t} + K \Phi = F \quad (6.6)$$

using the Bossak method leads to the system

$$\left(\frac{1-\alpha}{\beta(\Delta t)^2} M + \frac{\gamma}{\beta \Delta t} B + K \right) \Phi^{i+1} = F^{i+1} + M \left(\frac{1-\alpha}{\beta(\Delta t)^2} \Phi^i + \frac{\gamma}{\beta \Delta t} V^i + \frac{(1-\alpha)}{2\beta} A^i \right) + B \left(\frac{\gamma}{\beta \Delta t} \Phi^i + \left(\frac{\gamma}{\beta} - 1 \right) V^i + \left(1 - \frac{\gamma}{2\beta} \right) \Delta t A^i \right), \quad (6.7)$$

where

$$\begin{aligned} V^{i+1} &= V^i + \Delta t \left((1-\gamma) A^i + \gamma A^{i+1} \right), \\ A^{i+1} &= \frac{1}{\beta(\Delta t)^2} (\Phi^{i+1} - \Phi^i) - \frac{1}{\beta \Delta t} V^i + \left(1 - \frac{1}{2\beta} \right) A^i, \\ \beta &= \frac{1}{4}(1-\alpha)^2, \quad \gamma = \frac{1}{2} - \alpha. \end{aligned} \quad (6.8)$$

In the following the matrices M and B are referred to as the mass and damping matrix, respectively.

6.3 Keywords related to time discretization

All the keywords related to the time discretization may be given in Simulation section of the solver input file (.sif file). A number of keywords may also be given in Solver section, so that each system of field equations may be discretized using independently chosen time-stepping method. If keywords are not given in the Solver section, the values of the keywords are taken to be those given in the Simulation section. It should be noted that certain keywords such as those controlling the number of time steps, time-step sizes etc. may only be given in the Simulation section.

Timestepping Method *String*

This keyword is used to declare the time discretization strategy for the first order equations. The value of this keyword may be set to be either "BDF" or "Crank-Nicolson" and may be given in either Simulation section or Solver section of the solver input file.

BDF Order *Integer*

This keyword is used to define the order of the BDF method and may take values 1,...,5. This keyword may be given in either Simulation section or Solver section of the solver input file.

Time Derivative Order *Integer*

If a second order equation is discretized, this keyword must be given the value 2 in the Solver section of the solver input file. It should be noted that the second order time derivatives are always discretized using the Bossak method.

Bossak Alpha Real [-0.05]

This keyword is used to define the value for α in the Bossak method used in the time discretization of second order equations. This keyword may be given in either Simulation section or Solver section of the solver input file.

Timestep Intervals Integer array

This keyword is used to define the number of time steps. It may be array-valued so that different time-step lengths may be used for different time intervals of the entire simulation. For example, if one wishes to take first 50 time steps and then to use a different time-step length for the following 100 time steps, one may define

```
Timestep Intervals(2) = 50 100
```

and use the Timestep Sizes keyword to define time-step lengths for the two sets of time steps.

Timestep Sizes Real array

This keyword is used to define the length of time step. If the value of the Timestep Intervals keyword is array-valued, the value of this keyword must also be an array of the same size. For example, if one has defined

```
Timestep Intervals(2) = 50 100
```

the declaration

```
Timestep Sizes(2) = 0.1 1.0
```

sets the time-step length for the first 50 time steps to be 0.1 and for the remaining 100 time steps 1.0.

Timestep Size Real

Instead of using the Timestep Sizes keyword the length of time step may be defined by using this keyword (mind the singular). The value of this keyword is evaluated at the beginning of each time step. A variable time-step length may conveniently be defined using a MATC or Fortran function. This replaces the old keyword Timestep Function.

Output Intervals Integer array

This keyword is used to define the time-step interval for writing the results on disk. As in the case of the Timestep Sizes keyword the size of the value of this keyword must be compatible with that of the Timestep Intervals keyword. The value at a step m is saved if for the corresponding output interval $o \bmod (m-1, o) == 0$. An exception is output interval equal to zero for which output is not saved at all. However, the last step of the simulation is always saved.

Lumped Mass Matrix Logical [false]

The use of a lumped mass matrix may be activated by setting the value of this keyword to be True in the Solver section of solver input file. The default lumping is defined by

$$M'_{ii} = M_{ii} \frac{\sum_j M_{ij}}{\sum_i M_{ii}}. \quad (6.9)$$

The keywords related to the adaptive time-stepping may only be given in the Simulation section of the solver input file. When the adaptive time-stepping strategy is used, a set of trial time steps is defined using the keywords introduced above. The adaptive procedure is executed for each of these trial steps. Note that the adaptive time-stepping is possible only in the case of the first order BDF scheme.

Adaptive Timestepping Logical [false]

The value of this keyword must be set to be True if the adaptive time integration is to be used.

Adaptive Time Error Real

This keyword is used to define the threshold value for the criterion for determining whether the use of the current time step is accepted.

Adaptive Error Measure Real

Using this keyword one may define one's own measure for evaluating the difference between the computed results. This measure and the threshold value, which is given using the Adaptive Time Error keyword, may be used to define a user-defined criterion for determining whether the use of the current time step is accepted. The value of the Adaptive Error Measure keyword is evaluated twice for each trial time step. For the first time the value of the keyword is evaluated after the system is solved using the trial time step. The second time is after the system is solved using two time steps the lengths of which equal to the half of that of the trial time step. The absolute value of the relative difference between these two values is compared to the threshold value given by the Adaptive Time Error keyword to determine whether the use of the current time step is accepted. If several systems of field equations are solved, all the solutions must satisfy the similar criterion. If this keyword is not used, the default criterion is based on comparing the norms of the solution fields.

Adaptive Min Timestep Real

Using this keyword one can limit the subsequent division of the trial time steps by giving the minimum time-step length which is allowed.

Adaptive Keep Smallest Integer [1]

By default the adaptive scheme tries to double the length of the time step after the acceptable time step is found. If a value $n > 1$ is given for this keyword, the adaptive scheme tries to increase the step length after taking n steps which are at most as long as the step length accepted.

6.4 On the treatment of time derivatives in Elmer Solver code

In the following a number of issues that may be useful if one is writing a code to solve one's own application are explained.

By default Elmer Solver does not generate or use global mass or damping matrices in the solution of time-dependent systems. Mass and damping matrices need to be computed only element-wise, as the linear system resulting from the time discretization, such as (6.3), is first formed element-wise and this local contribution is later assembled to the global system. In the case of the first order equation (6.2) the local linear system may be formed by using the subroutine call

```
CALL Default1stOrderTime( M, K, F ),
```

where M is the element mass matrix, K is the element stiffness matrix and F is the element force vector. In a similar manner, in the case of the second order equation (6.6) one may use the subroutine call

```
CALL Default2ndOrderTime( M, B, K, F ),
```

where B is the element damping matrix.

Note that these subroutines must also be called for the local matrices and vectors that result from the discretization of neumann and newton boundary conditions. If the boundary conditions do not contain any time derivatives, the M and B matrices should be set to be zero before calling the above subroutines.

If the global mass matrix is required, it may be generated by using the subroutine call

```
CALL DefaultUpdateMass( M )
```

Similarly, the global damping matrix may be generated by using the subroutine call

```
CALL DefaultUpdateDamp( B ).
```

Global mass (and possibly damping) matrices are required, for example, in the solution of eigenvalue problems. One may also implement one's own time-stepping scheme at the global level using these matrices.

6.5 Predictor Corrector Scheme

The keyword `Timestep Intervals` is used to set the number of steps and the keyword `Timestep Sizes` is used to set the initial step size only. The step size will be changed automatically by the predictor-corrector scheme based on an error estimator and PI controller. The solver used for predictor-corrector error estimator is set by having

```
Exec Solver = "Predictor-Corrector"
```

There can be more than one solver with predictor-corrector scheme, but only the one with the smallest index will be used for the error estimate and the time step will be changed accordingly. In order to use a predictor-corrector scheme, one keyword should be given in the Simulation section and several keywords may also be given in the Solver section.

`Predictor-Corrector Control Logical`

The predictor-corrector scheme may be activated by setting the value of this keyword to be `True` in the Simulation section.

`Skip First Timestep Logical`

This keyword has to be set to `True` in the Solver section for most cases to get the solutions from the other solvers which are used as the initial condition for the predictor-corrector solver. If this keyword is `false`, the predictor-corrector solver will use the values from the Initial Condition section directly.

`Predictor method String`

This keyword is used to define the time scheme for the predictor in the Solver section. Currently, the value may only be `Adams-Bashforth`. But, there will be more options in the future release.

`Corrector method String`

This keyword is used to define the time scheme for the corrector in the Solver section. Currently, the value may only be `Adams-Moulton`. But, there will be more options in the future release.

`Predictor-Corrector Scheme Order Integer[2]`

This keyword is used to define the order of the predictor-corrector scheme. The value can be set to 1 or 2 which corresponds the first or second order semi-implicit schemes mentioned in [1].

`Predictor-Corrector Control Tolerance Real[1e-6]`

This keyword is used to define the tolerance ϵ for the PI control process in [1]. The default value is $1e-6$. A smaller value puts more restriction on the accuracy which may reduce the step sizes. A larger value of the tolerance allows larger step sizes, but the solution may not be stable in some cases.

`Predictor-Corrector Save Error Logical[false]`

This keywords is used to save the time steps and local truncation errors for each step.

The following two keywords are optional. They are related to the PI controller used for the time step control. The default values are given as in [1] for first and second order scheme.

`Predictor-Corrector Control Beta 1 Real`

`Predictor-Corrector Control Beta 2 Real`

Bibliography

- [1] Gong Cheng, Per Lötstedt, and Lina von Sydow. Accurate and stable time stepping in ice sheet modeling. *Journal of Computational Physics*, 329:29–47, 2017.

Chapter 7

Solving eigenvalue problems

7.1 Introduction

Eigenvalue problems form an important class of numerical problems, especially in the field of structural analysis. Also some other application fields may have eigenvalue problems, such as those in density functional theory. This manual, however, introduces eigenvalue computation in Elmer using terminology from elasticity.

Several different eigenvalue problems can be formulated in elasticity. These include the “standard” generalized eigenvalue problems, problems with geometric stiffness or with damping, as well as stability (buckling) analysis. All of the aforementioned problems can be solved with Elmer. The eigenproblems can be solved using direct, iterative or multigrid solution methods.

7.2 Theory

The steady-state equation for elastic deformation of solids may be written as

$$-\nabla \cdot \tau = \vec{f}, \quad (7.1)$$

where τ is the stress tensor. When considering eigen frequency analysis, the force term \vec{f} is replaced by the inertia term,

$$-\nabla \cdot \tau = -\rho \frac{\partial^2 \vec{d}}{\partial t^2}, \quad (7.2)$$

where ρ is the density.

The displacement can now be assumed to oscillate harmonically with the eigen frequency ω in a form defined by the eigenvector \vec{d} . Inserting this into the above equation yields

$$-\nabla \cdot \tau(\vec{d}) = \omega^2 \rho \vec{d}, \quad (7.3)$$

or in discretized form

$$Ku = \omega^2 Mu, \quad (7.4)$$

where K is the stiffness matrix, M is the mass matrix, and u is a vector containing the values of \vec{d} at discretization points. The equation (7.4) is called the generalized eigenproblem.

Including the effects of pre-stresses into the eigenproblem is quite straightforward. Let us assume that there is a given tension field σ in the solid. The tension is included by an extra term in the steady-state equation

$$-\nabla \cdot \tau - \nabla \cdot (\sigma \nabla u) = \vec{f}. \quad (7.5)$$

The pre-stress term includes a component K_G to the stiffness matrix of the problem and thus the eigenvalue equation including pre-stresses is

$$(K + K_G)u = \omega^2 Mu. \quad (7.6)$$

The pre-stress in Elmer may be a known pre-tension, due to external loading or due to thermal stress, for example. The stress tensor containing the pre-stresses σ is first computed by a steady-state analysis and after that the eigenvalue problem is solved. It should be noted though that the eigenvalue problem in a pre-stressed state is solved using first order linearization, which means that the eigenvalues are solved about the non-displaced state. If the pre-loading influences large deformations the eigenvalues are not accurate.

The eigenvalue problem with pre-stresses may be used to study the stability of the system. Some initial loading is defined and a pre-stress tensor σ is computed. This tensor is then multiplied by a test scalar λ . The critical load for stability, or buckling, is found by setting the force on the right hand side of the equation (7.5) equal to zero. The problem then is to solve λ from

$$Ku = -\lambda K_G u, \quad (7.7)$$

which again is formally an eigenvalue problem for the test parameter. The critical loading is found by multiplying the given test load with the value of λ . In other words, if $\lambda > 1$ the loading is unstable.

7.2.1 Damped eigenvalue problem

Finally, let us consider the damped eigenproblem, also called quadratic eigenvalue problem. In this case there is a force component proportional to the first time derivative of the displacement in addition to the inertia term

$$-\nabla \cdot \tau = -\delta \frac{\partial \vec{d}}{\partial t} - \rho \frac{\partial^2 \vec{d}}{\partial t^2}, \quad (7.8)$$

where δ is a damping coefficient. The problem is transformed into a more suitable form for numerical solution by using a new variable \vec{v}' defined as $\vec{v}' = \frac{\partial \vec{d}}{\partial t}$. This yields

$$-\nabla \cdot \tau = -\delta \vec{v}' - \rho \frac{\partial \vec{v}'}{\partial t}. \quad (7.9)$$

Working out the time derivatives and moving into the matrix form, the equation may be written as

$$Ku = -Dv - i\omega Mv, \quad (7.10)$$

or,

$$i\omega \begin{pmatrix} I & 0 \\ 0 & M \end{pmatrix} \begin{pmatrix} u \\ v \end{pmatrix} = \begin{pmatrix} 0 & I \\ -K & -D \end{pmatrix} \begin{pmatrix} u \\ v \end{pmatrix}, \quad (7.11)$$

where i is the imaginary unit, D is the damping matrix, and v a vector containing the values of \vec{v}' at the discretization points. Now the damped eigenproblem is transformed into a generalized eigenproblem for complex eigenvalues.

Finally, to improve the numerical behavior of the damped eigenproblem, a scaling constant s is introduced into the definition $s\vec{v}' = s \frac{\partial \vec{d}}{\partial t}$. In the matrix equation (7.11) this influences only the identity matrix blocks I to be replaced by sI . Good results for numerical calculations are found when

$$s = \|M\|_\infty = \max |M_{i,j}|. \quad (7.12)$$

7.3 Keywords related to eigenvalue problems

An eigenvalue analysis in Elmer is set up just as the corresponding steady-state elasticity analysis. An eigenvalue analysis is then defined by a few additional keywords in the Solver section of the `sif` file. The solver in question can be linear elasticity solver called Stress Analysis, linear plate elasticity solver, or even nonlinear elasticity solver, though the eigen analysis is, of course, linear.

Many of the standard equation solver keywords affect also the eigen analysis, *e.g.* the values given for Linear System Solver and Linear System Iterative Method in case of an iterative solver. More information about these settings is given in this Manual under the chapter concerning linear system solvers. The specific keywords for eigen analysis are listed below

Eigen Analysis Logical

Instructs Elmer to use eigensystem solvers. Must be set to True in all eigenvalue problems.

Eigen System Values Integer

Determines the number of eigen values and eigen vectors computed.

Eigen System Select String

This keyword allows the user to select, which eigenvalues are computed. The allowable choices are

- Smallest Magnitude
- Largest Magnitude
- Smallest Real Part
- Largest Real Part
- Smallest Imag Part
- Largest Imag Part

Smallest magnitude is the default.

Eigen System Shift Real

Gives an offset in the eigenvalues. May be useful if the user is interested in eigenvalues in a specific pre-defined frequency-regime, for example. In practice the shift ω_0^2 is achieved by defining $K := K - \omega_0^2 M$ and solving the eigenvalues of this modified system.

Eigen System Convergence Tolerance Real

The convergence tolerance for iterative eigensystem solver. The default is 100 times Linear System Convergence Tolerance.

Eigen System Max Iterations Integer

The number of iterations for iterative eigensystem solver. The default is 300.

Eigen System Complex Logical

Should be given value True if the eigen system is complex, *i.e.* the system matrices are complex. Not to be given in damped eigen value analysis.

Geometric Stiffness Logical

Defines geometric stiffness (pre-stress) to be taken into account in eigen analysis. This feature is only available with linear bulk elasticity solver.

Stability Analysis Logical

Defines stability analysis. This feature is only available with linear bulk elasticity solver.

Eigen System Damped Logical

Defines a damped eigen analysis. Damped eigen analysis is available only when using iterative solver.

Eigen System Use Identity Logical

If True defines the relation displacement and its derivative to be $s\vec{v}' = s\frac{\partial \vec{d}}{\partial t}$. The other possibility is to use $M\vec{v} = i\omega M\vec{u}$. The default is True.

Eigen System Lanczos Vectors Integer

Number of Lanczos vectors that are used to compute the eigenvalues. The default is $3N_e + 1$, where N_e is the number of eigenvalues.

Eigen System Compute Residuals Logical

Computes the residuals of the eigen value system.

7.4 Constructing matrices M and D in Solver code

In eigen analysis the mass matrix M and the damping matrix D have to be separately constructed. Usually in Elmer the different matrices are summed into a single matrix structure, since the final linear equation is of the form $Ax = b$, and there is no need for separate values of the mass matrix and the stiffness matrix.

The matrix is represented in Elmer using compressed row storage (CRS) format, as explained in chapter about Linear system solvers. The matrix structure holds also vectors for the values of the mass and damping matrices

```
TYPE Matrix_t
...
REAL(KIND=dp), POINTER :: MassValues(:), DampValues(:)
...
END TYPE Matrix_t
```

These arrays use the same Rows and Cols tables than the normal Values array.

The mass and damping matrices are constructed elementwise in a similar manner as the stiffness matrix. After each element the local contributions are updated to the equation matrices by the following subroutine calls

```
CALL DefaultUpdateEquations( STIFF, FORCE )

IF ( Solver % NOFEigenValues > 0 ) THEN
  CALL DefaultUpdateMass( MASS )
  CALL DefaultUpdateDamp( DAMP )
END IF
```

In this segment of code the variables STIFF, MASS, DAMP and FORCE store the local values of the stiffness matrix, the mass matrix, the damping matrix, and the right hand side of the equation, respectively. The integer NOFEigenValues in the Solver data structure gives the number of eigenvalues requested. Here it is used as an indicator of whether the mass and damping matrices need to be constructed.

The eigenvalues and eigenvectors are stored in the arrays Solver % Variable % EigenValues and Solver % Variable % EigenVectors,

```
TYPE Variable_t
...
COMPLEX(KIND=dp), POINTER :: EigenValues(:)
COMPLEX(KIND=dp), POINTER :: EigenVectors(:, :)
...
END TYPE Matrix_t
```

and Solver % Variable % EigenVectors(i, :) gives the eigenvector corresponding to the eigenvalue i .

Chapter 8

Dirichlet boundary conditions and limiters

Except for special cases, basically any partial differential equation (PDE) model needs boundary conditions. Often their definition is solver-specific in Elmer. This chapter focuses on constraints which are referred to as the Dirichlet conditions.

8.1 Dirichlet constraints

In the context of variational formulations of PDE models different kinds of boundary conditions are met. In connection with the second-order PDE models, boundary conditions of Dirichlet type refer to cases where the value of the field is set on the boundary. In engineering literature these are also known as the essential boundary conditions. The other types of boundary conditions involve derivatives of the solution and often correspond to prescribing a flux or force on the boundary. They are referred to as the Neumann boundary conditions or natural boundary conditions. More general boundary conditions of Robin type are also possible. For problem-specific examples of them the reader is directed to the Models Manual.

Technically the Dirichlet conditions in ElmerSolver are set through manipulating only the values of the linear system entries rather than the structure of the linear system. To be more specific, in setting the degree of freedom with index i , the i th row of the matrix is set to be zero, except for the diagonal entry which is set to be unity. When also the corresponding entry of the right-hand side (RHS) vector is set to have the desired value, the solution will satisfy the Dirichlet condition.

A handicap of the procedure described is that the symmetry of the original matrix is not maintained. This may affect the applicability of linear system solvers. There are two remedies to ensure that the matrix remains symmetric. The column may also be zeroed except for the diagonal entry and the contribution of the known values may be subtracted from the RHS vector. The second option is to eliminate all the rows and columns related to the known values. This reduces the size of the matrix but has an additional cost as a secondary matrix is created and the values are copied into it.

The standard Dirichlet conditions are given by referring to objects which have a lower dimension than the leading dimension in the geometry definition, i.e. for 3D problems values are usually fixed only at 2D faces. The Dirichlet conditions may hence be set for the degrees of freedom of existing boundary elements. Additionally Dirichlet-like conditions may be set for a set of nodes that is created on-the-fly. It is possible also to set Dirichlet-like conditions for bodies also. This may be particularly useful when the constraint is only conditional.

Sometimes a Dirichlet condition should depend on other variables in a way which defines whether or not to apply the condition at all. For example, the temperature at a boundary might only be defined in the case of inflow, while for outflow giving the temperature would not be physically justified. For this kind of purposes the user may define an additional real-valued variable that depends on the variable itself. Then, the Dirichlet condition can be applied conditionally in cases where the additional variable has a positive value.

Simulation

Set Dirichlet BCs by BC Numbering Logical

By default the Dirichlet conditions are set in the order that they just happen to appear in the boundary element list. This is fine when there are no conflicting boundary conditions. However, when such is the case the user may want to control the order of the conditions – the last Dirichlet condition always prevails. This flag enforces the BCs to be set in the order of boundary condition numbering in the `sif` file.

Boundary Condition `bc id`

Target Boundaries(n) Integer

The set of boundaries on which the Dirichlet conditions will be applied.

Target Nodes(n) Integer

Sets pointwise conditions on-the-fly. These points are referred via the absolute indexing of the nodes.

Target Coordinates(n,DIM) Real

Sets pointwise conditions via transforming the coordinate values into nodal indices corresponding to the nearest nodes at the time of the first call. Target groups defined by `Target Boundaries`, `Target Nodes`, and `Target Coordinates` should not reside in the same boundary condition definition.

Varname Real

The value of each variable which has an active equation associated may be set by giving its value in a boundary condition section. If the variable name is not listed in the keyword listing, the user should also define the value type which is always `Real`.

Varname i Real

For multicomponent fields the Dirichlet condition may be set for each field component separately.

Varname Condition Real

The Dirichlet condition related to the variable is set to be active only when the condition variable is positive.

Varname {e} i Real

This keyword is used to define the vector \mathbf{g} so that its tangential trace $\mathbf{g} \times \mathbf{n}$ is approximated by $\mathbf{u}_h \times \mathbf{n}$, with \mathbf{u}_h the edge element interpolating function corresponding to the solver variable. The value of this keyword defines the components g_i , $i \in \{1, 2, 3\}$, with respect to the global Cartesian coordinate system. For the edge finite elements see the appendix F.

Varname {e} Real

If the edge finite element has only degrees of freedom which are associated with element edges, this keyword can be used to define a scalar field \hat{g}_τ so that the values of the prescribed degrees of freedom are computed as (cf. the appendix F)

$$D = \int_{\mathcal{E}_{ij}} [\mathbf{g}(\mathbf{x}) \cdot \boldsymbol{\tau}(\mathbf{x}) + \hat{g}_\tau] \phi(\mathbf{x}) ds,$$

with \mathcal{E}_{ij} an element edge. This keyword should not be used if additional degrees of freedom associated with the element faces are present (that is, this feature should not be used for edge elements of the second degree or edge elements from the optimal family).

The Dirichlet-like conditions for bodies. It is also possible to set the values of an exported variable with exactly the same logic.

Body Force `body force id`

Varname Real

The setting of Dirichlet-like conditions for the whole body follows the same logic as for the boundaries. When the body force definition is assigned to a body, the variable values will be fixed as defined.

Varname Condition Real

The conditional Dirichlet-like condition may also be given for bodies.

The following may be used in conjunction with the Dirichlet conditions to affect the setup of the linear system equation.

Solver solver id

Linear System Symmetric Logical True

Make the matrix symmetric by eliminating the known values from the RHS and by zeroing the appropriate matrix entries.

Before Linsolve "EliminateDirichlet" "EliminateDirichlet"

Creates the secondary matrix of a reduced size by eliminating Dirichlet conditions and passing this to the linear system solver.

8.2 Soft limiters

The user may set soft lower and upper limits to the values of the field variable. For example, concentration can never be negative and hence a zero lower limit could be applied for it. The limits are applied in an iterative manner defining a contact set where Dirichlet-like conditions are applied. A node is included in the contact set when its value passes the limits, and a node is released from the contact set when the nodal load related to it is negative (or positive). The limiters may be applied to both boundary conditions and bodies.

Solver solver id

Apply Limiter Logical True

Activates the search for limiters and also activates the computation of nodal loads.

Limiter Value Tolerance Real

Defines the tolerance of the field value by which a node is added to the contact set.

Limiter Load Tolerance Real

Defines the tolerance of the nodal load to determine whether a node is removed from the contact set.

Limiter Load Sign Negative Logical

The default sign used for determining the contact set is derived from the default discretization of the Poisson equation. If the equation would be multiplied by -1 , the load would also change its sign. With that in mind, a possibility to influence the rules is given with this keyword.

Boundary Condition bc id

Varname Lower Limit Real

The lower limit of field variable.

Varname Upper Limit Real

The upper limit of field variable.

Body Force body force id

Varname Lower Limit Real

The lower limit of field variable.

Varname Upper Limit Real

The upper limit of field variable.

Chapter 9

Enforcing shared conditions between boundaries

Special classes of boundary conditions arise from treating constraints that involve two boundaries simultaneously. For example enforcing the periodicity of the solution is considered here. We also consider conditions relating to the application of the mortar finite element method, which offers a general strategy to approximate the solution continuity across a shared boundary when the meshes of two independently discretized bodies are not conforming on the shared boundary (that is, the locations of the mesh nodes do not coincide on the shared boundary).

Historically the implementation of periodic boundary conditions in Elmer has been based on projectors which are derived directly using the interpolation scheme. However, over the time the development has evolved towards mortar finite element strategies which lead to projectors based on weak formulations. Currently both approaches are still used but the mortar finite element approach should be employed when striving for the best possible accuracy over nonconforming meshes.

9.1 Periodic conditions

Periodic BCs may be considered to be a special case of Dirichlet conditions where the fixed value is given as a linear combination of other unknown values. The periodic boundary conditions in Elmer are very flexible. In fact they may even be antiperiodic.

Boundary Condition `bc id`

Periodic BC `Integer`

Periodic boundaries generally come in pairs. This command defines a pointer to another boundary condition definition which is considered to be the counterpart.

Periodic BC `Explicit Logical`

Sometimes the implicit periodic BCs (the default) leads to convergence problems. It also complicates the matrix structure by adding additional connections. Then the explicit type of periodic conditions may be a good alternative. Note that it requires iteration even for a linear operator.

Periodic BC `Translate(3) Real`

The correspondence of the periodic boundaries should be obtained by performing three different operations: translation, rotation and scaling. This generality is not usually needed and therefore default values are usually appropriate. For translation the default translation vector is obtained via moving in the normal direction of the primary boundary until the target boundary is hit. If this is not desired, the user may give another translation vector using this keyword.

Periodic BC `Rotate(3) Real`

By default no rotation is performed prior to the construction of the constraints. This keyword may be used to give the angles of rotation. There is no automatic rotation, but there is some detection of the correctness of the rotation.

`Periodic BC Scale(3) Real`

This keyword may be used to give a scaling vector if this is desired. If this is not given, an isotropic scaling will be used to map the bounding boxes of the two boundaries to the same size.

`Periodic BC VarName Logical`

The user should define the variables that are to be periodic in nature. This is done by attaching their names into logical expressions following the string `Periodic BC`.

`Anti Periodic BC VarName Logical`

The variable may be also antiperiodic, i.e. the absolute value is the same but the sign is different. Then this should be used instead.

`Periodic BC Use Lagrange Coefficients Logical`

By default the periodic BCs are handled by elimination. The other alternative is to use a Lagrange multiplier. This strategy is enforced by setting this keyword to be `True`.

9.2 Mortar conditions

The motivation for implementing mortar finite element framework into Elmer came from the simulation of rotating electrical machines. There the continuity of solution between rotor and stator must be approximated. Both parts usually have their own fixed meshes that move with respect to each other. The principal complication associated with the moving mesh approach is that the finite element meshes are then generally non-matching across the model boundary which represents the interaction surface Γ shared by the rotating and surrounding bodies. The usual continuity requirements to obtain a conforming finite element approximation over the entire domain cannot thus be satisfied. To this aim a monolithic discretization strategy based on the mortar finite element method has been implemented.

The implementation of the mortar conditions in Elmer is related to the periodic boundary conditions in that some common routines and techniques are shared. The mortar conditions may be dynamic, whereas the periodic conditions are assumed to be fixed.

The mortar finite element formulations approximate the solution continuity via a weak formulation and, in the case of the scalar unknown, lead to evaluating integrals of the type

$$\int (u_1 - u_2) \phi_S d\Gamma_S = 0 \quad (9.1)$$

where the integration is over the element(s) on the slave (non-mortar) side, u_1 the value of the matched quantity at the slave side and u_2 the quantity at the master (mortar) side. In addition, ϕ_S is a test function. The integration is accomplished in practice by searching the locations of the integration points within the mortar elements. If the set of the test functions coincides with the trace space of admissible solutions on the non-mortar side, the degrees of freedom at each side are related via the linear equations

$$Mu_1 = Nu_2, \quad (9.2)$$

with

$$\begin{aligned} M_{ij} &= \int \phi_S^j \phi_S^i d\Gamma_S, \\ N_{ij} &= \int \phi_M^j \phi_S^i d\Gamma_S. \end{aligned} \quad (9.3)$$

All the integration points of the non-mortar side must be locatable in the set of mortar elements. Therefore the mortar boundary may be larger than the non-mortar boundary but never vice versa. To enforce the constraint in the original system a constraint matrix block

$$C = [M \quad -N] \quad (9.4)$$

is constructed. This constraint block is then added to the original system by using the Lagrange multipliers as

$$\begin{bmatrix} A & C^T \\ C & 0 \end{bmatrix} \begin{bmatrix} u \\ \lambda \end{bmatrix} = \begin{bmatrix} f \\ 0 \end{bmatrix}. \quad (9.5)$$

Re-parametrization of the shared boundary

When we are studying problems where the potential contact surface (or line) is known, a re-parametrization of the surface meshes may be useful. For example, in the case of a cylindrical surface the element faces may be mapped from (x, y, z) coordinates to (θ, z) coordinates since the r -coordinate should be the same by construction. Finding the non-mortar points from the mortar side becomes then somewhat easier as we avoid the problem of two mismatching polygons. Similarly, (r, z) coordinates may be used instead of (x, y, z) coordinates when the solutions on two boundaries which are planes in the (r, z) coordinates are matched. The similar strategy works also for constraints that need to be expressed at certain coordinate level. For example, we may use the (x, y) -plane to express the constraints instead of the full 3D coordinates.

When using the reduction to (θ, z) plane, we may also consider periodicity of the projector. Assuming that the mortar nodes are in interval $[\theta_{min}, \theta_{max}]$ such that $2\pi/(\theta_{max} - \theta_{min})$ is an integer we may always map the angle ϕ to the existing interval from anywhere in $[0, 2\pi]$. Similar repeated projector could basically also be taking place in translation between two planes if their structure would be periodic.

Integration of mortar projector

When performing the integration, there are several ways how the integration and search is carried out. Which one is optimal depends greatly on how the mesh is configured.

If the mesh is basically conforming and the need for the projector is caused by the internally created discontinuity (using keyword `Discontinuous Boundary`), we have the luxury that there is an exact correspondence between a master and a slave element. Then an exact projector may be computed with a minimal effort. By default the projector is of mass matrix type but also a diagonal projector may be created by mass lumping. This branch of projectors is invoked if the corresponding `Mortar BC` is given the value zero. This approach works also in parallel, at least if the boundary halo is saved during the partitioning process.

In some cases of rotating machines created by extrusion we have the luxury that the nodes are only at specific coordinate levels. This makes it easier to have an accurate integration over the intersection of the elements. In the current implementation of `Level Projector` it is assumed that the surface elements are quadrilaterals and two of the element sides remain parallel. The integration is performed by checking whether any of master and slave elements cross each other and creating a temporal triangulation for their intersection. This is a loop of complexity $O(N^2)$ but for typical cases it is still not a bottle-neck. This approach may work in parallel when using a tailored partitioning algorithm that distributes the elements such that they by construction will be in the same partitioning as their counterparts.

A generic approach is used if we have no simplifications that could be utilized when integrating over the non-mortar elements. The integration utilizes oct-tree based search algorithm. Unfortunately the integration is not accurate (a temporal triangulation for improved integration accuracy is not generated in this case) which may result to accuracy and convergence problems when treating the resulting constrained linear system. Currently there is no way to perform parallel communication when performing the integration.

Biorthogonal basis for mortar projectors

There is a possibility to use biorthogonal basis functions when creating the mortar projector with some integrators. The biorthogonal basis is advantageous since either the non-mortar or mortar degrees of freedom may then be eliminated from the linear system, thereby significantly simplifying the solution.

Contact mechanics

ElmerSolver library includes some generic features for contact mechanics. The functionality relies heavily on contact mortar projectors which are used to ensure the match on the contact surface. The contact pressure, on the other hand, is used to determine the active contact set quite similarly as for the case of soft limiters. The contact features are applicable to all elasticity equations having the displacement degrees of freedom at each node. The contact mechanics is still under development and the documentation is therefore not complete. There are severe limitations when it comes to parallel operation and conflicting contact conditions on different boundaries.

Keywords for mortar projector

Solver `solver id`

Apply Mortar BCs `Logical`

Enforces the application of mortar BCs. This is not by default set to be `True` partly because at the time of writing this machinery is still under development. If this flag is set to be true, then the mortar BCs are applied to the variable of this solver. Otherwise the mortar projectors are not applied to the solver variable.

Edge Basis `Logical`

Activates the use of edge element basis functions for the mortar conditions.

Boundary Condition `bc id`

Mortar BC `Integer`

This is set for the non-mortar (slave) boundary to refer to the boundary condition index of the corresponding mortar boundary. The mortar boundaries always come in pairs, but the settings are only given in the context of the non-mortar boundary.

Rotational Projector `Logical`

The interface lying in 3D space may be flattened to 2D plane by several ways. Using the flattened geometry results to better interpolation accuracy since faceting is not needed. This keyword activates the use of the (θ, z) coordinate system when computing the projector. The angle may have a discontinuity around 180 degrees. The `Level Projector` option offers a special way to deal with cases that make a full circle.

Cylindrical Projector `Logical`

Similar to `Rotational Projector` except there is automatic fitting to get a cylinder surface such that the cylinder axis is determined via computation. Also, if the angle is not full 2π , it will not be enforced.

Anti Rotational Projector `Logical`

As the previous one but if the module or the periodic cycle is odd use a negative coefficient in enforcing the periodicity. This way the mortar projector becomes antiperiodic.

Sliding Projector `Logical`

Perform dimensional reduction so that only two of the coordinates are accounted for. By default this means mapping to (x, y) plane.

Anti Sliding Projector `Logical`

As the sliding projector except if the mapping of the sliding projector cycle is odd use a negative coefficient in the mortar projector.

Radial Projector `Logical`

Use the (r, z) coordinate system when computing the projector.

Mortar BC Scaling `Real`

The coefficient with which the values are scaled when enforcing continuity over the mortar interface. The default is one.

Level Projector `Logical`

Level projector assumes that the re-parametrized surfaces consist of quadrilaterals with two parallel edges at constant z -levels. This projector may use either strong (interpolation based) or weak strategies to find the constraints for edge and nodal degrees of freedom. There are many additional keywords that apply to the level projector.

Plane Projector `Logical`

This is similar to `Level Projector` except the plane over which the projection will be made is fitted automatically such that the best fit of a plane associated with either slave or master boundary will be used.

Level Projector Nodes `Strong Logical`

This uses strong (interpolation based) projector for the nodal degrees of freedom. The default is `False`, i.e. a weak projector is used.

Level Projector Plane Edges Strong Logical

For extruded geometry some edges will always be aligned with the plane. For these edges it seems natural to use only the other edges on the same constant z-plane to create the projector. Then this keyword should be set `True`. The default is `False`.

Level Projector Extruded Edges Strong Logical

The extruded edges are the ones that are created when the 2D mesh is extruded into the third direction. So they are aligned with the z-axis. This keyword controls how the edge element degrees of freedom are treated in the z-direction. Using the default value `False` is recommended.

Projector Skip Nodes Logical

The level projector may skip to create the projector for the nodal degrees of freedom. This would also be the choice when there are no nodal degrees of freedom for the solver. Also there could be a physical reason that would make them obsolete.

Projector Skip Edges Logical

Also the edge element degrees of freedom may be skipped. They will be treated by default if there are degrees of freedom associated with the edges.

Save Projector Logical

Save the projector values to an external file for mainly debugging purposes. Default is `False`.

The following keywords are related to the use of biorthogonal basis.

Biorthogonal Basis Logical [False]

This will enable the elimination of slave degrees of freedom by making the constraint matrix block corresponding to the slave entries diagonal.

Biorthogonal Dual Master Logical [True]

This option also makes the constraint matrix block corresponding to the slave entries diagonal. The value of this keyword defines whether the nodal basis or the biorthogonal basis should be used on the master side. The default is the latter.

Biorthogonal Dual LCoeff Logical [False]

Here nodal basis functions are used for expressing the traces of the fields to be matched, while the Lagrange multiplier is treated with the biorthogonal basis functions.

The following keywords are related to the use of contact mechanics.

Solver solver id

Apply Contact BCs Logical

Enforces the solver to apply contact mechanics conditions for the boundaries where the flag is given. When doing contact mechanics, it may be useful to set `Linear System Residual Mode = Logical True` so that only the difference is solved making the solution process more robust.

Boundary Condition bc id

Contact Depth Offset Initial Real

Sometimes the contact set is initially empty which may result to rigid body motion making the system indefinite. Therefore it may advantageous to introduce some initial offset to the contact set detection such that the initial contact set is increased for the first iteration.

Contact Depth Offset Real

This has the same effect as the previous one but is used for each iteration. This could be used to model a tight fit of cylinders, for example.

Contact Active Condition Real

This keyword may be used to enforce active contact somewhere. For positive values the contact is active, otherwise the flag has no effect.

`Contact Passive Condition Real`

This keyword may be used to enforce no-contact somewhere. For positive values the contact is passive, otherwise the flag has no effect.

`Stick Active Condition Real`

This keyword may be used to enforce active stick condition somewhere. For positive values the contact is active, otherwise the flag has no effect.

`Stick Passive Condition Real`

This keyword may be used to enforce slip contact somewhere. For positive values the slip contact, otherwise the flag has no effect.

`Static Friction Coefficient Real`

This keyword is used to give the static friction coefficient on a surface. The magnitude of the tangential force will then be the friction coefficient times the normal contact force.

`Dynamic Friction Coefficient Real`

This keyword is used to give the dynamic friction coefficient on a surface.

`Stick Contact Coefficient Real`

This keyword is used to give the stick coefficient that allows a small tangential movement even if stick model is used.

`Contact Velocity(dim) Real`

For steady state contact the velocity cannot be retrieved from history. Then the user may give the relative velocity between the slave and master surfaces. The velocity magnitude is not important for linear friction models but the direction of the velocity is used to set the direction of the friction forces.

`Normal-Tangential VarName Logical`

If the normal of the surface is not suitably aligned with some coordinate direction, it is recommended to use a normal-tangential coordinate system except for the tie contact. Here the Varname would typically be Displacement.

`Contact Active Set Minimum Integer`

If the contact set is empty, there may be numerical problems. Therefore a minimum size for the contact set may be defined. Then the requested number of nodes with the shortest distance are added to the contact set.

`Tie Contact Logical`

Contact type where the contact set is active on the whole surface. Hence there is no testing performed on the direction of the contact pressure.

`Stick Contact Logical`

Contact type where the tangential contact set is the same as the normal contact set.

`Slide Contact Logical`

Contact type where the tangential contact set is empty.

`Friction Contact Logical`

Contact type where the tangential contact set is tuned after some conditions. This flag is not fully operational yet.

`Contact Type String [stick, tie, friction, slide]`

Instead of using the above flags to set the contact type, the user may also give a string that chooses between the different options.

Chapter 10

Linear systems with constraints

10.1 Introduction

There are a number of cases where a set of linear equations under additional constraints has to be solved. For example, in conjunction with the mortar finite element methods the application of the mortar projection introduces additional constraints that are dealt with introducing Lagrange multipliers.

There are two different ways to treat the linear constraints. The first one is to use a large monolithic system where new constraint equations are added. This is generally applicable but may result in numerically difficult problems. The other approach uses elimination to reduce the size of the system. It is numerically favorable method but has some limitations related to its generality.

10.1.1 Solving linear systems with constraints

The Lagrange multiplier procedure leads to a matrix that is indefinite. Any Krylov method capable of handling indefinite matrices should be applicable to the constrained system. In practice the GCR and BiCGStab(L) methods in Elmer has been used successfully in applications. The constraint equations can be effectively preconditioned using ILU (for example).

In the case of (9.5) the current implementation in Elmer uses a preconditioning matrix

$$P = \begin{bmatrix} P_0 & C^T \\ C & 0 \end{bmatrix} \quad (10.1)$$

For the ungauged edge element discretization of the A-V formulation we select

$$P_0 = I \quad (10.2)$$

for magnetostatic cases, and

$$P_0 = \begin{bmatrix} I & 0 \\ 0 & A_V \end{bmatrix} \quad (10.3)$$

for dynamic cases, with A_V the diagonal block corresponding to the scalar potential of the A-V formulation. As we are not using any gauging for the curl-curl systems currently, the direct solvers are not available for solution of these systems.

Currently the mortar projector assumes that each node on the master boundary can be found in the slave boundary of the same partition. This limits the current parallelization strategy such that the ownership of the mortar boundaries must be carefully considered when making the partitioning. If the ownership issues is honoured, standard iterative parallel linear solvers may be used for solving the constrained system.

10.1.2 Eliminating linear constraints

Sometimes the constraint block C allows for applying an elimination procedure. The elimination may be highly beneficial since it removes the saddle-point nature of the problem and makes it much easier to solve with standard linear solvers.

10.2 Keywords related to treatment of linear constraints

The following keywords should be put in the solver section of the solver in question.

`Apply Mortar BCs` `Logical`

This flag introduces mortar BCs that results in linear constraints. More information on mortar BCs are provided in Section 9.2.

`Apply Contact BCs` `Logical`

This flag introduces contact BCs that results in linear constraints.

`Export Lagrange Multiplier` `Logical`

If linear constraints are involved, this defines whether the related Lagrange multipliers should be saved as a field. Often this field has a useful physical interpretation.

`Lagrange Multiplier Name` `String`

Name of the Lagrange multiplier to be exported.

`Use Transpose values` `Logical`

Should we use the constraint matrix in a transposed form.

`Nonlinear System Convergence Without Constraints` `Logical`

Taking into account the linear constraints (i.e. the Lagrange multipliers) in the nonlinear system norm may have a large impact on the value. Often it is preferable just to study the norm of the solution and neglect the additional dofs associated with the constraints. This is enforced by setting this keyword to `True`. The same effect is obtained also by setting `NonLinear System Consistent Norm` to `True`.

`Eliminate Linear Constraints` `Logical`

This specifies whether the Lagrange multipliers are eliminated. A sufficient condition for the elimination is that the slave part of the constraint matrix is diagonal. This may be obtained by either using a biorthogonal basis with the accurate integration or in conjunction with a discontinuous boundary.

`Eliminate Slave` `Logical`

This enforces also the elimination of slave degrees of freedom using the constraint equations. This is not possible if the constraint matrix is not diagonal (for example when using `Biorthogonal Dual LCoeff = True`), or there are diagonal entries in the constraint equation rows.

`Eliminate From Master` `Logical`

Instead of eliminating the Lagrange multiplier using the slave part use the master part. This could be applicable if we use the keywords `Discontinuous Boundary` and `Mortar BC Diag` and the constraint equation involves the diagonal. In this way both the Lagrange multiplier and the slave dofs can be eliminated.

Chapter 11

Solving constraint modes

11.1 Introduction

Elmer includes a possibility to compute constraint modes associated with a selected part of the boundary referred to as the attachment boundary. These are solution modes that are obtained by successively applying a unit field value to each degree of freedom associated with the attachment boundary. The results may be used for model lumping, for example. Together with solving eigenproblems where the degrees of freedom associated with the attachment boundary are fixed, the constraint modes may be employed to perform the Craig-Bampton component mode synthesis.

11.2 Keywords related to constraint modes analysis

`Solver solver id`

`Constraint Modes Analysis Logical`

When true, this instructs Elmer to compute constrained modes using the field equation of the solver.

`Constraint Modes Lumped Logical`

This defines whether the constraint modes are lumped such that all the dofs associated with the attachment boundary are given the fixed value, or whether all the dofs have a separate permutation. The default is that the system is not lumped and each dof has a separate permutation. Setting this flag to `True` enforces lumping.

`Constraint Modes Fluxes Logical`

Computes fluxes between the constraint modes. The fluxes are used to determine a conductivity matrix between the constraint modes. For example, if the equation in question is electrostatic equation, the fluxes will be surface charges, and the resulting lumped matrix is the capacitance matrix.

`Constraint Modes Fluxes Symmetric Logical`

This is used to assume that fluxes between constraint modes are symmetric in order to enforce the symmetry of the conductivity matrix.

`Boundary Condition bc id`

`Constrained Modes Varname String`

This keyword can be used to specify the attachment boundary and to specify that the field variable `Varname` is used for obtaining the constraint modes.

Chapter 12

Working with nodal loads

Usually source data for equations are specified as spatial fields. Elmer then internally transforms these definitions to entries of the right-hand side vector by following the weak formulation of the problem. This is convenient as the specification of data is independent of the underlying mesh.

However, it may sometimes be attractive to define discrete loads directly as entries of the right-hand side vector. Similarly, after the solution of the discrete system has been accomplished, it may be useful to consider discrete reactions which occur at nodes to maintain the balance. Both options are supported in Elmer. In addition, integrals of the basis functions can be computed to generate weights which can then be employed to reproduce a spatial data field approximately corresponding to the entries of the given right-hand side vector.

12.1 Setting nodal loads

Similarly to the Dirichlet constraints one may also set nodal loads, i.e. the entries of the right-hand side vector in the linear system. This has to be done with care as proper values generally depend on the mesh. There are anyhow some uses for this option. For example, in multiphysical couplings it may sometimes be a good choice to transfer the forces directly in the nodal form as this is the consistent way to compute the forces resulting from the solution of the discrete system.

```
Body Force bf id
```

```
Varname Load Real
```

This can be used to specify directly the entries of the right-hand side vector in the linear system. The nodal loads are given exactly as the Dirichlet conditions except that a string `Load` is attached to the name of the variable occurring in the place of `Varname`.

12.2 Computing nodal loads

It is possible to evaluate the nodal loads after the solution has been computed. This however requires that the original matrix A_0 is maintained without eliminating the rows and columns that correspond to the known Dirichlet conditions. Then the nodal forces may be computed from

$$f = A_0 x - b. \quad (12.1)$$

It should be noted that the nodal values depend on the mesh. If SI units are employed, for heat equation they are expressed in Watts and for electrostatic equation in Coulombs, for example.

```
Solver solver id
```

```
Calculate Loads Logical True
```

This keyword activates the computation of nodal loads. The resulting values will be saved to the variable the name of which is derived from the primary variable name by adding the suffix `Loads` to it.

12.3 Computing nodal weights

While the computation of the nodal loads often gives the best approximation of boundary fluxes, for example, they are however in cumbersome units as the nodal loads depend on the mesh. It may thus be more ideal to transform the nodal loads into a distributed field. To this aim there is a possibility to compute the values of the integrals

$$w_i = \int_{\Omega} \phi_i d\Omega \quad (12.2)$$

where ϕ_i is the basis function associated with the mesh node i . When the basis forms a partition of unity, the sum of all weights w_i gives the volume (or area) of the domain occupied by the elements. An approximate reproduction of the data field \bar{g} corresponding to the nodal load f may be obtained as $\bar{g} = f_i/w_i$.

Solver solver id

Calculate Weights Logical True

This keyword activates the computation of the nodal weights w_i .

12.4 Energy norm

When the initial matrix is known, the energy norm of x may be computed as

$$E = x^T A_0 x. \quad (12.3)$$

Solver solver id

Calculate Energy Norm Logical True

This activates the computation of the energy norm. The result will be written to the `Simulation` block with name `res: VarName Energy Norm` so that the `SaveScalars` subroutine may be used to save the result. The energy norm may only be computed when the loads are also computed.

Chapter 13

Generic solver utilities

When the solvers use the default procedure for solving the partial differential equations, there are a number of generic features that may be used with any model. This chapter describes these features.

13.1 Solver activation

There is a large number of different ways how solvers need to be activated and deactivated. Mostly these needs are related to different kinds of multiphysical coupling schemes. There is also a large number of auxiliary solvers that are needed, for example, only when results are saved or inspected. In the `Solver` section one may give the following keywords.

`Exec Solver String`

The value options are here *never*, *always*, *before all*, *before timestep*, *after timestep*, *before saving*, *after saving*, *after all*. If nothing else is specified, the solver is called every time in its order of appearance. The saving instance refers to the one defined by `Output Intervals` which controls the saving of the results.

`Exec Intervals Integer`

This keyword gives an interval at which the solver is active. At other intervals the solver is not used. This should have the same length as `Timestep Intervals` in the `Simulation` section.

`Start Time Real`

The start time after which the execution of the solver is started.

`Stop Time Real`

The stop time after which the execution of the solver is halted.

`Exec Condition Real`

Given this keyword, activate the solver only upon a positive value for the keyword.

13.2 Variable names

The variable name is presented in the `Solver` section by keyword `Variable`, for example

`Variable = Varname`

This name is used when setting Dirichlet conditions and initial conditions. Also the name is used as a basis for other features appending it with suffixes such as `Load`, `Condition` and `Passive`, described later in this chapter.

Sometimes one may want to rename the components of the primary variable. This may be done by defining the component names in the brackets. For example the command

```
Variable = Flow[Velo:2 Pres:1]
```

declares that variable `Flow` consists of `Velo` with two components and `Pres` with one component. If the number of components is 2 or 3, the variable will be treated as a vector in the ElmerPost files.

If one does not require output for a certain variable, one may prevent it with the `-nooutput` option, e.g.

```
Variable = -nooutput Dummy
```

One may also use the `-dofs` option to define the number of components in a variable, e.g.

```
Variable = -dofs 3 Flow
```

By default variables are assumed to be field variables. However, they may also be scalars which have globally the same value. These variables may be introduced with the `-global` option, e.g.

```
Variable = -global Frequency
```

After defining a global variable, it may be used similarly to `time` in giving dependencies.

These different options should not be fully mixed.

13.3 Exported and derived variables

The intent of exported variables is to enable automatic allocation and treatment of additional data that may usually be derived from the primary fields. Often this is done within a solver and many times the machinery is used transparently from the user.

Each solver that has a primary variable (defined by the `Variable` keyword) may also have exported variables. It is in principle way to define easily additional variables within the same solver. By default The exported variable inherits the permutation of the primary variable. The declaration would simply be

```
Exported Variable 1 = myvar
```

However, it is also possible that the variable is of different type. The currently allowed types include

- `-nodal`: default field associated with nodal finite elements
- `-global`: variable is defined globally (state variable)
- `-elem`: variable having a constant value within the element
- `-dg`: variable that is defined separately for nodes in each elements
- `-ip`: variable existing on the default integration points of the solver

For code developers it is worth noting that the permutation for the different variable types is quite different than for standard ones.

- `-nodal`: for node p the degree of freedom is `Perm(p)`
- `-global`: no permutation associated
- `-elem`: for element t the global dof is `Perm(t)`
- `-dg`: for element t and node p the global dof is `Perm(Element % DGIndexes(p))`
- `-ip`: for element t and integration point p the global dof is `Perm(t) + p`

The exported variable may be masked such that it will be defined in a domain smaller than the primary variable. Note that the domain may not be larger! This is achieved with a given mask `"secname: maskname"`

Here the `secname` can be any of `body`, `body force`, `material`, `boundary condition` and `maskname` refers to any keyword. The mask may be either logical keyword and then its value will be used. If the keyword is non-logical its presense will be used. So one can use some suitable keyword for masking purposes. For example,

The exported variable types may in most parts be used also in the `ListGet` operations. The main limitation is that the `-ip` variable can only have dependencies when the `ListGetElementReal` routine is used and the `GaussPoint` is given as the optional parameter. The `ListGetReal` routine does not know how to use the Gaussian points to return nodal values. So to fully use the features one needs to change the calling routines.

Upon request the exported variables may also be defined by the user in the `Body Force` and `Boundary Condition` sections. The operation is set behind keywords to circumvent unwanted definitions. The exported variables also follow the initial conditions.

Also upon request, the exported variables may be made transient, and optionally be numerically differentiated to obtain the first or second derivatives.

There can also be derivatives of the primary variable that are treated here since they are treated pretty similarly as the exported variables.

`Solver solver id`

`Exported Variable i -vartype Varname`

A name for an additional variable created by the solver, $i=1,2,3,\dots$. The `-vartype` may be any of `nodal`, `global`, `dg`, `elem`, `ip`.

`Exported Variable i Mask Section: Maskname`

An optional mask for defining the exported variable in a reduced domain.

`Update Exported Variables Logical`

Update the exported variable values of the current solver after leaving the iterative solution for the next solver, i.e. in the steady-state iteration level.

`Nonlinear Update Exported Variables Logical`

Update the exported variables within the nonlinear solution to get the current iterate.

`Calculate Velocity Logical`

Calculate the simple first order time differential for the primary variable. Only applicable to time-dependent simulations. Calculation is performed after each steady-state iteration.

`Nonlinear Calculate Velocity Logical`

As the previous keyword except calculation of the velocity is performed after each nonlinear iteration.

`Calculate Acceleration Logical`

Calculate the simple second order time differential for the primary variable. Only applicable to 2nd order equation in time. Calculation is performed after each steady-state iteration.

`Nonlinear Calculate Acceleration Logical`

As the previous keyword except calculation of the acceleration is performed after each nonlinear iteration.

`VarName Calculate Acceleration Logical`

Calculate the simple second order time differential of exported variable `VarName` creating new variable `VarName Acceleration`. Also enforces the field to be transient automatically.

`VarName Transient Logical`

Make the exported variable `VarName` transient such that space is allocated for field values at previous timesteps.

`VarName Calculate Velocity Logical`

Calculate the simple first order time differential of exported variable `VarName` creating new variable `VarName Velocity`. Also enforces the field to be transient automatically.

`VarName Calculate Acceleration Logical`

Calculate the simple second order time differential of exported variable `VarName` creating new variable `VarName Acceleration`. Also enforces the field to be transient automatically.

13.4 Active and passive elements

In Elmer it is possible to define certain areas of the modeled geometry to be passive during the solution. This feature also allows for deactivating and reactivating of the elements. Being a passive element means that its contribution is not included in the global matrix assembly. One could, for example, model two separate bodies heated with different heating power and connect them with a third structure only after suitable time has elapsed. This all could be modeled within a single simulation.

The geometry of the whole system is meshed as usual, and the passive elements are only omitted in the assembly of equations. The passive definition is done solverwise and elementwise. The former means that for example the temperature may be passive and the displacements active at the same element. The passive property of elements is defined in the `Body Force` section using a real valued parameter with the name constructed from the name of the variable followed by `Passive`. When the parameter obtains a value greater than zero, the element is passive.

```
Body Force  body force id
```

```
Varname Passive  Real
```

If this variable obtains a positive value, the element is set passive. Note that it is not possible to control components of vector-valued variables separately.

13.5 Timing the solvers

Often it is of interest to time the performance of the different steps in the solution sequence. For that purpose there are some keywords for activating timers for the solvers.

```
Solver  solver id
```

```
Linear System Timing  Logical True
```

This keyword activates the timing for the linear system.

```
Linear System Timing Cumulative  Logical True
```

This keyword sums up the cumulative time used in solving the linear systems.

```
Solver Timing  Logical True
```

This keyword activates the timing for the whole solver including iteration over nonlinear and linear systems etc. The time used for assembly may roughly be estimated as the difference between the linear solution time and the total time used in the solver.

```
Solver Timing Cumulative  Logical True
```

As the previous one but sums up the cumulative time used in the solver.

13.6 Consistency testing of solvers

Each solver may be tested for consistency. The way the testing is performed is that it is assumed that the correct solution will have a predefined norm. After the simulation has been carried out the computed norm and predefined norm are compared. If they are the same with a given tolerance the test is assumed to be passed.

The result of a test is outputted on the screen. Additionally a file named `TEST.PASSED` is saved in the working directory. If all the tests of the case are passed then the number 1 is written to the file in ascii format. If any of the tests fails then the number 0 is written to the file, correspondingly.

```
Reference Norm  Real
```

The target norm for the solver. The computed norm is tested against this one.

```
Reference Norm Tolerance  Real
```

The tolerance for testing. The default is 1.0e-6. A sloppier tolerance could be used if the case is known to be sensitive against numerical fluctuations.

Chapter 14

Meshing Utilities

ElmerSolver includes some internal possibilities to affect the mesh. These include the ability to split the mesh repeatedly on a partitioned level (mesh multiplication), the possibility to extrude meshes within the code, and the possibility to create discontinuous meshes on-the-fly, or to generate discontinuity using discontinuous galerkin method. These features make it possible to eliminate some meshing bottle-necks and to implement more versatile physical boundary conditions, for example.

14.1 Coordinate transformation

By default Elmer read in coordinates $\vec{r} = (x, y, z)$ as they are in the mesh files. The user may perform scaling such that $\vec{r} := C \vec{r}$. Here C may be either a scalar or a diagonal matrix that maps each coordinate differently. It could be used to scale the units of the mesh back to meters if they were initially given in mm.

The coordinates (x, y, z) may be mapped to any order by a suitable permutation which is any ordered list of numbers 1, 2 and 3. The mapping could be useful, for example, if the initial coordinate system has the rotational axis at x -axis whereas Elmer always assumes that the rotational axis is the y -axis.

Coordinate transformations can be used to map the Cartesian coordinates to some other coordinate system and vice versa. Currently supported transformations are from Cartesian to cylindrical,

$$r = r_0 + \sqrt{x^2 + y^2} \quad (14.1)$$

$$\phi = \arctan(y, x) \quad (14.2)$$

$$z = z \quad (14.3)$$

and cylindrical to Cartesian,

$$x = \cos(\phi)(r + r_0) \quad (14.4)$$

$$y = \sin(\phi)(r + r_0) \quad (14.5)$$

$$z = z. \quad (14.6)$$

The following keywords are related to coordinate transformations

Simulation

Coordinate Scaling Real

The coefficient used to scale the coordinate.

Coordinate Mapping(dim) Integer

The permutation of the coordinates.

Coordinate Transformation String

Options are "cartesian to cylindrical", "cylindrical to cartesian" and "previous".

Coordinate Transformation Radius Real

The offset r_0 in radius.

Note that the `Coordinate Transformation` may also be given in the Solver section which can then be used to invoke different coordinate systems for different solvers.

14.2 Mesh multiplication

Mesh multiplication is the process where each mesh edge is splitted into two resulting to an increased number of elements. In 2D each mesh multiplication increases the number of elements by factor of 4, and in 3D by factor of 8 correspondingly. For hexahedral elements this results to eight elements that are of original shape. For tetrahedral elements the innermost octahedron may be split into four tetrahedra in three ways of which only one preserves the minimum angle condition (Delaunay splitting) and this is naturally implemented in Elmer.

The mesh multiplication may be used recursively. This makes it possible to create meshes with a very large number of elements. The functionality of mesh multiplication may also be used in conjunction with the geometric multigrid of Elmer. Saving the meshes of different density gives an optimal hierarchy of meshes. The main handicap of mesh multiplication is that the geometric description will not be improved during the mesh multiplication. Therefore curved boundaries will be as faceted as they were in the coarsest mesh level.

Elmer includes also a feature for maintaining mesh grading. For simple geometries the mesh grading may be conserved with the following scheme:

1. Solve at the coarsest level with Galerkin method the nodal mesh parameter from the equation $h = V_e^{1/m}$, where h is the element size factor, V_e is the element volume and m is the dimension of grading.
2. When splitting an edge for the new mesh, inherit the size factor and coordinate from $h' = \sqrt{h_1 h_2}$ and $\vec{r}' = (h_1 \vec{r}_2 + h_2 \vec{r}_1) / (h_1 + h_2)$.

The procedure may be applied recursively. It provides ideal grading for structured boundary layers with $m = 1$ and homogeneous Delaunay meshes with $m = \text{dim}$.

The functionality of mesh multiplication is implemented also in parallel and is very efficient since it requires only quite little communication.

The following keywords are related to the mesh multiplication.

Simulation

Mesh Levels Integer

The number of mesh levels when using the equal split feature.

Mesh Keep Integer

The user may choose to optionally keep more than one level. This could be needed for coarser postprocessing, for example.

Mesh Keep Grading Logical

When creating meshes using equal split, the elements are by default split to equally sized pieces. When this flag is on, the solver tries to maintain the mesh grading.

Mesh Grading Power Real

The mesh grading is evaluated from the element sizes of the initial mesh. The size is a scalar function and cannot therefore handle very complicated meshes. For boundary layer type of meshes the optimal value is one, while for meshes of Delaunay type the optimal value is the space dimension.

14.3 Mesh extrusion

The feature of mesh extrusion was initially implemented for the needs of computational glaciology where the meshing had become a serious bottle-neck. The current implementation of mesh multiplication takes a 1D or 2D mesh and extrudes it to the desired direction. In the extrusion process the dimension of the elements also grow. Line segments are extruded into quadrilaterals, triangles into wedges, and quadrilaterals into hexahedrons, correspondingly.

When the mesh is extruded, the boundary condition will also be extruded. New boundary conditions may be set to the bottom and top of the mesh. These boundaries will get the first free boundary condition index and may thereby be referenced in the `sif` file.

The mesh density of the extrusion is defined only once and this same division is copied for all nodes. So each node gives rise to exactly similar 1D mesh stride in the desired direction.

Because only one division needs to be defined it can be more costly than would be otherwise feasible. Therefore also an iterative process for setting the division has been implemented. There the user can give an arbitrary mesh parameter function $h(x)$ and a 1D mesh is created so that the mesh position as accurately obeys the desired distribution. The algorithm is an iterative one based on a Gauss-Seidel type of iteration.

First initialize the distribution:

```
do i = 0, n
    w_i = i/n
end
```

Then iterate the distribution until it has converged:

```
do while( max |dw_i| > ε )
    do i = 1, n
        h_i = f((w_i + w_{i-1})/2)
    end
    do i = 1, n - 1
        w_i = (w_{i-1}h_{i+1} + w_{i+1}h_i)/(h_i + h_{i+1})
    end do
    do i = n - 1, 1, -1
        w_i = (w_{i-1}h_{i+1} + w_{i+1}h_i)/(h_i + h_{i+1})
    end do
end do
```

The following keywords are used to control the internal extrusion

Simulation

Extruded Mesh Levels Integer

The number of node layers in the extruded mesh. This keyword not only defines the number but also activates the whole feature. The number of node layers is the same over the whole extruded mesh.

Extruded Coordinate Index Integer

The index (1,2,3) of the coordinate to be extruded. The default direction is `dim+1`. So for 2D mesh the extrusion is by default carried out to the z -direction.

Extruded Mesh Ratio Real

The ratio of the first and last element when extruding the mesh. The element size is assumed to follow a geometric division. The same division is applied for all strides resulting from extruding a node.

Extruded Mesh Density Real

This keyword may be used to define an arbitrary functional dependence that defines the local mesh parameter h . The function will be automatically called at the interval $[0, 1]$. The same division is applied for the strides.

Extruded Min Coordinate Real

The mesh is by default extruded to interval $[0, 1]$. This keyword may be used to set a new minimum coordinate for the extrusion.

Extruded Max Coordinate Real

The mesh is by default extruded to interval $[0, 1]$. This keyword may be used to set a new maximum coordinate for the extrusion.

Preserve Baseline Integer

Sometimes the user wants to give a boundary condition for the initial baseline of the 2D mesh.

This keyword may be used to preserve it for setting the boundary condition. Otherwise only the extruded version of the baseline, i.e. higher dimensional creature will be available.

14.4 Discontinuous mesh for internal jumps

Discontinuous interfaces are sometimes needed for internal boundary conditions. For example, in the heat equation a poor thermal contact between objects may lead to a temperature jump at the interface. Such a jump may be difficult to model with standard elements since they cannot reproduce a discontinuity in the mesh. For this purpose a second set of nodes at the interface may be created. All the elements that share nodes at the interface must then be carefully checked to decide on which side of the interface the elements are located, and set the new element indices accordingly.

For some nodes on the boundary there is some room for interpretation when creating a discontinuous boundary. By default all nodes that have even just one hit on the boundary are doubled, i.e. a greedy method is used. The other alternative is to not include nodes which are attached also to other than boundary nodes and their immediate parents. This could be applied mainly to structured meshes with some cut through the mesh. For triangular and tetrahedral meshes the greedy method should be used since in these meshes elements will always have some orphan nodes on the boundary.

After a discontinuity has been created into the mesh the solution will by default be totally discontinuous, i.e. no flux will pass through the boundary. To allow again continuous solutions mortar and periodic projectors may be used to ensure continuity. Some field variables may be continuous and some discontinuous over the interface. The projectors with related jump and offset terms will allow this.

The following keywords are used to control the creation of discontinuous boundaries.

Simulation

`Discontinuous Boundary Greedy Logical [True]`

Whether to use the greedy method when defining the discontinuous boundary. The user may choose to set the flag to `False` freeing some corner nodes that are located in the junction of continuous and discontinuous interface from being treated as discontinuous.

`Discontinuous Bulk Greedy Logical [True]`

Similar as the one for boundary element but now the nodes are treated as continuous if they can be associated with some non-parental bulk element. This probably works correctly only for structured meshes. In unstructured meshes there are usually many nodes at the boundary associated with non-parent elements.

Boundary Condition `bc_id`

`Discontinuous Boundary Logical`

Create discontinuous boundary on-the-fly. The nodes on the boundary are doubled and heuristics is used to determine for each element which of the doubled nodes it should use. Discontinuous boundaries may only be created for internal boundaries. Currently there can be only one discontinuous boundary at a time.

`Discontinuous Target Bodies Integer`

This keyword may be used to control the owner of the discontinuous boundary (i.e. the "left" parent element). If not specified the ownership will be given to the parent element with a higher body index. There may be several discontinuous target bodies. Then the size of the array must also be specified.

`Discontinuous BC Integer`

This keyword may be used to activate the creation of the other side of the discontinuous boundary. The other side will be needed if the two boundaries, for example, move with respect to each other. The value of this keyword relates to the boundary condition index of the new boundary. Note that also `Mortar BC` and `Periodic BC` keywords will have the same effect but only when they are nonzero! If mortar or periodic boundary condition is requested with a zero argument, the discontinuous boundary will not be doubled. Instead a cheaper way to create the mapping between the boundary nodes is used.

14.5 Discontinuous Galerkin for internal jumps

The discontinuous interfaces described above have some challenges. The most eminent is that fact that as the mesh is made discontinuous to enable physical jumps in one condition, the discontinuous mesh needs to be considered also for all other equations. This is cumbersome particularly as the additional constraints make the resulting linear system more difficult to solve. Also, the parallel operation of the discontinuous mesh is not very robust currently. Therefore, using Discontinuous Galerkin method offers a more general method to enable discontinuities in the mesh.

Assume that we only have discontinuities at the interfaces. Then it does not really make sense to have all possible discontinuities enabled as in standard `Discontinuous Galerkin` method. Instead, we create a reduced basis where discontinuities are eliminated using the permutation vector. The end effect is much smaller system.

The Discontinuous Galerkin framework has the additional advantage that the other equations living on the same mesh are unaffected. This is ideal for multiphysical problems where jumps only related to some phenomena.

For the reduction we have different techniques that define how many discontinuous groups there are in the mesh, and among which of them discontinuities are enabled.

The following keywords are used to control the generation and preservation of discontinuities between bodies.

`Solver solver id`

`Discontinuous Galerkin Logical [True]`
Enable the Discontinuous Galerkin solver.

`DG Reduced Basis Logical`
Should we used reduction of the dofs for minimal basis that still conserves the wanted discontinuities.

`DG Reduced Basis Mapping Integer`
This defines the mapping such that the initial bodies may be mapped to smaller number of groups and thereby eliminate the discontinuities among bodies belonging to the same group.

`DG Reduced Basis Master Bodies Integer`

`DG Reduced Basis Slave Bodies Integer`
Master and slave bodies create two body groups and the discontinuities are preserved only between them. If there are bodies which do not belong to either group then they are not affected by the discontinuity.

`DG Indirect Connections Logical`
This keyword may be used to optionally generate also indirect connections over the discontinuous interface. This makes in possible to use various elimination techniques for the jump conditions.

Chapter 15

Adaptive Solution

15.1 Introduction

A posteriori error analysis and adaptive mesh refinement are nowadays standard tools in finite element analysis when cracks, boundary layers, corner singularities, shock waves, and other irregularities are present. A posteriori error indicators can be used to reveal flaws in finite element discretizations and well-designed adaptive mesh refinements can reduce the computational costs drastically.

15.2 Theory

Let us consider equilibrium equations of the form

$$-\nabla \cdot q = f \text{ in } \Omega, \quad (15.1)$$

$$q \cdot n = g \text{ on } \Gamma, \quad (15.2)$$

where q is either a flux vector or a second order stress tensor, Ω is a computational domain, Γ is a boundary part, f is an external source or body force, g is an external flux or traction and n is the unit outward normal to the boundary.

Most symmetric steady state problems described in the model manual of Elmer [] fit in the above framework of equilibrium equations. To fix ideas, suppose that q is the heat flux satisfying Fourier's law $q = -k\nabla T$, where T is the temperature and k is the heat conductivity of the material. We could also think of q as the stress tensor of linear elasticity. In this case Hooke's law states that $q = \mathcal{E} : \varepsilon$, where \mathcal{E} is the fourth order tensor of elastic coefficients, $\varepsilon = \text{sym}(\nabla u)$ is the linearized strain tensor and u is the displacement vector.

15.2.1 A posteriori estimate

Let us denote the finite element approximation of q by q_h and measure the error $q - q_h$ as

$$ERROR = \sqrt{\int_{\Omega} |q - q_h|^2 d\Omega} \quad (15.3)$$

Our primary goal is to ensure the accuracy of the solution by imposing the condition

$$ERROR \leq TOLERANCE \quad (15.4)$$

where $TOLERANCE > 0$ is an error tolerance prescribed by the user.

In practise, the goal must be replaced by a stronger condition

$$ESTIMATE \leq TOLERANCE \quad (15.5)$$

where $ESTIMATE$ is a computable functional (of all available data) satisfying

$$ERROR \leq ESTIMATE \quad (15.6)$$

Then, if (15.5) holds, (15.4) is satisfied and the quality of the numerical solution is guaranteed.

In Elmer the a posteriori estimate (15.5) is computed from local residuals of the finite element solution as a weighted sum over the elements,

$$ESTIMATE = \sqrt{\sum_E \eta_E^2}, \quad (15.7)$$

where η_E is the local error indicator for an individual element E :

$$\begin{aligned} \eta_E^2 = & \alpha_E h_E^2 \int_E |\nabla \cdot q_h + f|^2 d\Omega \\ & + \beta_E \sum_{e \text{ in } \Omega} h_e \int_e |[q_h \cdot n_e]|^2 d\Gamma \\ & + \gamma_E \sum_{e \text{ on } \Gamma} h_e \int_e |q_h \cdot n_e - g|^2 d\Gamma \end{aligned} \quad (15.8)$$

Here α_E , β_E , and γ_E , are local positive constants. The values of these constants depend, among other things, on the problem to be solved, and must be estimated carefully case by case [1].

The first sum in (15.8) is taken over all edges e of E inside the computational domain, the second sum is taken over all edges on the boundary part Γ , $[\![\cdot]\!]$ is the jump in (\cdot) across e , and n_e is a unit normal to the edge. h_E is the size of the element and h_e is the size of the edge.

The first term on the right-hand-side of (15.8) measures the local residual of the finite element solution with respect to the equilibrium equation (15.1). The second term measures the discontinuity in the numerical flux inside Ω and the third term the residual with respect to the boundary condition (15.2).

15.2.2 Adaptivity

The secondary goal of our numerical computations is to find a solution satisfying (15.4) as efficiently as possible. A nearly optimal solution strategy is obtained by utilizing the property (here we need to impose some minor restrictions on f and g , see [1])

$$LOCAL ERROR \geq \eta_E \quad (15.9)$$

where

$$LOCAL ERROR = \sqrt{\int_E |q - q_h|^2 d\Omega} \quad (15.10)$$

The estimate suggests that the error in the numerical solution should be reduced efficiently if the mesh is refined locally where the indicators η_E are large. Naturally, we can think of coarsening the mesh where the values of the indicators are small.

The adaptive mesh refinement strategy of Elmer is based on the local estimate (15.9) and on the following additional assumptions and heuristic optimality conditions:

- The local error behaves as

$$\eta_E = C_E h_E^{p_E} \quad (15.11)$$

for some constants C_E and p_E .

- In the optimal mesh the error is uniformly distributed over the elements:

$$\eta_E = TOLERANCE / N_{elements} \quad (15.12)$$

The constants C_E and p_E in (15.11) can be solved locally for each element if the local errors and the local mesh sizes are known from at least two different solutions. The second rule (15.12) can then be applied to extrapolate a new nearly optimal mesh density for the subsequent calculations.

The mesh refinements can be performed either by splitting the existing elements into smaller using the so called RGB-refinement strategy described in [], or by performing a complete remeshing of the computational domain using the built-in unstructured mesh generators that produce high quality Delaunay triangulations. In the latter alternative not only mesh refinement is possible, but also local adaptive coarsening.

15.3 Keywords related to the adaptive solution

The adaptive solver of Elmer is activated and controlled by the following keywords in the Solver block of the solver-input-file.

Adaptive Mesh Refinement Logical

If set to true, then after the solution of the linear system the program computes residual error indicators for all active elements, estimates the global error, computes a new mesh density and refines the mesh accordingly.

Adaptive Remesh Logical

If set to true, then a complete remeshing is performed after error estimation using the Mesh2D or Mesh3D generators. The new mesh density is written in file “bgmesh”. If set to false, then the RGB-splitting strategy for triangles is applied to perform the refinements.

Adaptive Save Mesh Logical

If set to true, the subsequent meshes are stored in directories RefinedMeshN, where N is the number of the adaptive iterate.

Adaptive Error Limit Real

Error tolerance for the adaptive solution.

Adaptive Min H Real

Imposes a restriction on the mesh size. Default is zero.

Adaptive Max H Real

Imposes a restriction on the mesh size. Default is infinite.

Adaptive Max Change Real

Controls the change in local mesh density between two subsequent adaptive iterates. Using this keyword the user can restrict the refinement/coarsening to stabilize the adaptive solution process.

15.4 Implementing own error estimators

Suppose that we are given a subroutine called `MySolver` for solving the Poisson equation, and we would like to enhance the code by implementing an a posteriori error indicator for adaptive mesh refinement. The first thing to do is to take the module `Adaptive` in use, and define the local error indicators as functions in an interface block. The beginning of the subroutine should look like the following:

```
SUBROUTINE MySolver( Model, Solver, Timestep, TransientSimulation )
  USE DefUtils
  USE Adaptive

  INTERFACE
    FUNCTION InsideResidual( Model, Element, Mesh, &
      Quant, Perm, Fnorm ) RESULT( Indicator )
      USE Types
```

```

TYPE(Element_t), POINTER :: Element
TYPE(Model_t) :: Model
TYPE(Mesh_t), POINTER :: Mesh
REAL(KIND=dp) :: Quant(:), Indicator, Fnorm
INTEGER :: Perm(:)
END FUNCTION InsideResidual

FUNCTION EdgeResidual( Model, Edge, Mesh, &
    Quant, Perm ) RESULT( Indicator )
USE Types
TYPE(Element_t), POINTER :: Edge
TYPE(Model_t) :: Model
TYPE(Mesh_t), POINTER :: Mesh
REAL(KIND=dp) :: Quant(:), Indicator
INTEGER :: Perm(:)
END FUNCTION EdgeResidual

FUNCTION BoundaryResidual( Model, Edge, Mesh, &
    Quant, Perm, Gnorm ) RESULT( Indicator )
USE Types
TYPE(Element_t), POINTER :: Edge
TYPE(Model_t) :: Model
TYPE(Mesh_t), POINTER :: Mesh
REAL(KIND=dp) :: Quant(:), Indicator, Gnorm
INTEGER :: Perm(:)
END FUNCTION BoundaryResidual
END INTERFACE

```

After these fixed declarations we may proceed normally by defining the local variables, allocate memory for local tables, integrate the stiffness matrix, set boundary conditions, and solve the problem. Error estimation and adaptive mesh refinements are then performed by calling the subroutine `RefineMesh`, which should appear in the code right after the function `DefaultSolve`.

```

Norm = DefaultSolve()

IF ( ListGetLogical( Solver % Values, 'Adaptive Mesh Refinement' ) ) &
    CALL RefineMesh( Model, Solver, Potential, Permutation, &
        InsideResidual, EdgeResidual, BoundaryResidual )

```

The functions `InsideResidual`, `EdgeResidual` and `BoundaryResidual` defined in the interface block should finally be contained in `MySolve`, and return the values of the error indicators described in the previous section.

As an example, suppose that we are using linear triangles or tetrahedra for solving the Poisson equation. In this case it holds $\nabla \cdot q_h = 0$ on each element E , and the contribution of the first term in (7.1) is simply

$$\text{InsideResidual} = h_E \sqrt{\int_E |f|^2 d\Omega} \quad (15.13)$$

The function that computes the value of the inside residual could be written as follows.

```

FUNCTION InsideResidual( Model, Element, Mesh, &
    Quant, Perm, Fnorm ) RESULT( Indicator )
IMPLICIT NONE
TYPE(Model_t) :: Model

```

```

INTEGER :: Perm(:)
REAL(KIND=dp) :: Quant(:), Indicator, Fnorm
TYPE( Mesh_t ), POINTER :: Mesh
TYPE( Element_t ), POINTER :: Element

TYPE(GaussIntegrationPoints_t), TARGET :: IP
TYPE(ValueList_t), POINTER :: BodyForce
REAL(KIND=dp) :: f, hK, detJ, Basis(MAX_NODES), &
    dBasisdx(MAX_NODES,3), ddBasisddx(MAX_NODES,3,3), &
    Source(MAX_NODES)
LOGICAL :: stat
INTEGER :: n

Indicator = 0.0d0
Fnorm = 0.0d0
hK = element % hK

BodyForce => GetBodyForce( Element )
Source = GetReal( Element, 'Source' )

IP = GaussPoints( Element )
DO n = 1, IP % n
    stat = ElementInfo( Element, Nodes, IP % u(n), IP % v(n), &
        IP % w(n), detJ, Basis, dBasisdx, ddBasisddx, .FALSE. )
    f = SUM( Source * Basis )
    Fnorm = Fnorm + f**2 * detJ % IP % s(n)
    Indicator = Indicator + f**2 * detJ * IP % s(n)
END DO

Fnorm = SQRT( Fnorm )
Indicator = hK * SQRT( Indicator )

END FUNCTION Inside Residual

```

For the boundary and edge residuals refer to the example `Poisson.f90` in the tutorial manual of Elmer.

Chapter 16

Parallel runs

16.1 Introduction

In times of even simple desktop PCs containing multiple CPUs or at least multiple cores, parallel computing is a necessity to exploit the complete potential of those architectures. As on multi-core architectures multi-threading (e.g., OpenMP) would be a feasible concept, Elmer utilizes the well established Message Passing Interface standard for inter-process communication. This approach makes it possible to run Elmer on both, multi-core as well as multi processor environments.

16.1.1 Parallel Implementation in Elmer

The general concept of a parallel run within Elmer is displayed in Fig. 16.1. Elmer uses domain decompo-




Figure 16.1: The principle steps to be taken for a parallel run of Elmer

sition for distributing the load to multiple processes that are being run on either different cores or CPUs. To that end, the initial mesh has to be split into parts that – with respect to the applied models – lead to similar loads of the processors¹. This will be discussed in section 16.2.

The solver stage mostly will demand from serial runs differing numerical techniques, as solution strategies have to take care of the by the domain boundaries limited possibilities of memory access. In general, convergence of linear systems are more difficult to achieve compared to serial runs. These issues will be addressed in section 16.3.1.

Finally, as also the output of the parallel runs is split into domains, the post-processing usually demands an additional step of unifying the split results. Alternatively, new visualization software is capable to do that either on the fly or to also visualize the results using multiple processes. Especially the latter method inevitably will gain importance with the increasing size of models that cannot be handled on a single CPU/core due to memory and computational constraints. Concepts of post-processing parallel results are discussed in section 16.4.

16.2 Preprocessing of Parallel Runs

In order to utilize the decomposition, the mesh has to be split into the same amount of partitions, N , as there are different processes within the parallel computation. The plain and easy way is to start from a mesh for a serial run. The typical structure of a single domain mesh database is the following:

```
meshdirectoryname |
                  |-mesh.header
                  |-mesh.nodes
                  |-mesh.elements
                  |-mesh.boundary
```

The mesh consists of a header file, containing the basic information (e.g., numbers of nodes and elements), a file containing all nodes and two further files defining the bulk- and boundary-elements.

The parallel mesh consisting of 2 partitions is written under the same directory within the sub-directory `partitioning.2`:

```
meshdirectoryname |
                  |-mesh.header
                  |-mesh.nodes
                  |-mesh.elements
                  |-mesh.boundary
                  |-partitioning.2 |
                                |-part.1.header
                                |-part.1.nodes
                                |-part.1.elements
                                |-part.1.boundary
                                |-part.1.shared
                                |-part.2.header
                                |-part.2.nodes
                                |-part.2.elements
                                |-part.2.boundary
                                |-part.2.shared
```

These files basically reflect the structure of a single domain mesh on the partition level. Additionally, a file names `part.N.shared` (with N being the partition number) is introduced. It contains – as the name suggests – information on between domains shared nodes.

¹Currently Elmer is not able to perform internal load balancing.

16.2.1 Partitioning with ElmerGrid

Provided, a single domain mesh exists, the corresponding ElmerGrid command to create a with respect to the x-direction split mesh (in our case $2 \times 1 \times 1 = 2$ partitions) would read as

```
ElmerGrid 2 2 meshdirectoryname -partition 2 1 1 0
```

There are different methods of partitioning built into ElmerGrid. they are summarized in table 16.1

option	purpose	parameters
-partition N_x N_y N_z F	partition with respect to directions	$N_{x/y/z}$... number of partitions in x/y/z-direction, $F = 0$... element-wise partitioning, 1 ... node-wise partitioning
-partorder n_x n_y n_z	(optional in addition to previous) direction of ordering	$n_{x/y/z}$... components of normal vector of ordering
-metis N M	uses metis library for partitioning	N ... number of partitions, M ... method
		$M=0$... PartMeshNodal
		$M=1$... PartMeshDual
		$M=2$... PartGraphRecursive
		$M=3$... PartGraphKway
		$M=4$... PartGraphVKway

Table 16.1: Partition methods for ElmerGrid

Depending on what partitioning method was used, additional parameters may be used for adaption of mesh specifications Those parameters and their purpose are listed in 16.2

option	purpose	parameters
-halo	create halo for the partitioning	
-indirect	create indirect connections in the partitioning	
-periodic F_x F_y F_z	declare the periodic coordinate directions for parallel meshes and sets periodic points into same partitions	$F_{x,y,z} = 1 \dots$ periodic, 0 ... not periodic
-partoptim	apply aggressive optimization to node sharing	
-partbw	minimize the bandwidth of partition-partition couplings	
-parthypre	hypr type numbering (number the nodes continuously partition-wise)	

Table 16.2: Additional mesh generation options for ElmerGrid

Figure 16.2 shows the different distribution of partitions obtained with two different methods. In gen-




Figure 16.2: Distribution of four partitions using the options `-partition 2 2 1` (left) and `-metis 4 1` (right). It comes clear that the partitioning to the left contains more partition-partition boundaries and consequently will perform worse in a parallel run

eral, the user should utilize the `metis` options, if more complex geometries (like in Fig. 16.2) are to be decomposed. This ensures that the number of shared nodes and consequently also the amount of inter-process communication during the parallel computation is minimized. More simple objects, especially those using regular meshes, can be split according to the axis using the `partition` option without having to compromise on communication speed.

Halo Elements

One of the additional options displayed in Tab. 16.2 are so called halo elements. As displayed in Fig. 16.3, halo-elements are additional elements that do not belong to the partition (i.e., they are not contributing in the domain-wise solution procedure), but rather are replicas of the neighbor elements of adjoining partitions. Thus, in a parallel run, the values of variables as well as the geometry of the whole element are at disposition withing the domain partition. These may be needed by a specific FE method, such as the Discontinuous




Figure 16.3: The concept of halo-elements. Each partition contains information on the neighbor elements along the domain boundary (red) of the adjoining partitions thus leading to a redundant stripe of elements (light-gray) that is shared between the domains

Galerkin method or by specific solvers/functions that need additional geometric information from the other domain (e.g., element-averaged surface normals).

16.3 Parallel Computations in Elmer

As mentioned before, Elmer utilizes Message Passing Interface (MPI) for inter-process communication while doing a parallel computation. To that end, a special parallel executable that is linked to a MPI library (the minimum requirement). The compilation process for the MPI version is shortly explained in chapter 17 of this guide. The executable file of the parallel version of Elmer has a to the serial call different name, `ElmerSolver_mpi`. Typically it is executed as an argument to an additional call that is specific to the parallel (MPI) environment of the platform. For instance, in a typical MPI installation (OpenMPI, MPICH) the command

```
mpirun -np 4 ElmerSolver_mpi
```

will run a four-process parallel Elmer run. The typical screen-output upon launching `ElmerSolver_mpi` indicating the number of processes is

```
ELMER SOLVER (v 5.5.0) STARTED AT: 2009/09/09 10:28:28
ELMER SOLVER (v 5.5.0) STARTED AT: 2009/09/09 10:28:28
ELMER SOLVER (v 5.5.0) STARTED AT: 2009/09/09 10:28:28
ELMER SOLVER (v 5.5.0) STARTED AT: 2009/09/09 10:28:28
ParCommInit: Initialize #PEs: 4
```

```

MAIN:
MAIN: =====
MAIN:  E L M E R  S O L V E R  S T A R T I N G
MAIN:  Library version: 5.5.0 (Rev: 4195)
MAIN:  Running in parallel using 4 tasks.
MAIN:  HYPRE library linked in.
MAIN:  MUMPS library linked in.
MAIN: =====
MAIN:
MAIN:

```

It is – unlike in the serial version of Elmer – not possible to explicitly add the Solver Input File (SIF, suffix `*.sif`) as an argument to the calling command. The user rather has to provide a file called `ELMERSOLVER_STARTINFO` containing the file name. Additionally, the in the SIF declared mesh-directory has to contain a mesh with – in this specific case – four partitions.

16.3.1 Numerical Strategies in Parallel

The concept of domain decomposition means that ElmerSolver is run on $N > 1$ separate parts of a domain that are interlinked at the boundaries. If no special solvers (see later in this section) are utilized, this inherently means that iterative methods have to be used in order to achieve convergence for the linear(ized) system solution procedure. The selection of available iterative methods, which all fall within Krylov subspace methods, is to be found in section 4.3. These methods in general have similar convergence compared to a single process run. The biggest difference introduced by domain decomposition is, that preconditioning strategies are altered. To give an example: As only applied to the local matrix, the LU factorization of a parallel run in comparison to a serial drops the gray zones indicated in Fig. 16.4. This not necessarily will,




Figure 16.4: Difference of ILU factorization between serial and domain decomposition runs. If the factorization is applied only locally within the domain, contributions from the light-gray zones are not accounted for in the latter

but can negatively affect the convergence of the iterative method.

Hypre

Hypre is a library for solving large, sparse linear systems of equations on massively parallel computers. Hypre was developed at the Center for Applied Scientific Computing (CASC) at Lawrence Livermore National Laboratory. Hypre is distributed under the GNU Lesser General Public License and thus not included in the Elmer distribution. It rather has to be downloaded, compiled and linked together with the Elmer sources.

The principle keyword for utilizing Hypre is given in the solver section

```
Linear System Use Hypre Logical
    if set to True, Hypre will be used to solve the linear system.
```

In Elmer, the only Krylov sub-space method being implemented is the Conjugate Gradient Stabilized (BiCGStab) method, which is taken into use by

```
Linear System Solver = "Iterative"
Linear System Iterative Method = "BiCGStab"
```

In combination with the BiCGStab method, the following preconditioner can be taken into use. If ILU method for preconditioning, the following settings have to be set in the solver section (here with ILU fill-in level 1):

```
Linear System Preconditioning String "ILUN"
    with  $N$  being the fill-in level (just in the built-in Elmer preconditioner). The only significant difference to Elmer's built-in ILU preconditioner is, that in case of Hypre, the missing parts (illustrated in Fig. 16.4) are now being passed from one domain to the other. In other words, the preconditioner should behave exactly as if it would be applied in a serial, single domain run. This can improve convergence, but comes at the cost of increased inter-processor communication.
```

```
Linear System Preconditioning String "ParaSails"
    ParaSails is a sparse approximate inverse preconditioner it is preconditioner for sparse matrix systems. It has the following additional parameters
```

ParaSails Threshold is a Real value that determines the typical value for matrix entries being dropped. Its suggested range for direct input (positive sign) is from 0.0 to 0.1, with lower values demanding higher accuracy and consequently computing time/memory. Alternatively, if negative values are entered, they are interpreted as the fraction of nonzero elements that are being dropped (e.g., -0.9 leads to 90/

ParaSails Filter is a Real value that determines the typical value for matrix entries in the computed approximate inverse that are dropped. Its suggested range for direct input (positive sign) is from 0.0 to 0.05. Alternatively, if negative values are entered, they are interpreted as the fraction of nonzero elements that are being dropped (see earlier item).

ParaSails Maxlevel is an Integer value that determines the accuracy of the preconditioner. Usually a value of 0 or 1 is within the applicable frame

ParaSails Symmetry is an Integer value that determines the nature of the original matrix. the following settings are to be found from the Hypre manual:

- 0 non-symmetric and/or indefinite problem, non-symmetric preconditioner
- 1 Semi Positive Definite (SPD) problem, SPD (factored) preconditioner
- 2 non-symmetric definite problem, SPD (factored) preconditioner

A typical section for the Navier-Stokes solver being solved with BiCGStab and ParaSails could look as follows

```
Solver 1
    Equation = "Navier-Stokes"
    Optimize Bandwidth = Logical True
    Linear System Solver = "Iterative"
```

```

Linear System Iterative Method = "BiCGStab"
Linear System Max Iterations = 500
Linear System Convergence Tolerance = 1.0E-06
Linear System Abort Not Converged = True
Linear System Preconditioning = "ILU1"
Linear System Residual Output = 1
Linear System Use Hypre = Logical True
Linear System Preconditioning = "ParaSails"
ParaSails Threshold = Real -0.95
ParaSails Filter = Real -0.95
ParaSails Maxlevel = Integer 1
ParaSails Symmetry = Integer 0
Stabilization Method = Stabilized
Nonlinear System Convergence Tolerance = 1.0E-04
Nonlinear System Max Iterations = 30
Nonlinear System Newton After Iterations = 1
Nonlinear System Newton After Tolerance = 1.0E-04
Steady State Convergence Tolerance = 1.0E-03
End

```

Linear System Preconditioning String "BoomerAMG"

The multi-level procedure given by BoomerAMG can also be utilized as a preconditioner. See the following part for the particular parameters that can be passed to BoomerAMG

Additionally, the Algebraic MultiGrid (AMG) solver, BoomerAMG, can be used directly to solve the system by

```

Linear System Solver = "Iterative"
Linear System Iterative Method = "BoomerAMG"

```

The following parameters for BoomerAMG can be set

BoomerAMG Relax Type Integer

Defines the smoother on the fine grid and the up and the down cycle. Possible choices

- 0 Jacobi
- 1 Gauss-Seidel, sequential
- 2 Gauss-Seidel, interior points in parallel, boundary sequential
- 3 hybrid Gauss-Seidel or SOR, forward solve (default value)
- 4 hybrid Gauss-Seidel or SOR, backward solve
- 5 hybrid chaotic Gauss-Seidel (does not work with MPI!)
- 6 hybrid symmetric Gauss-Seidel or SSOR
- 9 Gaussian elimination (only on coarsest level)

The solver on the coarsest level is set to Gaussian elimination

BoomerAMG Coarsen Type Integer

Sets the parallel coarsening algorithm to be used. Possible options are

- 0 CLJP-coarsening (a parallel coarsening algorithm using independent sets (default)
- 1 classical Ruge-Stüben coarsening on each processor, no boundary treatme
- 3 classical Ruge-Stüben coarsening on each processor, followed by a third pass, which adds coarse points on the boundaries
- 6 Falgout coarsening (uses 1 first, followed by CLJP using the interior coarse points generated by 1 as its first independent set)
- 7 CLJP-coarsening (using a fixed random vector, for debugging purposes only)
- 8 PMIS-coarsening (a parallel coarsening algorithm using independent sets, generating lower complexities than CLJP, might also lead to slower convergence)
- 9 PMIS-coarsening (using a fixed random vector, for debugging purposes only)
- 10 HMIS-coarsening (uses one pass Ruge-Stüben on each processor independently, followed by PMIS using the interior C-points generated as its first independent set)
- 11 one-pass Ruge-Stüben coarsening on each processor, no boundary treatment

BoomerAMG Num Sweeps *Integer*

sets the number of sweeps on the finest level (default value = 1)

Boomeramg Max Levels *Integer*

sets maximum number of MG levels (default value = 25)

BoomerAMG Interpolation Type *Integer*

Sets parallel interpolation operator. Possible options are

- 0 classical modified interpolation (default)
- 1 LS interpolation (for use with GSMG)
- 2 classical modified interpolation for hyperbolic PDEs
- 3 direct interpolation (with separation of weights)
- 4 multipass interpolation
- 5 multipass interpolation (with separation of weights)
- 6 extended classical modified interpolation
- 7 extended (if no common C neighbor) classical modified interpolation
- 8 standard interpolation
- 9 standard interpolation (with separation of weights)
- 10 classical block interpolation (for use with nodal systems version only)
- 11 classical block interpolation (for use with nodal systems version only) with diagonalized diagonal blocks
- 12 FF interpolation
- 13 FF1 interpolation

BoomerAMG Smooth Type *Integer*

For the use of more complex smoothers. possible options are

- 6 Schwarz smoothers (default and recommended)
- 7 Pilut
- 8 ParaSails
- 9 Euclid

BoomerAMG Cycle Type *Integer*

For a V-cycle (default) give the value 1 for a W-cycle 2

```
BoomerAMG Num Functions Integer
      has to be equal to the value given in Variable DOFs
```

Usually, the default values deliver a good performance and should hence be used as a reference constellation. Mind also, that BoomerAMG would have more (partly obsolete) option that have not directly been made accessible through its Elmer interface.

MUMPS

The only implementation of a direct solver in the parallel version of Elmer is MUMPS (<http://mumps.enseeiht.fr/>), a sparse matrix multi-frontal parallel direct solver. MUMPS is not licensed under the GNU license terms and hence is not able to be distributed together with the Elmer sources or binaries, but has to be downloaded by the user. How MUMPS is integrated into Elmer is explained in chapter 17.

There are certain limits what comes to memory consumption, as for the time being the analysis phase is done on a master process. As a rule of thumb, about 100k Elements distributed on a quad core machine with 2 Gb of memory per process still works, but does not work fast.

MUMPS is invoked by choosing `Linear System Solver = Direct` in combination with `Linear System Direct Method = Mumps`. A typical call of the Navier-Stokes solver reads like:

```
Solver 1
  Equation = "Navier-Stokes"
  Optimize Bandwidth = Logical True
  Linear System Solver = Direct
  Linear System Direct Method = Mumps
  Linear System Convergence Tolerance = 1.0E-06
  Steady State Convergence Tolerance = 1.0E-03
  Stabilization Method = Stabilized
  Nonlinear System Convergence Tolerance = 1.0E-03
  Nonlinear System Max Iterations = 1
  Nonlinear System Newton After Iterations = 1
  Nonlinear System Newton After Tolerance = 1.0E-03
End
```

Mind, that MUMPS will not work in serial runs. There the `UMFPack` library should be applied in order to utilize the same method.

16.4 Post-processing of Parallel Runs

During a parallel run of Elmer the results are also written in parallel. That means, a run with $N > 1$ partitions/processes produces $N > 1$ output files. If the base name of the output file is `parallelrun`, given in the SIF as

```
Post File = "parallelrun.ep"
```

the results of a $N = 4$ parallel run will be written into the mesh-directory as

```
parallelrun.ep.0
parallelrun.ep.1
parallelrun.ep.2
parallelrun.ep.3
```

These files contain the results of each domain (starting with zero) for ElmerPost. Similar, if the result file base name is given as

```
Post File = "parallelrun.result"
```

the results of a $N = 4$ parallel run will be written into the mesh-directory as

```
parallelrun.result.0  
parallelrun.result.1  
parallelrun.result.2  
parallelrun.result.3
```

From these files, a new Elmer run (on the same partitioned mesh) can be restarted.

16.4.1 Combination of Results with ElmerGrid

The traditional way of post-processing Elmer results is to view them using ElmerPost. It is possible to directly load one of the domain ElmerPost output files into ElmerPost, viewing only the part of the mesh corresponding to the part of the domain decomposition. Usually, the user wants to combine those splitted results into a single. With the results as displayed above, the correct ElmerGrid command to achieve that results is

```
ElmerGrid 15 3 parallelrun
```

This will go through all timesteps (transient) or steady state levels (steady state) of all partitions and add them to the file `parallelrun.ep` (default output name can be changed using the `-out` option). If lesser timesteps/steady-state levels have been needed for convergence of the run, ElmerGrid still would try to add the maximum given number, thus filling in zeros into the combined file. This can be avoided by giving the additional option `-saveinterval first last interval`. For instance, if the user knows that the run has converged at steady-state level number 8, and is just interested to combine the converged result the command

```
ElmerGrid 15 3 parallelrun -saveinterval 8 8 1
```

executed in the mesh-directory would deliver the result. ElmerGrid assumes that all partitions in the directory have to be unified into the single output file. If for instance a run of 4 partitions has been run with the same output name after a run with 6 partitions, the files `parallelrun.ep.4` and `parallelrun.ep.5` would still reside in the mesh directory, but should not be added. This can be achieved by the additional option `-partjoin 4`. This is also useful if only some parts of the full result should be included.

16.4.2 User Defined Functions and Subroutines in Parallel

In principle, the user does not have to distinguish between a user defined function/solver in parallel and serial runs. Communication between processes is being taken care of within the built-in Elmer routines, such that the single solver does not really "see" anything from the parallel environment it is running on. In special cases a few lines of code might be necessary to deal with differences imposed by parallel runs. Those are explained in section 18.4.3.

Chapter 17

Compilation and Linking

The main Elmer distribution channel is currently GitHub. The source code may hence be found at

```
https://github.com/ElmerCSC/elmerfem
```

and a local copy of the repository can be made by giving the command

```
git clone git://www.github.com/ElmerCSC/elmerfem
```

or

```
git clone https://www.github.com/ElmerCSC/elmerfem
```

with the files placed under the working directory.

There are different branches in the repository. The most important branch `devel` contains the latest updates of Elmer and should preferably be used for compilation. In order to see what is the current branch use the command

```
git status
```

If the branch is not currently `devel`, one may switch to it by using the command

```
git checkout devel
```

The following simple instructions are intended mainly for Unix/Linux users. Additional compilation instructions may be found via the `www`-site of Elmer and Wiki pages. For example compiling a parallel version is not covered in the following.

17.1 Compiling with CMake

To compile, one should first create a separate directory for building. Now, if the working directory is the place where `git clone` command was given, one may create the building directory and proceed to it as

```
mkdir build
cd build
```

CMake (version 2.8.9 or newer is needed) may then be called to generate makefiles. The following command may already produce what is needed:

```
cmake -DWITH_ELMERGUI:BOOL=FALSE -DWITH_MPI:BOOL=FALSE
      -DCMAKE_INSTALL_PREFIX=../install ../elmerfem
```

The command line may include additional options. For example the place of BLAS and LAPACK libraries can be defined by using options `-DBLAS_LIBRARIES` and `-DLAPACK_LIBRARIES`.

If no errors are encountered in connection with running the `cmake` command, one may proceed to compilation:

```
make install
```

Finally, one may test the fresh version of Elmer as

```
ctest
```

or just perform a specific test (cf. the contents of `../elmerfem/fem/tests` for existing tests) for example as

```
ctest -R elasticity
```

In order to be able to run Elmer from anywhere, one may define the following environment variables:

```
export ELMER_HOME=$HOME/elmer/install
export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$ELMER_HOME/lib
export PATH=$PATH:$ELMER_HOME/bin
```

17.2 Compiling a user-defined subroutine

The `elmerf90` command is provided to help compiling your own solvers. It is a wrapper script to the compiler that was used to produce the Elmer installation and may be used as

```
elmerf90 -o MySolver MySolver.f90
```

In the MinGW system in Windows the suffix `.dll` should preferably be used

```
elmerf90 -o MySolver.dll MySolver.f90
```

After successful compilation, the file `mysolver.dll` should be found in the local directory. In the file-name declaration of the `Procedure`-keyword given in the solver input file, the suffix `.dll` can be omitted, so one should be able to define

```
Solver 1
  Procedure = "MySolver" "SubroutineName"
  ...
End
```

Chapter 18

Basic Programming

18.1 Introduction

The Elmer distribution contains a large set of different solvers and also the possibility to declare dependence of material properties or boundary conditions on certain variables (e.g., using the MATC language). Nevertheless, there always may occur certain physical problems that are too difficult to be handled via the solver input file. Such problems can be coped by introducing new user functions or even complete new solvers. Elmer is very flexible if certain models have to be added by the user providing her/his own code.

This chapter shall provide a brief introduction to basic programming of user functions as well as solvers that can be added to Elmer. This will be done by mainly using certain examples and explaining the programming steps occurring in these.

The Elmer Solver source is written in the programming language Fortran 90. Since the Elmer Solver binaries are compiled as shared objects, it is sufficient to just newly compile the code contributed by the user as an executable of a shared object (`.so` in UNIX and `.dll` in Windows) that dynamically is linked to the rest of the distribution. In order to provide Elmer with the needed information to be able to load an external function or solver, the following entry in the solver input file (suffix `.sif`) has to be given:

```
Procedure "filename" "procedure"
```

Where the file `filename` is the above mentioned executable that should contain the Fortran 90 subroutine or function `procedure`. The file `filename` should reside in the same directory where the solver input file is. Else, the relative or absolute path to that file should be added in front of the entry `filename`.

18.2 Basic Elmer Functions and Structures

In order to provide a better understanding for the following mainly example-based explanation some of the most often needed functions and routines provided by Elmer shall be discussed in this section. Most of these routines and functions are defined in the Fortran 90 module `DefUtils`. It has to be included in the code provided by the user giving the keyword

```
USE DefUtils
```

It is important to notice that – due to the nature of the Finite Element Method – the basic data structure in the Elmer Solver is the single element, rather than single points. That simplifies data manipulation in solver subroutines, but makes things a little bit more difficult if dealing with the coding of pointwise defined boundary and initial condition as well as body forces and parameter functions. In the Elmer Solver the type `Element_t` contains information on elements.

18.2.1 How to Read Values from the Input File

In the Elmer Solver the entries of each section of the solver input file – such as material, body force and initial condition – are accessed via pointer of the defined data type `ValueList_t`, further referred to as “list”. A list provides access to all the information that has been passed to the Elmer Solver from the solver input files, related to the specific section.

The principal connection between the solver input file and the access from a user function is depicted in Fig. 18.1

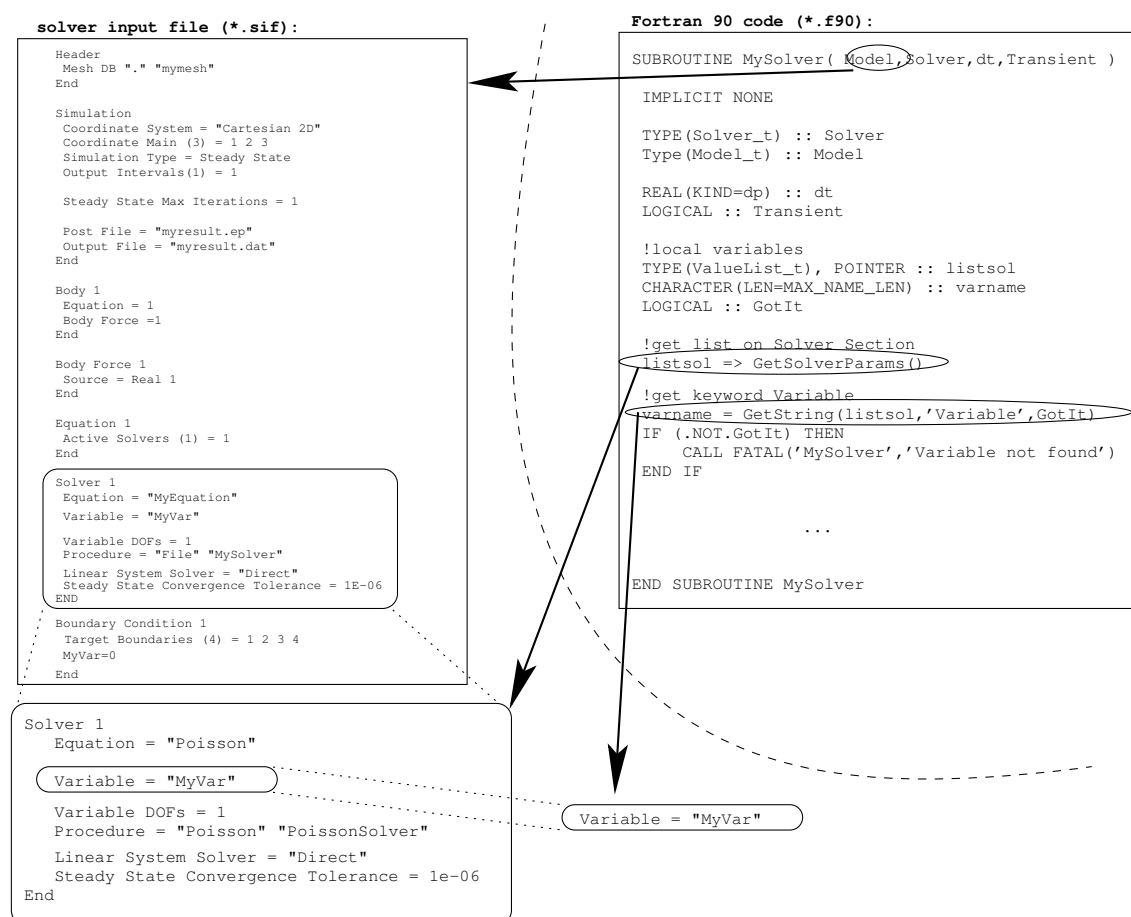


Figure 18.1: Scheme of the access of structures in the solver input file from a Fortran 90 user subroutine. The example shows, how a string is read in from the `Solver` section.

How to Access Different Sections

The following table shows the definition of the functions defined in the module `DefUtils` to provide the correct list for parameters and constants concerning the simulation and solvers

function	corresponding section
<code>GetSimulation()</code>	Simulation
<code>GetConstants()</code>	Constants
<code>GetSolverParams()</code>	Solver 1, ...

For instance, the following source code lines provide access to the entries in the simulation section of the solver input file

```
! variable declaration for pointer on list
TYPE(ValueList_t), POINTER :: Simulation
...
! assign pointer to list
Simulation => GetSimulation()
...
```

Lists that provide information connected to a certain element are

function	corresponding section
GetMaterial(Element, Found)	Material 1,...
GetBodyForce(Element, Found)	Bodyforce 1,...
GetEquation(Element, Found)	Equation 1,...
GetBC(UElement)	Boundary Condition 1,...

In the first three of these functions shown above the optional variable `Found` of type `LOGICAL` is set to `.TRUE.` upon successful search in the solver input file. Hence, it can be used for error handling. The arguments `Element` and `UElement` are of type `Element_t`. If writing a solver, the current element is known and hence can directly be passed to the functions listed above. Else, this argument may also be omitted. However, Elmer Solver needs to have the information upon the element in order to inquire the number of the material/bodyforce/equation/boundary condition section from the solver input file. Hence, if this function argument is dropped, Elmer Solver falls back to the structure `CurrentModel % CurrentElement`, which by the active solver has to be assigned to the address of the current element (see section 18.4).

The functions for input of different values from the solver input file need the assigned pointer to the corresponding to the specific section.

Reading Constants from the Solver Input File

The following value types are defined for the solver input file:

Value in Input File	Variable in Elmer Solver
Real	Real (KIND=dp)
Integer	INTEGER
Logical	LOGICAL
String	CHARACTER (LEN=MAX_NAME_LEN)
File	CHARACTER (LEN=*)

The defined interface of such a function is

```
FUNCTION FunctionName( List, Name, Found ) Result(x)
  TYPE(ValueList_t), POINTER :: List
  CHARACTER(LEN=*) :: Name
  LOGICAL, OPTIONAL :: Found
```

The arguments have the following purpose

List	List from which the value has to be read. This pointer has to be obtained by one of the previously introduced functions
Name	The keyword in the particular section for the value
Found	Optional boolean variable that contains the value <code>.TRUE.</code> upon successful read in

The type of the returned of value, `x`, is depending on the function. The following functions are declared in the `DefUtils` module:

- A value of type `REAL` is read in using the function


```
REAL(KIND=dp) :: r
...
r = GetConstReal( List, Name, Found )
```

- A variable of type INTEGER is read in using the function

```
INTEGER :: i
...
i = GetInteger( List, Name, Found )
```

- A string is read into a user function or solver by the following code line

```
CHARACTER(LEN=MAX_NAME_LEN) :: str
...
str = GetString( List, Name, Found )
```

It is important to note that these routines are only meant for reading in constant values. Consequently, these values must not be dependent on other variables.

Reading Mesh-values from the Solver Input File

The previously introduced function `GetConstReal` is defined for reading in a constant value of type `REAL(KIND=dp)`. In the case if values have to be obtained for nodes of an element defined on the mesh (e.g., an initial condition, a boundary condition or a material parameter), the following function has to be used

```
FUNCTION GetReal( List, Name, Found, UElement ) RESULT(x)
  TYPE(ValueList_t), POINTER :: List
  CHARACTER(LEN=*) :: Name
  LOGICAL, OPTIONAL :: Found
  TYPE(Element_t), OPTIONAL, TARGET :: UElement
  REAL(KIND=dp) :: x(CurrentModel % CurrentElement % Type % NumberOfNodes)
```

The returned value, `x`, is a one-dimensional array of type `REAL(KIND=dp)` with entries for every node of the either given element `UElement` or alternatively the default structure `CurrentModel % CurrentElement`. For instance, reading in the material parameter `Viscosity` from an already assigned pointer of type `ValueList_t` for a given element, `Element`, is done by the following code lines

```
REAL(KIND=dp), ALLOCATABLE :: viscosity(:)
INTEGER :: NoNodes
TYPE(ValueList_t), POINTER :: Material
TYPE(Element_t), POINTER :: Element
LOGICAL :: Found
...
allocate viscosity, set pointers Material and Element
...
NoNodes = GetElementNOFNodes(Element)
...
viscosity(1:NoNodes) = GetReal(Material, 'Viscosity', Found, Element)
```

The user has to make sure that the array that later contains the nodal values is of sufficient size. This, for instance, can be guaranteed by allocating it to the maximal occurring number of nodes for an element in the model

```
ALLOCATE(viscosity(CurrentModel % MaxElementNodes))
```

Physical Time as Argument of User Function

If a user function needs physical time as an input, it can be passed as an argument. For instance, if a boundary condition for the normal component of the velocity would have the physical time as the input variable, the function call in the solver input file then would look as follows (see section 18.3 for more details on user functions)

```
Boundary Condition BCNo
  Name = "time_dependent_outlet"
  Target Boundaries = BoundaryNo
  Normal-Tangential Velocity = True
  Velocity 2 = 0.0
  Velocity 1
    Variable Time
    Real Procedure "executable" "timeOutletCondition"
  End
End
```

Here the entries *BCNo* and *BoundaryNo* have to be replaced by the correct boundary condition and boundary target number. The file *executable* should contain the compiled user function *timeOutletCondition*.

18.2.2 How to Communicate with Structures Inside Elmer Solver

Often it is necessary to get information from inside the Elmer Solver, such as mesh coordinates or field variables associated to another solver procedure. If writing a solver subroutine, all information of that kind is accessible via the type `TYPE(Solver_t) :: Solver`. In the case of a user function (boundary condition, initial condition, material parameter), the default structure `CurrentModel % Solver` has to be used.

Inquiring Information on the Element

As mentioned earlier, most of the pre-defined functions and subroutines inside Elmer Solver apply on the whole element rather than on single nodes. Information on elements can be accessed via the pre-defined type `Element_t`. We list the functions/subroutines for the mostly needed purposes:

- Setting the active element (bulk):

```
TYPE(Element_t), POINTER :: Element
Type(Solver_t), Target :: Solver
INTEGER :: ElementNumber
...
Element => GetActiveElement(ElementNumber)
```

The argument *Solver* is optional. If it is not given, `CurrentModel % Solver` is used. This function also automatically sets the pointer `CurrentModel % CurrentElement` to the element with the given element number *ElementNumber*. This is important if sub-sequentially called functions rely on this default value to be set.

The total number of active bulk elements for a specific solver is to be inquired using the value `Solver % NumberOfActiveElements`.

- Setting the active element (boundary):

```
TYPE(Element_t), POINTER :: BoundaryElement
INTEGER :: ElementNumber
...
BoundaryElement => GetBoundaryElement(ElementNumber)
```

This routine also sets the structure `CurrentModel % CurrentElement` to the boundary element.

In contrary to the domain (i.e., bulk) it is a priori not known which boundary element is part of a boundary condition of a specific solver. This information may be obtained using the function

```
Type(Element_) :: BoundaryElement
LOGICAL :: IsActiveBoundary
...
IsActiveBoundary = BoundaryElement(BoundaryElement, Solver)
```

where both arguments are optional. If they are omitted, Elmer Solver takes the values `CurrentModel % CurrentElement` and `CurrentModel % Solver`, respectively. The boundary element number, `ElementNumber` may vary between 1 and the maximum value

```
Solver % Mesh % NumberOfBoundaryElements
```

- Inquire number of nodes in an element:

```
INTEGER :: N
TYPE(Element_t) :: Element
...
N = GetElementNOFNodes( Element )
```

The argument `Element` is optional. The default value is `CurrentModel % CurrentElement`

- Get nodal coordinates for element:

```
TYPE(Nodes_t) :: ElementNodes
TYPE(Element_t) :: Element
TYPE(Solver_t) :: Solver
...
CALL GetElementNodes( ElementNodes, Element, Solver )
```

The arguments `Element` and `Solver` are optional. The default values are `CurrentModel % CurrentElement` and `CurrentModel % Solver`, respectively. The argument `ElementNodes` is of the pre-defined type `Nodes_t`. The different components of the coordinates for the *i*-th node can be accessed by

```
REAL(KIND=dp) :: Xcoord, Ycoord, Zcoord
...
Xcoord = ElementNodes % x(i)
Ycoord = ElementNodes % y(i)
Zcoord = ElementNodes % z(i)
```

They correspond to the axes of the defined coordinate system in the solver input file.

- Get local coordinates of the *i*-th node for assigned element:

```
REAL(KIND=dp) :: U, V, W
TYPE(Element_t), POINTER :: Element
INTEGER :: i
...
U = Element % Type % NodeU(i)
V = Element % Type % NodeV(i)
W = Element % Type % NodeW(i)
```

Local coordinates are corresponding to the position inside the prototype element that is used inside the Elmer Solver. They are important if parameter values have to be obtained by summation over the basis functions.

- Get normal vector at the i -th node of the assigned boundary element:

```
REAL(KIND=dp) :: U, V, Normal(3)
TYPE(Element_t), POINTER :: BoundaryElement
LOGICAL :: CheckDirection
...
U = BoundaryElement % Type % NodeU(i)
V = BoundaryElement % Type % NodeV(i)
Normal = NormalVector( BoundaryElement, Nodes, U, V, CheckDirection )
```

The function needs the boundary element as well as the local coordinates of the point, where the surface (edge) normal shall be evaluated. The optional last argument, *CheckDirection*, is a boolean variable. If set to *.TRUE.*, the direction of the normal is set correctly to the rules given in section 18.3.2. The surface normal is returned in model coordinates and is of unity length.

Inquiring Nodal Values of Field Variables

Nodal values for an element of a scalar variables are read by the subroutine

```
SUBROUTINE GetScalarLocalSolution( x,name,UElement,USolver )
  REAL(KIND=dp) :: x(:)
  CHARACTER(LEN=*) , OPTIONAL :: name
  TYPE(Solver_t) , OPTIONAL, TARGET :: USolver
  TYPE(Element_t), OPTIONAL, TARGET :: UElement
```

The returned value is an array containing the nodal values of the variable *name*. If this variable name is not provided, it is assumed that the corresponding solver *USolver* has only one variable with a single degree of freedom. If the optional parameters *USolver* and *UElement* are not provided, then the default values *CurrentModel % Solver* and *CurrentModel % CurrentElement*, respectively, are used.

For instance, the following code lines read in the nodal element values for the variable *Temperature* (from the heat solver)

```
REAL(KIND=dp), ALLOCATABLE :: localTemp(:)
ALLOCATE(localTemp(CurrentModel % MaxElementNodes))
...
CALL GetScalarLocalSolution(localTemp, 'Temperature')
```

In the case of a vector field variable, the analog function *GetVectorLocalSolution* has to be used. For instance, if the user wants to read in the local velocity of an deforming mesh (from the *MeshSolver*), the following syntax has to be applied

```
REAL(KIND=dp) , ALLOCATABLE :: localMeshVelocity(:, :)
ALLOCATE(localMeshVelocity(3, Solver % Mesh % MaxElementNodes))
...
CALL GetVectorLocalSolution( MeshVelocity, 'Mesh Velocity')
```

Inquiring Values of Field Variables for the Whole Mesh

Sometimes, the user also would like to have values for a field variable of the complete domain accessible. This is done by assigning a pointer to the variable using the function *VariableGet*

```
VariablePointer => VariableGet( Solver % Mesh % Variables, 'Variable' )
```

The argument *Variable* has to be replaced by the variable name. The returned pointer is of type *Variable_t*. This type contains the following components

component	purpose
INTEGER, POINTER :: Perm(:)	Contains permutations for the variable. Since Elmer Solver tries to optimize the matrix structure, the numbering of the nodal values of the variable usually does not coincide with the numbering of the mesh-nodes. The is used to identify the mesh-node for a variable-entry. Hence, the field <code>VariablePointer % Perm(i)</code> contains the nodal number of the <i>i</i> -th value of the field variable <i>Variable</i> .
INTEGER :: DOFs	Contains the amount of vector components of the variable. DOFs is 1 in case of a scalar, 2 or 3 in case of a two- or three-dimensional vector field.
REAL(KIND=dp), POINTER :: Values(:)	contains the values of the field variable

For instance, in order to get access to the temperature field (similar as in the example above), the following code lines may be used

```

TYPE(Variable_t), POINTER :: TempVar
INTEGER, POINTER :: TempPerm(:)
REAL(KIND=dp), POINTER :: Temperature(:)
INTEGER :: ElementNo, N
REAL(KIND=dp), ALLOCATABLE :: localTemp(:)
ALLOCATE(localTemp(CurrentModel % MaxElementNodes))
TYPE(Element_t), POINTER :: Element
...
TempVar => VariableGet( Solver % Mesh % Variables, 'Temperature' )
IF ( ASSOCIATED( TempVar ) ) THEN
    TempPerm    => TempVar % Perm
    Temperature => TempVar % Values
!!!! stop if temperature field has not been found !!!!
ELSE IF
    CALL Fatal('MyOwnSolver', 'No variable Temperature found')
END IF
...
DO ElementNo = 1, Solver % NumberOfActiveElements
    Element => GetActiveElement(ElementNo)
    N = GetElementNOFNodes()
    NodeIndexes => Element % NodeIndexes
    localTemp(1:N) = Temperature(TempPerm(Element % NodeIndexes))
END DO

```

It is recommended to check whether the pointer to the variable has been assigned correctly. In our little example the call of the routine `Fatal` would stop the simulation run if the assessment would lead to a negative result.

Inquiring the Current Time

In certain situations in transient simulations the physical time might be needed in a user function. In Elmer Solver the physical time is treated as a variable and hence has to be read in using the type `Variable_t`

```

TYPE(Variable_t), POINTER :: TimeVar
Real(KIND=dp) :: Time
...
TimeVar => VariableGet( Solver % Mesh % Variables, 'Time' )
Time = TimeVar % Values(1)

```

How to Post Messages

Including PRINT or WRITE statements to stdio in numeric-codes can lead excessive output (large number of iterations and/or large model sizes) and consequently to a reduction in performance. It is recommended to use stdio-output routines provided by Elmer Solver, for which to a certain extend the amount of output can be controlled via the solver input file. The three pre-defined subroutines for output of messages are:

- Info is used for displaying messages (e.g., upon convergence) on the screen. The syntax is

```
CALL Info('FunctionName','The displayed message', level=levelnumber)
```

The first string shall indicate which function the displayed message comes from. The second entry is a string that contains the message itself. The integer value *levelnumber* indicates the importance of the message, starting from 1 (most important). The maximum level of messages being displayed can be determined in the simulation section of the solver input file

```
max output level = 3
```

- Warn is used for displaying warnings. Warnings are always displayed and should be used if conditions in the code occur that might lead to possible errors

```
CALL Warn('FunctionName','The displayed warning')
```

- Fatal is used to terminate the simulation displaying a message. Consequently, it is used in conditions in the code where a continuation of the computation is impossible

```
CALL Fatal('FunctionName','The displayed error message')
```

Of course the strings do not have to be hard-coded but can be composed during run-time, using the WRITE command. The string variable Message is provided for that purpose

```
WRITE(Message,formatstring) list, of, variables
CALL Info('FunctionName',Message, level=3)
```

The format-string has to be set according to the list of variables.

18.3 Writing a User Function

User functions can be used to provide a pointwise (not element wise!) value for a certain property. They are used for setting values of material properties inside the domain and/or to set values for certain variables on boundary conditions at the domain boundary.

The defined interface for a user function looks as follows

```
FUNCTION myfunction( model, n, var ) RESULT(result)
  USE DefUtils
  IMPLICIT None
  TYPE(Model_t) :: model
  INTEGER :: n
  REAL(KIND=dp) :: var, result

  ... contents of the function
  leading to a value for variable result ...

END FUNCTION myfunction
```

The function returns the value `result`. If this is supposed to be a nodal value of a material property or a variable condition, the variable type in Elmer has to be real double precision, hence, `REAL (KIND=dp)`.

The first argument passed to the function, `model`, is of the declared type `Model_t`. It is basically the interface to all information accessible at runtime, such as variable values, mesh coordinates and boundary conditions.

The second argument, `n`, is the integer number of the node for which the function - i.e., the value `result` - is evaluated. Through the last argument, `var`, a value of an input variable on which the result depends is passed to the function. In the solver input file this variable is indicated using the `Variable`-keyword. For instance (see examples later in this chapter), if the user wants to provide the function above with the result being the density which – in this particular case – depends on the temperature, the syntax looks as follows

```
Density = Variable Temperature
Procedure "filename" "myfunction"
```

Mind that there always has to be an input variable to be given using this keyword. In the case that there is no need for an input, this may occur as a dummy argument in the function.

18.3.1 User Functions for Domain Properties

In the following we will give an outline of the main issues concerning the preparation of a user function for a domain property. This might be of use if a material parameter (from material section of the solver input file), an initial condition (from solver input file initial condition section) or a body force (from solver input file body force section) of somewhat higher complexity has to be defined for the domain.

Some basic aspects concerning the syntax of such functions shall be explained by the following examples:

Example: Viscosity as a Function of Temperature

This example demonstrates the following topics:

- definition of a material property dependent on one input variable
- how to read in a value from the material section of the solver input file

We want to implement the following relation for the dynamic viscosity, μ , of a fluid

$$\mu(T) = \mu_{293K} \exp \left[C \cdot \left(\frac{293}{T} - 1 \right) \right] \quad (18.1)$$

where T is the temperature. The parameters μ_{293K} (i.e., the reference viscosity at 293 Kelvin) and C have to be read into our function from the material section of the solver input file. Thus, the material section looks as follows:

```
! Material section (ethanol)
! -----
Material 1
...
Viscosity = Variable Temperature
Procedure "fluidproperties" "getViscosity"
...
Reference Viscosity = Real 1.2E-03
Parameter C = Real 5.72
...
End
```

The values $\mu_{293K} = 1.2 \cdot 10^{-3}$ and $C = 5.72$ are the numerical values of the parameter occurring in (18.1) for pure ethanol. The executable containing the function named `getViscosity` will be called `fluidproperties`. The header – including the variable declaration – of the function then reads as follows:

```

!-----
! material property function for ELMER:
! dynamic fluid viscosity as a function of temperature
!-----
FUNCTION getViscosity( model, n, temp ) RESULT(visc)
  ! modules needed
  USE DefUtils

  IMPLICIT None

  ! variables in function header
  TYPE(Model_t) :: model
  INTEGER :: n
  REAL(KIND=dp) :: temp, visc

  ! variables needed inside function
  REAL(KIND=dp) :: refVisc, C
  Logical :: GotIt
  TYPE(ValueList_t), POINTER :: material

```

In order to get the pointer to the correct material-list, we use the function `GetMaterial`

```

! get pointer on list for material
material => GetMaterial()
IF (.NOT. ASSOCIATED(material)) THEN
  CALL Fatal('getViscosity', 'No material found')
END IF

```

We immediately check, whether the pointer assignment was successful. In the case of the NULL-pointer being returned, the pre-defined procedure `Fatal` will stop the simulation and `stdio` will display the message: `(getViscosity): No material-id found`

The next step is to read in the numerical values for the parameter:

```

! read in reference viscosity
refvisc = GetConstReal( material, 'Reference Viscosity', GotIt)
IF(.NOT. GotIt) THEN
  CALL Fatal('getViscosity', 'Reference Viscosity not found')
END IF

! read in parameter C
C = GetConstReal( material, 'Parameter C', GotIt)
IF(.NOT. GotIt) THEN
  CALL Fatal('getViscosity', 'Parameter C not found')
END IF

```

The variable `GotIt` contains the state of success of the input. In the case of unsuccessful read-in (i.e., the variable `GotIt` shows the value `.FALSE.`), the simulation will be stopped by the routine `Fatal`.

Finally, after a check upon the absolute temperature being larger than zero, the viscosity is computed according to equation (18.1)

```

! compute viscosity
IF (temp <= 0.0D00) THEN ! check for physical reasonable temperature
  CALL Warn('getViscosity', 'Negative absolute temperature.')
  CALL Warn('getViscosity', 'Using viscosity reference value')
  visc = refVisc(1)
ELSE

```



```

        visc = refVisc * EXP (C *(2.93D02/temp - 1.0D00))
    END IF

```

```

END FUNCTION getViscosity

```

In the case of negative absolute temperature, the reference value will be returned. The pre-defined routine Warn will display a specific warning on stdio.

Example: Body Force as a Function of Space

For the use as body force for the solver presented in 18.4 (i.e. heat source distribution for heat conduction equation), we want to write a function that provides a scalar in the domain as a function of space. This example demonstrates the following topics:

- definition of a scalar in the domain as function of space in two dimensions
- how to inquire the dimension of the problem

We want to implement the following two-dimensional spatial distribution for the scalar h :

$$h(x, y) = -\cos(2\pi x) \cdot \sin(2\pi y), \quad x, y \in [0, 1] \quad (18.2)$$

where x corresponds to `Coordinate 1` and y to `Coordinate 2` of the solver input file.

Since the function given in (18.2) is dependent on two input variables, the single argument that is able to be passed to the function is not sufficient. Hence it will just be used as dummy argument. Consequently, the user can provide any (existing!) variable as argument in the solver input file. The header reads as follows

```

!-----
! body force function for ELMER:
!   scalar load as function of coordinates x and y
!               -cos(2*pi*x)*sin(2*pi*y)
!-----
FUNCTION getLoad( model, n, dummyArgument ) RESULT(load)
    ! modules needed
    USE DefUtils

    IMPLICIT None

    ! variables in function header
    TYPE(Model_t) :: model
    INTEGER :: n
    REAL(KIND=dp) :: dummyArgument, load

    ! variables needed inside function
    INTEGER :: DIM
    REAL(KIND=dp) :: x, y
    Logical :: FirstVisited = .TRUE.

    ! remember these variables
    SAVE DIM, FirstVisited

```

Further we want to check whether the problem is two-dimensional and hence suitable for our function. This is done only the first time the function is called, since - hopefully - the dimension of the problem does not change during all the following calls. The function returning the problem dimension is `CoordinateSystemDimension()`

```

    ! things done only the first time the function is visited
    IF (FirstVisited) THEN

```

```

! check whether we are dealing with a 2d model
DIM = CoordinateSystemDimension()
IF (DIM /= 2) THEN
    CALL FATAL('getLoad', 'Dimension of model has to be 2d')
END IF
FirstVisited = .FALSE.
END IF

```

The next step to inquire the coordinates of the current point the function is evaluated for. The structure `model` contains that information

```

! get the local coordinates of the point
x = model % Nodes % x(n)
y = model % Nodes % y(n)

```

Finally, the result is computed

```

! compute load
load = -COS(2.0D00*PI*x) * SIN(2.0D00*PI*y)

```

```
END FUNCTION getLoad
```

Figure 18.2 shows the result of a simulation using the solver defined in section 18.4 together with the function `getLoad` as body force. The entry in the body force section of the solver input file reads as follows:




Figure 18.2: Result obtained with the routine `getLoad` as body force input for the heat conduction solver presented in 18.4. The z -coordinate is set proportional to the result obtained in the x - y plane.

```

Body Force 1
Heat Source
Variable Temp      !just a dummy argument

```

```
Real Procedure "load" "getLoad"
End
```

where the shared object file has been given the name `load`. All boundaries are set to adiabatic boundary condition, i.e., $\nabla T \cdot \vec{n} = 0$. This is possible if – and only if – the integral over the load vector over the whole domain balances to zero, like in our case. Since no Dirichlet condition is set, the result is not unique and contains an undetermined offset T_0 .

18.3.2 User Functions for Boundary Conditions

As for the user functions for bulk properties presented in section 18.3.1, the function for a boundary property is evaluated pointwise. Hence, the identical function interface is used. The major difference between bulk elements and elements on boundaries are, that the latter can be associated to more than one body. That is the case on boundaries between bodies. This is important in such cases where the boundary condition is dependent on properties inside a specific body to which the boundary is attached. Possible configurations of boundaries are depicted in Fig. 18.3.




Figure 18.3: Possible constellations for boundaries and the normal vectors, \vec{n} at element nodes. Mind, that for a body-body interface (case b) the default orientation of the surface normal may vary from element to element. At boundary points that have a discontinuous first derivative of the surface (i.e. at corners and edges), the evaluation of the normal at the same point for two adjacent boundary elements leads to different surface normals. Parent elements of boundary elements for the specific body are marked as filled.

If the keyword for checking the directions in the function `NormalVector` (see section 18.2.2) is set to `.TRUE.`, the following rules apply:

- In the case of an outside boundary the surface normal, \vec{n} , is always pointing outwards of the body.
- By default on a body-body boundary, the orientation is such that the normal always is pointing towards the body with lower density, ϱ , which is evaluated by comparison of the values given for the keyword `Density` in the corresponding material sections of the adjacent bodies.

In certain cases, if densities on both sides are either equal or varying functions of other variables, this may lead to a varying orientation of the surface. This effect and the effect of different directions of surface normal for the same point at surface corners and edges is depicted in Fig. 18.3. Whereas the latter effect can only be dealt with by either producing smooth surfaces or averaging surface normals, the first problem can

be overcome by providing the keyword `Normal Target Body` in the solver input file. This keyword defines the direction of the surface normals for the pre-defined subroutine `NormalVector`. For instance, the following sequence fixes the surface normal at boundary condition number 1 to point into body number 2

```
Boundary Condition 1
  Target Boundaries = 1
  Normal Target Body = Integer 2
...
End
```

Example: Idealized Radiation Law for External Heat Transfer

As an illustrative example we want to show how to implement a very idealized radiation boundary condition for heat transfer. This example explains:

- how to identify the type of boundary according to Fig. 18.3
- how to get material parameters from domain parent elements of the boundary element
- how to identify the local node number in an element
- how to inquire boundary normals

The net flux into the body at such a boundary shall be approximated by

$$q_n = \epsilon q_{\text{ext}} - \epsilon \sigma \cdot (T^4 - T_{\text{ext}}^4), \quad (18.3a)$$

where T_{ext} is the external temperature, σ stands for the Stefan-Boltzmann constant and ϵ is the emissivity. The external heat flux shall be defined as

$$q_{\text{ext}} = \begin{cases} -I \vec{s} \cdot \vec{n}, & \vec{s} \cdot \vec{n} < 0, \\ 0, & \text{else,} \end{cases} \quad (18.3b)$$

where I is the intensity and \vec{s} the direction of the incoming radiation vector. The surface normal, \vec{n} , is assumed to point outwards the body surface.

Since we are planning to use this boundary condition in connection with the solver presented in section 18.4.2, we have to provide the load vector $l = q_n / (c_p \varrho)$ occurring in the force vector of (18.7). This means that we have to inquire the material parameters c_p and ϱ for the parent element from the material section of the adjacent body.

The header of the function reads as

```
!-----
! boundary condition function for ELMER:
!   simplified radiation BC
!-----
FUNCTION simpleRadiation( model, n, temp ) RESULT(load)
  ! modules needed
  USE DefUtils

  IMPLICIT None

  ! variables in function header
  TYPE(Model_t) :: model
  INTEGER :: n
  REAL(KIND=dp) :: temp, load

  ! variables needed inside function
```

```

REAL(KIND=dp), ALLOCATABLE :: Density(:), HeatCapacity(:), ExtTemp(:)
REAL(KIND=dp), POINTER :: Work(:, :)
REAL(KIND=dp) :: radvec(3), Normal(3), NormDir, U, V, &
    Emissivity, normalabsorbtion, emission, StefanBoltzmann
INTEGER :: DIM, other_body_id, nboundary, nparent, &
    BoundaryElementNode, ParentElementNode, istat
Logical :: GotIt, FirstTime=.TRUE., Absorption = .TRUE.
TYPE(ValueList_t), POINTER :: ParentMaterial, BC
TYPE(Element_t), POINTER :: BoundaryElement, ParentElement
TYPE(Nodes_t) :: ElementNodes

SAVE FirstTime, Density, HeatCapacity, ExtTemp
! -----

```

The boundary element and the pointer to the list of the corresponding Boundary Condition-entry in the solver input file are assigned

```

! -----
! get element information
! -----
BoundaryElement => CurrentModel % CurrentElement
IF ( .NOT. ASSOCIATED(BoundaryElement) ) THEN
    CALL FATAL('simpleRadiation', 'No boundary element found')
END IF

BC => GetBC()
IF ( .NOT. ASSOCIATED(BC) ) THEN
    CALL FATAL('simpleRadiation', 'No boundary condition found')
END IF

```

Thereafter, a case distinction between the two possible constellations in Fig. 18.3 (i.e., outside or body-body boundary). A outside boundary is indicated by the value -1 of BoundaryElement % BoundaryInfo % outbody. In the case of a boundary-boundary interface the surface normal is supposed to point outwards. I.e., the body hosting the parent element is taken the one for which ParentElement % BodyId does not coincide with BoundaryElement % BoundaryInfo % outbody

```

other_body_id = BoundaryElement % BoundaryInfo % outbody
! only one body in simulation
! -----
IF (other_body_id < 1) THEN
    ParentElement => BoundaryElement % BoundaryInfo % Right
    IF ( .NOT. ASSOCIATED(ParentElement) ) &
        ParentElement => BoundaryElement % BoundaryInfo % Left
    ! we are dealing with a body-body boundary
    ! and assume that the normal is pointing outwards
    ! -----
ELSE
    ParentElement => BoundaryElement % BoundaryInfo % Right
    IF (ParentElement % BodyId == other_body_id) &
        ParentElement => BoundaryElement % BoundaryInfo % Left
END IF

! just to be on the save side, check again
! -----
IF ( .NOT. ASSOCIATED(ParentElement) ) THEN
    WRITE(Message, *) &

```

```

        'No parent element found for boundary element no. ', n
        CALL FATAL('simpleRadiation',Message)
    END IF

```

After that procedure the pointer `ParentElement` is set on the adjacent element of the boundary element inside the body for which the radiation boundary condition is evaluated.

We further need the total number of nodes in the boundary element and the parent element given by `nboundary` and `nparent`, respectively. Also the corresponding number of the boundary node number, `n`, in the parent element, `ParentElementNode`, as well as in the boundary element itself, `BoundaryElementNode`, is evaluated. At the end of this code sequence, the routine `GetElementNodes` sets the information on the nodes of the boundary element

```

! get the corresponding node in the elements
! -----
nboundary = GetElementNOFNodes(BoundaryElement)
DO BoundaryElementNode=1,nboundary
    IF ( n == BoundaryElement % NodeIndexes(BoundaryElementNode) ) EXIT
END DO
nparent = GetElementNOFNodes(ParentElement)
DO ParentElementNode=1,nboundary
    IF ( n == ParentElement % NodeIndexes(ParentElementNode) ) EXIT
END DO

! get element nodes
! -----
CALL GetElementNodes(ElementNodes, BoundaryElement)

```

The needed space for reading material parameter from the parent element as well as boundary condition parameters for the boundary element is allocated. In the case of the function being re-visited, we first do a deallocation, since the values of `nboundary` or `nparent` may change from element to element (hybrid mesh)

```

! -----
! Allocations
! -----
IF (.NOT.FirstTime) THEN
    DEALLOCATE(Density, HeatCapacity, ExtTemp)
ELSE
    FirstTime = .FALSE.
END IF
ALLOCATE(Density( nparent ), HeatCapacity( nparent ), &
        ExtTemp(nboundary), stat=istat)
IF (istat /= 0) CALL FATAL('simpleRadiation', 'Allocations failed')

```

The following code lines read the values for the parameters associated with the boundary element and the Stefan-Boltzmann constant from the solver input file

```

! -----
! get parameters from constants section
! and boundary condition section
! -----
DIM = CoordinateSystemDimension()
StefanBoltzmann = ListGetConstReal( Model % Constants, &
        'Stefan Boltzmann',GotIt)
IF (.NOT. GotIt) THEN ! default value in SI units
    StefanBoltzmann = 5.6704D-08
END IF

```

```

Emissivity = GetConstReal( BC, 'Emissivity', GotIt)
IF ((.NOT. GotIt) .OR. &
  ((Emissivity < 0.0d0) .OR. (Emissivity > 1.0d0))) THEN
  load = 0.0d00
  CALL WARN('simpleRadiation', 'No Emissivity found.')
  RETURN ! no flux without or with unphysical emissivity
END IF

ExtTemp(1:nboundary) = GetReal( BC, 'External Temperature', GotIt)
IF (.NOT. GotIt) THEN
  WRITE (Message, *) 'No external temperature defined at point no. ', n
  CALL Warn('simpleRadiation', Message)
  ExtTemp(1::nboundary) = temp
END IF

Work => ListGetConstRealArray( BC, 'Radiation Vector', GotIt)
IF ( GotIt ) THEN
  radvec = 0.0D00
  NormDir = SQRT(SUM(Work(1:DIM,1) * Work(1:DIM,1)))
  IF (NormDir /= 0.0d00) THEN
    radvec(1:DIM) = Work(1:DIM,1)*Work(4,1)/NormDir
  ELSE ! no radiation for weird radiation vector
    Absorption = .FALSE.
  END IF
ELSE ! no absorption without radiation vector
  Absorption = .FALSE.
END IF

```

If absorption of an incoming radiation vector has to be computed, the surface normal has to be inquired

```

! -----
! get surface normals ( if needed )
! -----
IF (Absorption) THEN
  U = BoundaryElement % Type % NodeU(BoundaryElementNode)
  V = BoundaryElement % Type % NodeV(BoundaryElementNode)
  Normal = NormalVector(BoundaryElement, ElementNodes, U, V, .TRUE.)
END IF

```

Thereafter, the needed material parameters are read from the material section of the solver input file that associated with the body for which the radiation boundary condition is computed

```

! -----
! get material parameters from parent element
! -----
ParentMaterial => GetMaterial(ParentElement)
! next line is needed, if the consequently read
! parameters are themselves user functions
! -----
CurrentModel % CurrentElement => ParentElement
Density(1:nparent) = GetReal(ParentMaterial, 'Density', GotIt)
IF (.NOT. GotIt) Density(1:nparent) = 1.0d00
HeatCapacity(1:nparent) = GetReal(ParentMaterial, 'Heat Capacity', GotIt)
IF (.NOT. GotIt) HeatCapacity(1:nparent) = 1.0d00

! put default pointer back to where it belongs

```

```
! -----
CurrentModel % CurrentElement => BoundaryElement
```

Since these material parameters themselves may be given in form of user functions, the default pointer `CurrentModel % CurrentElement` has to be set to the parent element upon call of the function `GetReal`.

Finally the two parts of the total normal heat flux are evaluated. The external load is obtained by dividing this `Laue` by the inquired values for `Density` and `HeatCapacity`.

```
!-----
! compute flux and subsequently load
!-----
IF (Absorption) THEN
    normalabsorbtion = -1.0D00 & ! since normal pointing outwards
    * Emissivity * DOT_PRODUCT(Normal, radvec)
    IF (normalabsorbtion < 0.0d0) &
        normalabsorbtion = 0.0d00
ELSE
    normalabsorbtion = 0.0d00
END IF

emission = StefanBoltzmann * Emissivity * &
    ( temp**4 - ExtTemp(BoundaryElementNode)**4) &
    / (HeatCapacity(ParentElementNode) * Density(ParentElementNode))

load = normalabsorbtion + emission

END FUNCTION simpleRadiation
```

Figure 18.4 shows the result of the heat conduction solver presented in section 18.4 in combination with the function `simpleRadiation` as boundary condition on two oppositely directed boundaries (boundary condition no. 1). Since the radiation vector is aligned with the x -direction and hence perpendicular with respect to these two boundaries, the absorption part vanishes for one of these boundaries. For the sake of simplicity, the material parameters ϱ , c_p and k have been set to unity. The corresponding entries of the solver input file for the boundary conditions of the case shown in Fig. 18.4 are:

```
Boundary Condition 1
    Target Boundaries(2) = 1 2
    Boundary Type = String "Heat Flux"
    External Load
        Variable Temp
        Real Procedure "radiation_flux.exe" "simpleRadiation"
    External Temperature = Real -20.0E00
    Radiation Vector(4) = Real -1.0E00 0.0E00 0.0E00 1.0E01
    Emissivity = Real 0.5
End

Boundary Condition 2
    Target Boundaries = 3
    Boundary Type = String "Given Temperature"
    Temp = Real 0
End

Boundary Condition 3
    Target Boundaries = 4
    Boundary Type = String "Adiabatic"
End
```





Figure 18.4: Result of the heat conduction solver applying the simplified radiation boundary condition described in this section. The coordinate directions as well as the number of the corresponding boundary condition section are shown. The latter can be compared to the solver input file entries shown in this section.

18.4 Writing a Solver

In Elmer an additional solver may be provided by dynamically linked subroutines. The pre-defined interface of a solver subroutine is

```
SUBROUTINE mysolver( Model, Solver, dt, TransientSimulation )
  TYPE(Model_t) :: Model
  TYPE(Solver_t) :: Solver
  REAL(KIND=dp) :: dt
  LOGICAL :: TransientSimulation
```

The first argument, `Model`, is the same structure also passed to a user function (see section 18.3). The second argument, `Solver`, is of type `Solver_t` and contains all information upon options set for this particular solver. The timestep size, `dt`, and a boolean variable, `TransientSimulation`, indicating whether the solver is to be run as a part of a transient (value `.TRUE.`) or steady state (value `.FALSE.`) simulation are further arguments of the solver subroutine.

18.4.1 Structure of a Solver

The well known structure of a linearized partial differential equation (PDE) for the scalar T in the variational formulation is

$$M_{ij} \frac{\partial T_j}{\partial t} + A_{ij} T_j = F_i, \quad (18.4)$$

with the mass matrix, M_{ij} , the stiffness matrix, A_{ij} and the force vector, F_i .

Time-stepping and the coupled solver iteration – i.e., the steady state or time level iteration of several solvers of the simulation – is taken care of by the main part of the Elmer Solver. The task of the user supplying a user defined solver subroutine is to linearize an eventually nonlinear PDE and to supply the Elmer Solver with the element-wise components for the matrices as well as the force vector.

This leads to a principle structure of a user defined solver subroutine as it is depicted in Fig. 18.5. We further will explain the construction of a user solver subroutine by discussing an example case.

18.4.2 Example: Heat Conduction Equation

In order to provide a simple example, we want to explain how the solution of the purely convective heat transport equation is implemented in Elmer. This example explains:

- reading solver parameters from the solver input file
- assembly of the local element matrix components for the domain and a Neumann condition including explanation of the most important Elmer Solver routines needed for linear system matrix assembly and solution
- how to deal with non-linear and transient PDE's
- how to implement Dirichlet boundary conditions

For constant density, ϱ , and heat capacity, c_p this equation may be written as

$$\frac{\partial T}{\partial t} - \nabla \cdot \left(\frac{k}{c_p \varrho} \nabla T \right) = \frac{h}{c_p \varrho}, \quad (18.5)$$

where T stands for the temperature, k for the heat conductivity and h is the heat source.

The variational formulation of (18.5) reads after partial integration of the conduction term and application of Green's theorem

$$\int_V \frac{\partial T}{\partial t} \gamma_i dV + \int_V \frac{k}{c_p \varrho} \nabla T \cdot \nabla \gamma_i dV = \int_V \frac{h}{c_p \varrho} \gamma_i dV + \oint_{\partial V} \frac{1}{c_p \varrho} \underbrace{(k \nabla T) \cdot \vec{n}}_{=q_n} \gamma_i dA, \quad (18.6)$$

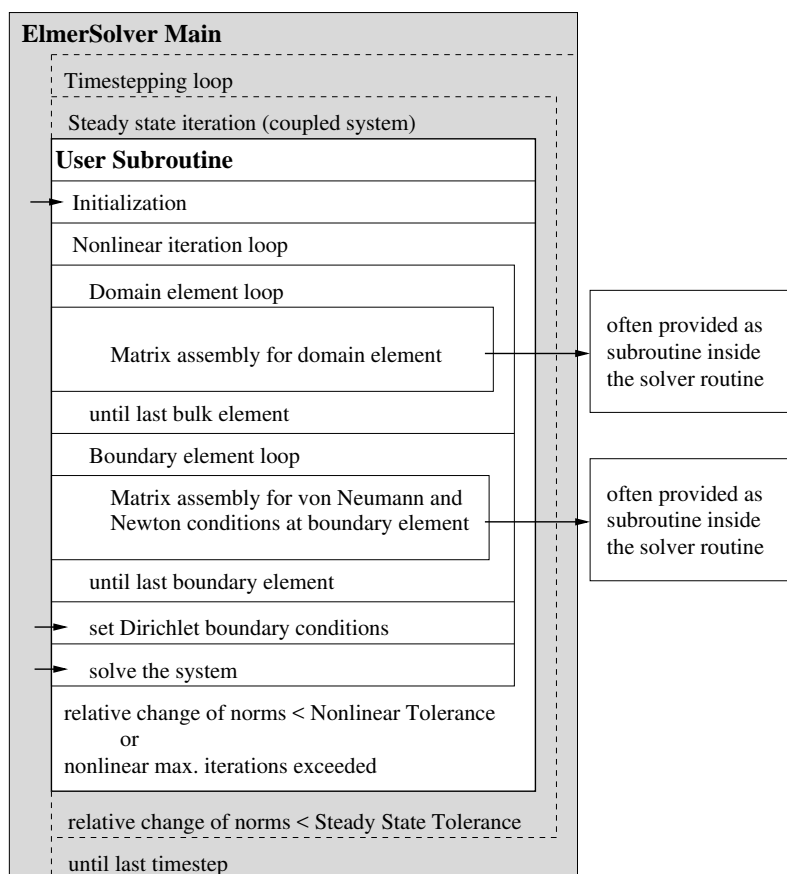


Figure 18.5: Flowchart of a user defined solver subroutine within Elmer Solver. The grey underlayed area indicates the tasks of a simulation that are provided by Elmer, whereas the white area contains the flowchart for the source code of the user subroutine. Arrows pointing into this field indicate needed subroutine/function calls to Elmer Solver.

where γ_i is the basis-function, V and ∂V is the element volume and its enclosing surface, respectively. The surface normal of ∂V is denoted by \vec{n} . According to the Galerkin method, the variable is given as $T = T_j \gamma_j$, with the sum taken over the index j . Comparing with (18.4) leads to the matrices and vectors

$$\begin{aligned} M_{ij} &= \int_V \gamma_j \gamma_i dV, \\ A_{ij} &= \int_V \frac{k}{c_p \varrho} \nabla \gamma_j \cdot \nabla \gamma_i dV, \\ F_i &= \int_V \frac{h}{c_p \varrho} \nabla \gamma_i dV + \oint_{\partial V} \underbrace{\frac{q_n}{c_p \varrho}}_{=l} \gamma_i dA. \end{aligned} \quad (18.7)$$

Although the external heat flux perpendicular to the surface normal, $q_n(T)$, in general is a function of the temperature we want to keep it formal as being prescribed. Hence, only a contribution in the force vector occurs in our formulation. Mind, that a linear or linearized expression of $q_n(T)$ directly could give a contribution to the stiffness matrix at the entries corresponding to boundary nodes. In our approach, even in the case of a linear dependency $q_n(T) \propto T$ we have to iterate the solution because of our treatment of the boundary condition.

The header contains the declaration needed variables – we tried to give them self explaining identifiers. Furthermore, allocations of the needed field arrays are done for the first time the subroutine is visited (checked by the boolean variable `AllocationsDone`). The variable names of these arrays then have to be included in the `SAVE`-statement at the end of the variable declaration block.

```
SUBROUTINE MyHeatSolver( Model, Solver, dt, TransientSimulation )

  USE DefUtils

  IMPLICIT NONE
!-----
  TYPE(Solver_t) :: Solver
  TYPE(Model_t) :: Model

  REAL(KIND=dp) :: dt
  LOGICAL :: TransientSimulation
!-----
! Local variables
!-----
  TYPE(Element_t), POINTER :: Element

  LOGICAL :: AllocationsDone = .FALSE., Found, Converged

  INTEGER :: n, t, istat, other_body_id, iter, NonlinearIter
  REAL(KIND=dp) :: Norm, PrevNorm=0.0d00, NonlinearTol, RelativeChange

  TYPE(ValueList_t), POINTER :: BodyForce, Material, BC, SolverParams
  REAL(KIND=dp), ALLOCATABLE :: MASS(:,:), STIFF(:,:), LOAD(:), FORCE(:)
  REAL(KIND=dp), ALLOCATABLE :: HeatCapacity(:), HeatConductivity(:), &
    Density(:), ExternalTemp(:), TransferCoeff(:), HeatFlux(:)

  CHARACTER(LEN=MAX_NAME_LEN) :: BoundaryType

  SAVE MASS, STIFF, LOAD, FORCE, &
    HeatCapacity, HeatConductivity, &
    Density, ExternalTemp, TransferCoeff, HeatFlux, &
```

```

        AllocationsDone, PrevNorm
!-----
!Allocate some permanent storage, this is done first time only:
!-----
IF ( .NOT. AllocationsDone ) THEN
    N = Solver % Mesh % MaxElementNodes !big enough for elemental arrays
    ALLOCATE( FORCE(N), LOAD(N), MASS(N,N), STIFF(N,N), STAT=istat )
    IF ( istat /= 0 ) THEN
        CALL Fatal( 'MyHeatSolve', &
            'Memory allocation error for matrix/vectors.' )
    END IF
    ALLOCATE( HeatCapacity(N), HeatConductivity(N), Density(N), &
        ExternalTemp(N), TransferCoeff(N), HeatFlux(N), STAT=istat )
    IF ( istat /= 0 ) THEN
        CALL Fatal( 'MyHeatSolve', &
            'Memory allocation error for parameter arrays.' )
    END IF
    AllocationsDone = .TRUE.
END IF

```

In the next step, information on the nonlinear iteration is being read from the solver section of the solver input file

```

!Read in solver parameters
!-----
SolverParams => GetSolverParams()
IF (.NOT. ASSOCIATED(SolverParams)) &
    CALL FATAL('MyHeatSolve', 'No Solver section found')
NonlinearIter = GetInteger(SolverParams, &
    'Nonlinear System Max Iterations', Found)
IF (.NOT.Found) NonlinearIter = 1
NonlinearTol = GetConstReal( SolverParams, &
    'Nonlinear System Convergence Tolerance', Found )
IF (.NOT.Found) NonlinearTol = 1.0D-03

```

Therafter, the nonlinear iteration loop (outermost loop of the white underlayed area in Fig. 18.5) is started and the linear system solver is initialized (routine DefaultInitialize)

```

!-----
! Nonlinear iteration loop
!-----
DO iter=1,NonlinearIter
    Converged = .FALSE.
    WRITE(Message,'(A,I5,A,I5)') 'Nonlinear iteration no.',iter,&
        ' of max. ', NonlinearIter
    CALL INFO('MyHeatSolve',Message,level=1)

    !Initialize the system and do the assembly:
    !-----
    CALL DefaultInitialize()

```

The next loop is over all elements in the simulation domain our solver has been assigned to (Solver % NumberOfActiveElements). The function GetActiveElement inquires the element associated with the element number t. This call at the same time also sets the default pointer CurrentModel % CurrentElement to that particular element, which is important if subsequently called functions rely on

this pointer to be set correctly (see section 18.3). After inquiring the number of nodes the nodal material parameter values $c_p \rightarrow \text{HeatCapacity}(1:n)$, $k \rightarrow \text{HeatConductivity}(1:n)$ and $\rho \rightarrow \text{Density}(1:n)$ are read in. If one of these parameters is not found (i.e., `Found == .FALSE.`), a default value of 1 will be set in order to avoid division by zero.

```
!-----
!   Assembly for the domain
!-----
      DO t=1,Solver % NumberOfActiveElements

        ! get element info
        !-----
        Element => GetActiveElement(t)
        n = GetElementNOFNodes()

        ! get material parameters
        ! -----
        Material => GetMaterial()
        IF (.NOT. ASSOCIATED(Material)) THEN
          WRITE(Message,'(A,I5,A)') &
            'No material for bulk element no. ',t,' found.'
          CALL FATAL('MyHeatSolve',Message)
        END IF
        HeatCapacity(1:n) = GetReal(Material, 'Heat Capacity', Found )
        IF (.NOT. Found) HeatCapacity(1:n) = 1.0D00
        HeatConductivity(1:n) = &
          GetReal(Material, 'Heat Conductivity', Found )
        IF (.NOT. Found) HeatCapacity(1:n) = 1.0D00
        Density(1:n) = GetReal(Material, 'Density', Found )
        IF (.NOT. Found) Density(1:n) = 1.0D00
```

In order to call the subroutine taking care of the composition of the element matrices and force vector (subroutine `LocalMatrix`), the load vector – in our case the heat source, $h \rightarrow \text{LOAD}(1:n)$ – has to be read from the body section of the solver input file. In the case of a transient simulation (indicated by `TransientSimulation == .TRUE.`) the first order time discretization is accounted for using the subroutine `Default1stOrderTime`. Mind, that also higher order time discretization routines would be at hand. The local matrix is added to the global coefficient matrix of Elmer Solver calling the subroutine `DefaultUpdateEquations`

```
      !Get load for force vector
      !-----
      LOAD = 0.0d0
      BodyForce => GetBodyForce()
      IF ( ASSOCIATED(BodyForce) ) &
        LOAD(1:n) = GetReal( BodyForce, 'Heat Source', Found )

      !Get element local matrix and rhs vector:
      !-----
      CALL LocalMatrix( MASS, STIFF, FORCE, LOAD, Element, n,&
        HeatCapacity, HeatConductivity, Density, TransientSimulation)

      !Update global matrix and rhs vector from local matrix & vector:
      !-----
      IF ( TransientSimulation ) THEN
        CALL Default1stOrderTime( MASS,STIFF,FORCE )
      END IF
```

```

CALL DefaultUpdateEquations( STIFF, FORCE )

!-----
END DO ! end Assembly for the domain
!-----

```

After the bulk elements, the contribution to the coefficient matrix and the force vector from a Neumann type of boundary condition has to be taken into account. To this end, we are looping over all boundary elements. Their total number is given by `Solver % Mesh % NumberOfBoundaryElements`. The routine `ActiveBoundaryElement` checks whether the previously inquired element is part of a boundary condition that has been assigned to our solver. If the value 1 is returned from the function `GetElementFamily` – i.e. we are dealing with boundary element given at a point element – the element also will be skipped, since Neumann conditions cannot be set on such elements. Finally, the list-pointer to the associated boundary condition section (`GetBC`) is set and the local matrices and vectors are initiated to zero.

```

!-----
!      assembly of Neumann boundary conditions
!-----
DO t=1, Solver % Mesh % NumberOfBoundaryElements

    ! get element and BC info
    ! -----
    Element => GetBoundaryElement(t)
    IF ( .NOT.ActiveBoundaryElement() ) CYCLE
    n = GetElementNOFNodes()
    ! no evaluation of Neumann BC's on points
    IF ( GetElementFamily() == 1 ) CYCLE
    BC => GetBC()

    FORCE = 0.0d00
    MASS = 0.0d00
    STIFF = 0.0d00

```

Since we have to define between different types of boundary conditions, we inquire the contents of a keyword `Boundary Type` from the solver input file. If this string is equal to 'heat flux', the variable with the name 'External Load' will be read in from the boundary condition list BC. Thereafter, the contribution to the force vector will be computed by the internal subroutine `BoundaryCondition` (see later in this code). Mind, that the external load is not the given heat flux, q_n , but its value divided by the heat capacity and the density, $l = q_n / (c_p \rho)$. This has to be taken care of if a numerical value or even a user function is provided in the solver input file (see section 18.3.2). In the case of no boundary type being found or an unknown string being detected, the natural boundary condition, zero flux perpendicular to the surface, will be assumed. This is equivalent to the 'adiabatic' boundary condition. In the case of 'given temperature' the natural boundary condition will be altered by the matrix manipulation arising from the Dirichlet boundary condition (see later in this code).

```

! check type of boundary and set BC accordingly
!-----
BoundaryType = GetString(BC, 'Boundary Type', Found)
IF (.NOT. Found) CYCLE
! natural boundary condition
IF ((BoundaryType == 'adiabatic') &
    .OR. (BoundaryType == 'given temperature')) THEN
    CYCLE
! external heat flux
ELSE IF (BoundaryType == 'heat flux') THEN
    ! get external load; mind that this is the heat flux

```

```

! divided by the density and heat capacity
LOAD(1:n) = LOAD(1:n) + GetReal(BC,'External Load', Found)
! do the assembly of the force vector
CALL BoundaryCondition(LOAD, FORCE, Element, n)
ELSE
  WRITE(Message,' (A,I3,A)') &
    'No boundary condition given for BC no. ',GetBCId(),&
    '. Setting it to adiabatic.'
  CALL WARN('MyHeatSolve',Message)
  CYCLE
END IF

```

The boundary element loop is closed after the component system matrix and vector are updated for the current boundary element.

```

IF ( TransientSimulation ) THEN
  MASS = 0.d0
  CALL Default1stOrderTime( MASS, STIFF, FORCE )
END IF

CALL DefaultUpdateEquations( STIFF, FORCE )
!-----
END DO ! end of assembly of Neumann boundary conditions
!-----

```

Before setting the Dirichlet conditions (i.e., given boundary temperature T) using the subroutine `DefaultDirichletBCs()` it is important to finish the element-wise assembly of the Elmer Solver system matrix calling `DefaultFinishAssembly`

```

CALL DefaultFinishAssembly()

! call Elmer Solver routine for Dirichlet BCs
! -----
CALL DefaultDirichletBCs()

```

The system is solved by the function call `DefaultSolve`, which returns the norm, N^n of the solution vector T_j for the n -th nonlinear iteration step. This is needed in order to inquire the change of the solution between two steps. If the relative norm

$$R = 2 \frac{|N^{n-1} - N^n|}{N^{n-1} + N^n},$$

is smaller than the given tolerance `Nonlinear System Tolerance` of the solver section, then the nonlinear iteration is taken to be converged.

```

! Solve the system
! -----
Norm = DefaultSolve()

! compute relative change of norm
! -----
IF ( PrevNorm + Norm /= 0.0d0 ) THEN
  RelativeChange = 2.0d0 * ABS( PrevNorm-Norm ) / (PrevNorm + Norm)
ELSE
  RelativeChange = 0.0d0
END IF

WRITE( Message, * ) 'Result Norm      : ',Norm

```



```

CALL Info( 'MyHeatSolve', Message, Level=4 )
WRITE( Message, * ) 'Relative Change : ',RelativeChange
CALL Info( 'MyHeatSolve', Message, Level=4 )

! do we have to do another round?
! -----
IF ( RelativeChange < NonlinearTol ) THEN ! NO
    Converged = .TRUE.
    EXIT
ELSE ! YES
    PrevNorm = Norm
END IF
!-----
END DO ! of the nonlinear iteration
!-----

```

After leaving the nonlinear iteration loop the status of convergence shall be displayed on stdio

```

! has non-linear solution converged?
! -----
IF ((.NOT.Converged) .AND. (NonlinearIter > 1)) THEN
    WRITE( Message, * ) 'Nonlinear solution has not converged', &
        'Relative Change=',RelativeChange,'>',NonlinearTol
    CALL Warn('MyHeatSolve', Message)
ELSE
    WRITE( Message, * ) 'Nonlinear solution has converged after ', &
        iter, ' steps.'
    CALL Info('MyHeatSolve', Message, Level=1)
END IF

```

In the code lines given above, the user could exchange the routine Warn by Fatal if the simulation should stop upon failed nonlinear iteration convergence.

Further we have to include the needed local subroutines using the Fortran 90 command

```

!-----
! internal subroutines of MyHeatSolver
!-----
CONTAINS

```

The subroutine LocalMatrix composes the local matrices and vectors for a bulk element. The header with the variable declaration reads as follows

```

!-----
SUBROUTINE LocalMatrix(MASS, STIFF, FORCE, LOAD, Element, n, &
    HeatCapacity, HeatConductivity, Density, TransientSimulation)
    IMPLICIT NONE
!-----
    REAL(KIND=dp), DIMENSION(:, :) :: MASS, STIFF
    REAL(KIND=dp), DIMENSION(:) :: FORCE, LOAD, &
        HeatCapacity, HeatConductivity, Density
    INTEGER :: n
    TYPE(Element_t), POINTER :: Element
    LOGICAL :: TransientSimulation
!-----
    REAL(KIND=dp) :: Basis(n), dBasisdx(n, 3), ddBasisddx(n, 3, 3)
    REAL(KIND=dp) :: detJ, LoadAtIP, &

```

```

        LocalHeatCapacity, LocalHeatConductivity, LocalDensity
LOGICAL :: Stat
INTEGER :: t,i,j,DIM
TYPE(GaussIntegrationPoints_t) :: IP

TYPE(Nodes_t) :: Nodes
SAVE Nodes

```

For the sake of simplicity we use the same identifiers as in the solver subroutine for the variables in the argument list.

The next step is to inquire the dimension of the coordinate system. Thereafter, we get the nodes of the element using the already introduced function `GetElementNodes`. Since the values in `CurrentModel` % `CurrentElement` and `CurrentModel` % `Solver` have been set, no additional arguments to the variable `Nodes` have to be set. After we have initialized the local matrix and vector components to zero, the information upon the Gauss-points needed for integration is inquired by the function `GaussPoints`. They returned variable `IP` is of type `GaussIntegrationPoints_t`.

```

DIM = CoordinateSystemDimension()

CALL GetElementNodes( Nodes )

STIFF = 0.0d0
FORCE = 0.0d0
MASS = 0.0d0

!Numerical integration:
!-----
IP = GaussPoints( Element )

```

The integration over the element is done by summing over all Gauss-points (from 1 to $N_{IP} \rightarrow IP \% n$). The square root of the determinant of the element coordinate system metric tensor $\sqrt{\det(\mathbf{J}^T \cdot \mathbf{J})} \rightarrow \text{DetJ}$ as well as the local basis functions, $\gamma_i \rightarrow \text{Basis}$, their derivatives, $\nabla \gamma_i \rightarrow \text{dBasisdx}$, are evaluated for every Gauss-point using the function `ElementInfo`. The variable `ddBasisddx` is just passed as a dummy argument, since the last argument, `getSecondDerivatives` is set to `.FALSE.`. The pointer to the element, `Element`, and its nodes, `Nodes` and the local variables of the Gauss-points `IP \% U(t)`, `IP \% V(t)` and `IP \% W(t)`, are needed as input.

```

!-----
!   Loop over Gauss-points (element Integration)
!-----
DO t=1, IP \% n
    !Basis function values & derivatives at the integration point:
    !-----
    getSecondDerivatives = .FALSE.
    stat = ElementInfo( Element, Nodes, IP \% U(t), IP \% V(t), &
        IP \% W(t), detJ, Basis, dBasisdx, ddBasisddx, &
        getSecondDerivatives)

```

Thereafter, the material parameters at the Gauss-points are evaluated, using the basis function. For instance, the local density, $\varrho|_{IP} \rightarrow \text{LocalDensity}$ at the Gauss-point is evaluated as follows:

$$\varrho|_{IP} = \varrho_i \gamma_i|_{IP},$$

with the sum taken over the nodal index i . The load vector $h/(\varrho c_p)|_{IP} \rightarrow \text{LoadAtIP}$ is evaluated in a similar way.

```

!Material parameters at integration point:
!-----
LocalHeatCapacity = SUM( HeatCapacity(1:n) * Basis(1:n) )
LocalHeatConductivity = SUM( HeatConductivity(1:n) * Basis(1:n) )
LocalDensity = SUM( Density(1:n) * Basis(1:n) )
!The source term at the integration point:
!-----
LoadAtIP = SUM( Basis(1:n) * LOAD(1:n) ) &
           / (LocalHeatCapacity * LocalDensity)

```

The force vector is obtained by the integral over the element, which is approximated by the sum over all Gauss-point contributions

$$F_j = \int_V \frac{h}{\varrho c_p} \gamma_j DV \approx \sum_{t=1}^{N_{IP}} \left(\sqrt{ds^2} \sqrt{\det(\mathbf{J}^T \cdot \mathbf{J})} \frac{h}{\varrho c_p} \gamma_j \right) |_{IP}.$$

The model coordinate system metric $\sqrt{ds^2} \rightarrow IP \% s(t)$ as well as the previously inquired element coordinate system metric $\sqrt{\det(\mathbf{J}^T \cdot \mathbf{J})} \rightarrow DetJ$ have to be taken into account.

The matrix components are evaluated in complete analogy to the force vector

$$M_{ij} = \int_V \gamma_j \gamma_i dV \approx \sum_{t=1}^{N_{IP}} \left(\sqrt{ds^2} \sqrt{\det(\mathbf{J}^T \cdot \mathbf{J})} \gamma_j \gamma_i \right) |_{IP},$$

$$A_{ij} = \int_V \frac{k}{c_p \varrho} \nabla \gamma_j \cdot \nabla \gamma_i dV \approx \sum_{t=1}^{N_{IP}} \left[\sqrt{ds^2} \sqrt{\det(\mathbf{J}^T \cdot \mathbf{J})} \frac{k}{c_p \varrho} (\nabla \gamma_j \cdot \nabla \gamma_i) \right] |_{IP},$$

where the dot product of the first derivatives of the basis function is implemented using the expression $\nabla \gamma_j \cdot \nabla \gamma_i \rightarrow \text{SUM}(\text{dBasisdx}(i, 1:\text{DIM}) * \text{dBasisdx}(j, 1:\text{DIM}))$

```

!-----
!      Loop over j-components of matrices and over force vector
!-----
DO j=1,n
    FORCE(j) = FORCE(j) + IP % s(t) * DetJ * LoadAtIP * Basis(j)
!-----
!      Loop over i-components of matrices
!-----
DO i=1,n
    !The mass matrix, if needed
    !-----
    IF (TransientSimulation) THEN
        MASS(i,j) = MASS(i,j) + IP % s(t) * DetJ * &
                        Basis(i)*Basis(j)
    END IF

    !Finally, the stiffness matrix:
    !-----
    STIFF(i,j) = STIFF(i,j) &
                + IP % s(t) * DetJ * LocalHeatConductivity &
                * SUM(dBasisdx(i,1:DIM) * dBasisdx(j,1:DIM)) &
                / (LocalDensity * LocalHeatCapacity)
!-----
END DO ! end Loop over i-components of matrices
!-----

```

```

        END DO ! end Loop over j-components of matrices and vector
!-----
        END DO ! end Loop over Gauss-points (element Integration)
!-----
        END SUBROUTINE LocalMatrix
!-----

```

The last two statements in the code sequence given above close the loop over the Gauss-points and provide the end statement of the local subroutine `LocalMatrix`.

The subroutine `BoundaryCondition` evaluates the contribution to the force vector at the boundary elements with given external load $l = q_n / (c_p \varrho) \rightarrow \text{LOAD}$

$$F_j = \oint_{\partial V} l dV \approx \sum_{t=1}^{N_{IP}} \left(\sqrt{ds^2} \sqrt{\det(\mathbf{J}^T \cdot \mathbf{J})} l \right) |_{IP}.$$

Since this is implemented in complete analogy to the assembly of the force vector in the previously discussed subroutine `LocalMatrix`, a detailed explanation can be omitted

```

!-----
SUBROUTINE BoundaryCondition(LOAD, FORCE, Element, n)
  IMPLICIT NONE
!-----
  REAL(KIND=dp), DIMENSION(:) :: FORCE, LOAD
  INTEGER :: n
  TYPE(Element_t), POINTER :: Element
!-----
  REAL(KIND=dp) :: Basis(n), dBasisdx(n,3), ddBasisddx(n,3,3)
  REAL(KIND=dp) :: detJ, LoadAtIP, &
    LocalHeatCapacity, LocalDensity
  LOGICAL :: stat, getSecondDerivatives
  INTEGER :: t, j
  TYPE(GaussIntegrationPoints_t) :: IP

  TYPE(Nodes_t) :: Nodes
  SAVE Nodes
!-----

  CALL GetElementNodes( Nodes )

  FORCE = 0.0d0

  !Numerical integration:
  !-----
  IP = GaussPoints( Element )
!-----
! Loop over Gauss-points (boundary element Integration)
!-----
  DO t=1, IP % n
    !Basis function values & derivatives at the integration point:
    !-----
    getSecondDerivatives = .FALSE.
    stat = ElementInfo( Element, Nodes, IP % U(t), IP % V(t), &
      IP % W(t), detJ, Basis, dBasisdx, ddBasisddx, &
      getSecondDerivatives)
    !The source term at the integration point:

```

```

!-----
LoadAtIP = SUM( Basis(1:n) * LOAD(1:n) )
!-----
!      Loop over j-components of matrices and over force vector
!-----
      DO j=1,n
        FORCE(j) = FORCE(j) + IP % s(t) * DetJ * LoadAtIP * Basis(j)
      END DO
    END DO
!-----
END SUBROUTINE BoundaryCondition
!-----

```

The Fortran 90 file is completed providing the end statement for the user solver subroutine `MyHeatSolver`.

```

!-----
END SUBROUTINE MyHeatSolver
!-----

```

18.4.3 User Defined Functions and Subroutines in Parallel

User defined functions (boundary and initial conditions as well as body forces) and routines (solvers) do not have to be especially adopted to be run within a `ElmerSolver_mpi` call. Nevertheless, there might be special situation that demand additional code-lines to avoid issues caused by parallel execution. To that end, one has to be able to retrieve certain information on the parallel environment. Therefore, a special type exist within Elmer with the following entries:

```

TYPE ParEnv_t
  INTEGER                :: PEs
  INTEGER                :: MyPE
  LOGICAL                :: Initialized
  INTEGER                :: ActiveComm
  LOGICAL, DIMENSION(:), POINTER :: Active
  LOGICAL, DIMENSION(:), POINTER :: IsNeighbour
  LOGICAL, DIMENSION(:), POINTER :: SendingNB
  INTEGER                :: NumOfNeighbours
END TYPE ParEnv_t

```

The following list shall give an overview of the most important operations in this context:

How many parallel processes/domains are there in the current run? This can be accessed by the construct `ParEnv % PEs`. For instance, the following line can be used to introduce a branch between instructions for a serial and a parallel Elmer run:

```

! serial run
IF ( ParEnv % PEs <= 1 ) THEN

ELSE ! parallel run

ENDIF

```

What parallel process/domain is the current task being run on? This number can be inquired via the construct `ParEnv % myPe`

What parallel process/domain my element belongs to? If the pointer `CurrentElement` is pointing to the actual element of interest, the host process/domain this `CurrentElement % partIndex`. For instance, in combination with the instruction in the previous point, this can be used to inquire whether an element actually is being processed within the domain (is part of the domain's matrix) or is just an halo-element:

```
! skip rest of routine, because element is not
! part of the domain (halo-element)
IF (ParEnv % myPe .NE. Element % partIndex) CYCLE
```

Is the parallel process/domain with index i a neighbour of my current process/domain? This is inquired by the logical value in `ParEnv % IsNeighbour(i)`

How many neighbour processes/domains does my current process/domain have? This is inquired by the logical value in `ParEnv % NumOfNeighbours`

```
+++++
```

18.5 Compilation and Linking of User Defined Routines/Functions

Elmer provides a script `elmerf90` (residing in `$ELMER_HOME/bin`) for compilation of external routines. This makes sure that the same compiler the Elmer installation has been compiled with is used for compilation of the additional routines. This is necessary, as usually modules of different Fortran compilers are incompatible. Additionally, the script contains the necessary compiler options in order to take care that the essential libraries are being linked to the resulting shared object.

Compilation of ones own code is quite straight forward. It is essential, that the wrapper `elmerf90` is in the path of the system (or alternatively called with its path preceded). If now the routines/functions are in the file `mySolver.f90` then the dynamically linked object is obtained under UNIX/Linux with the command

```
elmerf90 mySolver.f90 -o mySolver
```

Under Windows XP, compilation can be made via the ElmerGUI or manually using the command window (Windows XP: Start → Run, enter `cmd`). Preferably, the source file should be in the tree of the Elmer (e.g.,

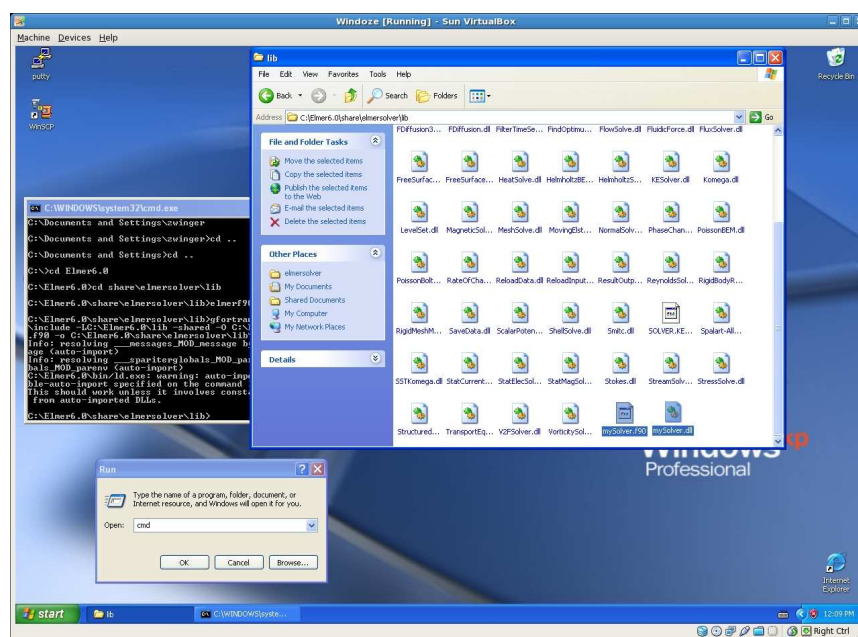


Figure 18.6: Manual compilation of additional modules in Windows XP

`C:\Elmer6.0` installation under the directory `C:\Elmer6.0\share\elmersolver\lib`. From within this directory, the command

```
elmerf90 myOwnRoutines.f90
```

should create the shared object `mySolver.dll` that should be found automatically in every Elmer run.

Appendix A

Format of mesh files

In this appendix the format of ElmerSolver mesh files is described. The mesh data are arranged into four separate files: `mesh.header`, `mesh.nodes`, `mesh.elements`, and `mesh.boundary`.

In parallel computations the names of the files are changed so that `mesh` is replaced by string `part.n` where `n` refers to the number of the partition. Additionally in parallel there is the `part.n.shared` file.

In the mesh files numeric codes are used for distinguishing different element types. For the element type codes and the node numbering order of the elements see also appendix D.

A.1 The format of header file

The header file `mesh.header` tells how many nodes and elements are present in the mesh. The lines of this file are organized as

```
nodes elements boundary-elements
nof_types
type-code nof_elements
type-code nof_elements
...
```

In the first line the numbers of nodes, elements, and boundary elements are given, while the count in the second line is the number of different element types used in the mesh. The lines which follow give the numbers of elements as sorted into different element types.

For example, the following header file

```
300 261 76
2
404 261
202 76
```

tells us that the mesh is composed of 300 nodes, 261 elements, and 76 boundary elements. Two different element types are used in the mesh: there are 261 elements of type code 404 (bilinear quadrilateral) and 76 elements of type code 202 (linear line element).

A.2 The format of node file

The file `mesh.nodes` contains node data so that each line defines one node. Each line starts with two integers followed by three real numbers:

```
n1 p x y z
n2 p x y z
...
nn p x y z
```

The first integer is the identification number for the node. The second integer is a partition index for parallel execution and is not usually referenced by the solver in the case of sequential runs. If the partition index is not of particular use, it may be set to be -1 (or 1). The real numbers are the spatial coordinates of the node. Three coordinates should always be given, even if the simulation was done in 1D or 2D. It is also noted that the nodes may be listed in any order.

A.3 The format of element file

The `mesh.elements` file contains element data arranged as

```
e1 body type n1 ... nn
e2 body type n1 ... nn
...
en body type n1 ... nn
```

Each line starts with an integer which is used for identifying the element. The integer `body` defines the material body which this element is part of. Then the element type code and element nodes are listed. For example, the element file might start with the following lines:

```
1 1 404 1 2 32 31
2 1 404 2 3 33 32
3 1 404 3 4 34 33
4 1 404 4 5 35 34
...
```

A.4 The format of boundary element file

The elements that form the boundary are listed in the file `mesh.boundary`. This file is similar to the `mesh.elements` file and is organized as

```
e1 bndry p1 p2 type n1 ... nn
e2 bndry p1 p2 type n1 ... nn
...
en bndry p1 p2 type n1 ... nn
```

The first integer is again the identification number of the boundary element. Next the identification number of the part of the boundary where this element is located is given. Whether the boundary element can be represented as the side of a parent element defined in the file `mesh.elements` is indicated using the two parent element numbers `p1` and `p2`. If the boundary element is located on an outer boundary of the body, it has only one parent element and either of these two integers may be set to be zero. It is also possible that both parent element numbers are zeros. Finally the element type code and element nodes are listed.

A.5 The format of shared nodes file

The nodes that are shared between different partitions are written in the `part.n.shared` file. The file is organized as

```
n1 np p1 p2 ... pm
n2 np p1 p2 ... pm
...
nn np p1 p2 ... pm
```

Here the first index on the row refers to a node that is shared by different partitions, the partition under study being one of the owners. Integer `np` gives the number of partitions sharing the node and this is followed by the number of the sharing partitions. By convention the owner partition is the first one of the partitions. Most shared nodes are shared just by two partitions but they can be also more owners on some nodes.

A.6 Exceptions on parallel mesh format

When the mesh is saved in parallel there is an exception that must be dealt with. Assume that you have a boundary condition for which you need to assign Dirichlet conditions. If it would happen to be so that some partition shares such boundary node without being the owner of the boundary element itself, the boundary nodes must still somehow be available for setting boundary conditions. In practice this is done so that new single-node elements of type 101 are added to the `mesh.n.boundary` list. These must be given the same `bndry` index as the initial boundary that it was deduced from.

Appendix B

Format of result output files

B.1 Format versions

Result files can be written as either ASCII text or in binary. This is controlled by the parameter

```
Binary output = logical true|false
```

in the 'Simulation' section of the .sif file. Default is `false`.

The format of the file is recorded on it's first line;¹ it's either

```
BINARY v.e
```

or

```
ASCII v
```

The *v* at denotes the version number of the format, and the *e* in the binary format denotes an endianness-marker; either L for little endian or B for big endian.

ElmerSolver can read files of older format versions for restarting, but all new files are written in the newest formats. The current formats documented here are ASCII version 1 and BINARY version 2.

B.2 General structure

Both binary and ASCII files have the following general structure. In the binary files, the header is separated from the rest by a null byte. The ASCII format has no such separator.

```
[File format version line]
[Header]
[<null byte> (binary only)]
[timestep 1]
[timestep 2]
[timestep 3]
.
.
.
[timestep n]
```

¹except for old ASCII files, that lack the format version line, and start with `"!File started at:"`

B.2.1 The header

The header looks the same for both binary and ASCII (ans is written in ASCII also for binary files):

```
!File started at: date time
Degrees of freedom:
variable 1          n1    :fs
variable 2          n2    :fs
variable 3          n3    :fs
.
.
.
Total DOFs:          nTotDOFs
Number Of Nodes:     nNodes
```

Note that in the list of variables, vectors appear both as vectors (DOF > 1) and separate components (DOF = 1).

B.2.2 Timesteps

For each time step, the time and the values of the variables are stored. For vector variables, the components are stored as separate variables, typically named *varname 1*, *varname 2*, etc.

If the parameter

```
Omit unchanged variables in output = logical true|false
```

in the 'Simulation' section of the .sif file is set to `true`, only variables whose values has changes since last time are saved. Default is `false`

For the binary format, the following type specifiers are used in this document:

```
<s(str)>    Null terminated string of characters.
<i(i4)>     4 byte integer.
<j(i8)>     8 byte integer.
<x(dbl)>    8 byte floating point number.
```

For this part of the file, the general structure of binary and ASCII files are essentially the same, with just some minor differences:

ASCII

Time: *ns nt t*
Variablename_1
 [Permutation table p_1]
 ⋮
 $Var_1(p_1(i))$
 ⋮
Variablename_2
 [Permutation table p_2]
 ⋮
 $Var_2(p_2(i))$
 ⋮
Variablename_3
 [Permutation table p_3]
 ⋮
 $Var_3(p_3(i))$
 ⋮

Binary

<Time: (str)>< $ns(i4)$ >< $nt(i4)$ >< $t(db1)$ >
 <*Variablename_1* (str)>
 [Permutation table p_1]
 ⋮
 < $Var_1(p_1(i))(db1)$ > Variable 1's values
 $\forall i \text{ s.t. } p_1(i) > 0.$
 ⋮
 <*Variablename_2* (str)>
 [Permutation table p_2]
 ⋮
 < $Var_2(p_2(i))(db1)$ > Variable 2's values
 $\forall i \text{ s.t. } p_2(i) > 0.$
 ⋮
 <*Variablename_3* (str)>
 [Permutation table p_3]
 ⋮
 < $Var_3(p_3(i))(db1)$ > Variable 3's values
 $\forall i \text{ s.t. } p_3(i) > 0.$
 ⋮

nt = time step number, ns = saved time steps number, t = time.

The permutation tables

The permutation tables can be stored in three different ways:

1. As an explicit table:

ASCII

Perm: *size np*
 ⋮
 $i \quad p(i)$
 ⋮

Binary

< $size(i4)$ >< $np(i4)$ >
 ⋮
 < $i(i4)$ >< $p(i)(i4)$ > Permutation indexes i
 and values $p(i)$
 $\forall i \text{ s.t. } p(i) > 0.$
 ⋮

$size$ = total size of the permutation table (> 0), and np = number of positive values in the table.

2. If the permutation table is the same as for the previous variable, there's no need to store it again. This case is written as

ASCII

Perm: use previous

Binary

<-1 (i4)>< $Pos(i8)$ >

Pos in the binary format is the position in bytes of the previous table.

3. No permutation table; corresponds to the case

$$size = np = nNodes, \text{ and } p(i) = i \quad \forall i.$$

This case is stored as

ASCII

Perm: NULL

Binary

<0 (i4)>

B.3 The positions file

For binary format, a positions file named '*outputfile.pos*' will be created. It contains the positions (in bytes) of the timesteps and variables in the result file, expressed as 8 byte integers. It has the following format ($nVar$ = number of variables):

```
<Endianness-marker (char)>
<nVar (i4)>
<varname 1 (str)>
<varname 2 (str)>
:
:
<varname nVar (str)>
<Pos. for Timestep 1 (i8)>
<Pos. for variable 1 (i8)>
<Pos. for variable 2 (i8)>
:
:
<Pos. for variable nVar (i8)>
<Pos. for Timestep 2 (i8)>
<Pos. for variable 1 (i8)>
<Pos. for variable 2 (i8)>
:
:
<Pos. for variable nVar (i8)>
<Pos. for Timestep 3 (i8)>
<Pos. for variable 1 (i8)>
<Pos. for variable 2 (i8)>
:
:
<Pos. for variable nVar (i8)>
:
:
```

Note: Positions are saved for *all* variables for every time step; even if an unchanged variable isn't saved to the result file for a time step, it will still have a position in the .pos file (i.e. the position of where it was saved last time). Because of this all timesteps has the same size of $(nVar + 1) \times 8$ bytes. Finding the position of the n :th variable of the m :th time step is therefore easy; it's found at the

$$(size-of-header + ((nVar + 1) \times (m - 1) + n) \times 8) : th$$

byte in the positions file.

Appendix C

Format of ElmerPost Input File

The lines of ElmerPost input file are organized as

```
nn ne nf nt scalar: name vector: name ...
x0 y0 z0
...           ! nn rows of node coordinates (x,y,z)
xn yn zn
group-name element-type i0 ... in
...           ! group data and element descriptions
group-name element-type i0 ... in
#time 1 timestep1 time1
vx vy vz p ...
...           ! nn rows
vx vy vz p ...
#time 2 timestep2 time2
vx vy vz p
...           ! nn rows
vx vy vz p ...
....
#time n timestepn timen
vx vy vz p ...
...           ! nn rows
vx vy vz p ...
```

The header

The file starts with the header line which contains the following information:

- nn: the total number of nodes
- ne: the total number of the elements including boundary elements
- nf: the total number of degrees of freedom, i.e. the total number of scalar unknowns in the model
- nt: the number of time steps for which solution data is stored
- scalar: name vector: name ... : the list which pairs variable names with their types.

The mesh and group data

After the header the node coordinates are given, each coordinate triplet on its own row. Three coordinates should be given even if the model was two-dimensional.

Group data consist of the following information:

- `group-name`: the name of the element group (having the same material, body etc.)
- `element-type`: the numeric code giving the element type; see also Appendix D.
- The numbers `i0 ... in` are the indices of the element nodes. The nodes are referenced using the row indices of the node coordinate array at the beginning of the file. The first node in the array has the index zero.

It is noted that there is also another level of element grouping that can be employed as follows

```
#group group1
  element-definitions
  ...
#group group2
  element-definitions
  ...
#endgroup group2
  element-definitions
  ...
#endgroup group1
```

The number of element groups is not restricted in any way.

The solution data

For each timestep the following solution data is written:

- `#time n t time`: `n` is the order number of the solution data set, `t` is the simulation timestep number, and `time` is the current simulation time.
- The next `nn` rows give the node values of the degrees of freedom. The values are listed in the same order as given in the header with the keywords `scalar:` and `vector:`

An example file

Here a very simple example file is given. There is only one element, three nodes, one variable, and the solution data are given for a single timestep:

```
3 1 1 1 scalar: Temperature
0 0 0
1 0 0
0 1 0
#group all
body1 303 0 1 2
#endgroup all
#time 1 1 0
1
2
3
```

Appendix D

Basic element types

The basic nodal element types which ElmerSolver can handle are the linear and quadratic elements in one, two, and three dimensions. In addition, cubic elements in one and two dimensions are also available. Each element type is associated with a unique code which consists of the element family (the count of hundreds) and the number of nodes.

- nodal element (101)
- linear (element type code 202), quadratic (203) and cubic (204) line segments
- linear (303), quadratic (306) and cubic (310) triangles, see Figure D.1




Figure D.1: The linear (303) and quadratic (306) triangular elements.

- bilinear (404), quadratic (408,409) and cubic (412,416) quadrilaterals, see Figure D.2
- linear (504) and quadratic (510) tetrahedrons, see Figure D.3
- linear (605) and quadratic (613) pyramids
- linear (706) and quadratic (715) wedges
- trilinear (808) and quadratic (820,827) hexahedrons (or bricks), see Figure D.4.

The element types are defined in the file `elements.def` and new element types may be added by adding their definitions to this file.




Figure D.2: The four-node (404) and eight-node (408) quadrilateral elements.




Figure D.3: The linear (504) and quadratic (510) tetrahedron elements.




Figure D.4: The 8-node (808) and 20-node (820) hexahedron elements.

Appendix E

Higher-order finite elements

Higher-order finite elements are elements for which the degree of basis functions is higher than 1. If better solution accuracy is sought by increasing the approximation degree, higher-order elements are also called p -elements. They differ from the usual Lagrange-type elements in that, in addition to nodal basis functions of the lowest degree, there exists additional basis functions which are associated with edges, faces and interiors of elements.

- **Side modes** get their values along some element edge. They vanish towards other edges and all nodal points of the element. Side modes are defined for all 2d and 3d elements.
- **Face modes** get their values along some element face. They vanish towards other faces and all edges and nodal points of the element. Face modes are only defined for 3d elements.
- **Internal modes** get their values inside element and vanish towards element faces, edges and nodal points. They are defined for all 1d, 2d and 3d elements.

Thus, for p -elements the usual element interpolant, now denoted as $u_{h,p}$, is the sum of nodal, edge, face and bubble interpolants as

$$u_{h,p} = u_{h,p}^v + u_{h,p}^e + u_{h,p}^f + u_{h,p}^b \quad (\text{E.1})$$

where $u_{h,p}^v$ is the lowest-order nodal interpolant while $u_{h,p}^e$, $u_{h,p}^f$ and $u_{h,p}^b$ are edge, face and bubble interpolants, respectively. Contribution of one p -element to the global system is similar to that of h -element. Naturally for higher-order elements the size of the local stiffness matrix is larger, because of the larger number of basis functions.

Properties for good p -elements include computational efficiency, at least partial orthogonality and hierarchy of basis functions. Hierarchy means that if basis for some element of a given degree p is \mathcal{B}^p it holds for $p+1$ that $\mathcal{B}^p \subset \mathcal{B}^{p+1}$. Orthogonality properties of basis functions ensure that condition number of the global stiffness matrix does not increase as dramatically as for the nodal (Lagrange) elements of higher order. This ensures good numerical stability. Some good references to higher-order finite elements in literature are [3] by Szabo and Babuska and [4] by Solin et al.

E.1 Higher-order elements in Elmer

Elements implemented in Elmer follow the ones presented in [3]. Now let us define some orthogonal polynomials based on Legendre polynomials $P_i(x)$, $i \geq 0$. The so-called Lobatto shape functions ϕ_k are defined as

$$\phi_k(\xi) = \sqrt{\frac{1}{2(2k-1)}} (P_k(\xi) - P_{k-2}(\xi)), \quad k = 2, 3, \dots \quad (\text{E.2})$$

where P_k are Legendre polynomials. Function $\phi_k(\xi)$ has two of its roots at $\xi = \pm 1$, so define now another function φ_k as

$$\varphi_k(\xi) = \frac{4\phi_k(\xi)}{1-\xi^2}, \quad k = 2, \dots, p \quad (\text{E.3})$$

Functions ϕ_k and φ_k are used to define higher order elements. Different element shapes and their basis functions are defined in Appendix E.3. The pyramidal element used in Elmer is based loosely to Devloo's representation in [2].

In Elmer elements of varying polynomial degree p may be used in the same mesh. The mesh may also consist of different element types, as the side and face modes for different element types are compatible with one another. Pyramidal and wedge elements of high order to connect tetrahedral and brick elements are also supported. To achieve the best possible converge the use of pyramidal elements in a mesh should be kept to a minimum.

The global continuity of the higher-order finite element approximation is enforced by the solver when the function `ElementInfo` is used for obtaining basis functions values for elements. To combine elements of varying degree a maximum rule is used. Thus, if two or more elements share an edge and have different polynomial degrees, the maximum of among the values is chosen as the degree of the global edge.

The approximation degree can be specified either in the solver section or in the equation section of the solver input file. The solver section may for example contain the command

```
Element = "p:4"
```

so that the spatial discretization is based on polynomials of degree 4. Basically the same command may also be used in the equation section if the intention is to apply the same basis functions for all active solvers, while a more advanced element definition command is needed if the user wants to specify the approximation degree solver-wise. One may define for example

```
Equation 1
  Active Solvers(2) = 1 2
  Element{1} = String "p:4"
  Element{2} = String "p:1"
End
```

so that the fourth-order approximation is applied in connection with the solver having the integer identifier 1, while the first-order approximation is used for the solver with the identifier 2.

A body-wise changing approximation degree can also be defined by a combined use of the body and equation sections. If we have for example split the entire computational domain into two bodies and we want to use basis functions of degree $p = 1$ in the body 1 and the polynomial degree $p = 3$ in the body 2, we may make the desired element choices by including the following definitions in the `sif`-file:

```
Body 1
  Equation = 1
  Material = 1
End
```

```
Body 2
  Equation = 2
  Material = 1
End
```

```
Equation 1
  Active Solvers(2) = 1 2
  Element{1} = String "p:1"
  Element{2} = String "p:1"
End
```

Equation 2

```
Active Solvers(2) = 1 2
Element{1} = String "p:3"
Element{2} = String "p:3"
```

End

It is worth noting, however, that all available solvers of Elmer may not support p -approximation without modifying the code. Usually this only consists of making local stiffness matrix and force vector larger. Nevertheless several solvers already support the approximation with high-order polynomials.

The actual number of degrees of freedom associated with the element edges, faces or interiors depends on the element type and the local polynomial degree p . Each degree of freedom is naturally associated with some basis function. The following table gives the number of degrees of freedom for elements used in Elmer:

Element	Nodes	Edges	Faces	Bubbles
Line	2	-	-	$p - 1$
Quadrilateral	4	$4(p - 1)$	-	$\frac{(p-2)(p-3)}{2}$
Triangle	3	$3(p - 1)$	-	$\frac{(p-1)(p-2)}{2}$
Brick	8	$12(p - 1)$	$3(p - 2)(p - 3)$	$\frac{(p-3)(p-4)(p-5)}{6}$
Tetrahedron	4	$6(p - 1)$	$2(p - 1)(p - 2)$	$\frac{(p-1)(p-2)(p-3)}{6}$
Wedge	6	$9(p - 1)$	-	$\frac{(p-2)(p-3)(p-4)}{6}$
(quad. face)	-	-	$\frac{3(p-2)(p-3)}{2}$	-
(triang. face)	-	-	$(p - 1)(p - 2)$	-
Pyramid	5	$8(p - 1)$	-	$\frac{(p-1)(p-2)(p-3)}{6}$
(quad. face)	-	-	$\frac{(p-2)(p-3)}{2}$	-
(triang. face)	-	-	$2(p - 1)(p - 2)$	-

A p -element passed to the Elmer Gaussian point generator `GaussPoints` contained in the module `Integration` returns enough integration points to integrate the worst-case product of two element basis functions. Here the worst case is the integration for two basis functions of degree $p_m = \max\{p_e, p_f, p_b\}$. As the Gaussian quadrature is accurate up to the degree $p = 2n - 1$, where n is the number of points used, the number of points for each element is calculated from

$$n = \frac{2p_m + 1}{2} \quad (\text{E.4})$$

and rounded up to the nearest integer. To get the final number of points for multiple integrals, n is raised to the power of element dimension. If integral includes a non-constant factor, i.e. $\int_K \alpha \phi_i \phi_j dK$ where α is a function of degree k , the default quadrature is not accurate and the number of integration points needs to be set manually. Now the minimum number of gaussian points to integrate accurately becomes

$$n = \frac{\min\{2p_m + k, 3p_m\} + 1}{2} \quad (\text{E.5})$$

which may again be rounded up to the nearest integer and raised to the power of element dimension to get the actual number of integration points.

E.1.1 Boundary conditions

From the point of view of programming, nonhomogeneous boundary conditions of Dirichlet type, e.g. $u = g$ on a boundary, are more difficult to handle for higher order elements. Even though the nodal values can be set easily, finding coefficients of additional basis functions such that the data is represented (approximately) as a linear combination of the basis functions over the element boundary is not as straightforward. The subroutine `DefaultDirichletBCs` solves the unknown coefficients of higher-order basis functions by seeking the best approximation of the data over the boundary. Thus, from the viewpoint of the user, Dirichlet boundary conditions need no extra actions as compared to the use of the conventional elements. Neumann

and Robin boundary conditions can also be treated in the same way as done in the case of the standard Lagrange elements.

E.1.2 Some practical aspects

Generally using p -elements yields a better approximation than using normal linear elements. In fact, convergence for p -elements is exponential when there are no singularities inside or on the boundary of the solution domain. Typical singular points in the solution are points where boundary conditions or material parameters change abruptly or vertex-type singularities (such as the inner corner of a L-shaped beam or a crack tip). When there are singular points inside the domain, convergence is algebraic. If the singular point is a nodal point, convergence rate is twice that of h -method, otherwise it is equal to the h -method.

Therefore, if possible, the mesh should be designed such that near nodal singularities small low-degree ($p = 1$) elements were used. In other parts of the solution domain, where the solution is smoother, large elements with high polynomial degree are advised. As Elmer is not hp -adaptive so that the element degree is not modified by the solver, mesh design issues must be taken into account for computational efficiency.

It is well known that it is possible to reduce the size of the linear system by the elimination of bubble functions. This procedure is often called the static condensation. In Elmer the static condensation for the local stiffness matrix may be done (and is advised to be done) when the bubble part need not be retrieved afterwards (for example time stepping algorithms need to know the complete solution). Whether the bubble degrees of freedom are assembled into the global system is controlled by the keyword `Bubbles in Global System`.

E.2 ElmerSolver services for higher-order elements

This section describes some of the services related to p -elements, which are included in different parts of the Solver.

E.2.1 Properties of p -elements

There are several utilities for determining p -element properties. First of all, it is possible to check if some element is a p -element by checking the `isPElement` flag of the element. If the flag is set to be true, element is a p -element. Functions

```
isPTriangle( Element )
isPTetra( Element )
isPPyramid( Element )
isPWedge( Element )
```

check if the given element is a p -type triangle, tetrahedron, pyramid or wedge. They are implemented because the reference p -element for triangles, tetrahedra, pyramids and wedges are different from those defined for the Lagrange elements. The routines

```
getEdgeP( Element, Mesh )
getFaceP( Element, Mesh )
```

return the maximum polynomial degrees associated with the element edges or faces when `Element` and finite element mesh `Mesh` are given.

E.2.2 Fields related to p -elements

In the module `Types`, the type `Element_t` has the following p -element related fields

```
INTEGER :: TetraType
LOGICAL :: isPElement
LOGICAL :: isEdge
```

```
INTEGER :: localNumber
INTEGER :: GaussPoints
```

`Tetratype` defines the type of the tetrahedral p -element: for non-tetrahedral elements `Tetratype=0`, for tetrahedral elements `Tetratype={1,2}`.

`isPElement` defines if an element is of higher order: `isPElement=.TRUE.` for p -elements, otherwise `.FALSE.`.

`isEdge` defines if the element is an edge for some higher entity, i.e. an edge or a face of a 2d or 3d element. If `isEdge=.TRUE.`, the element is an edge, `.FALSE.` otherwise.

For boundary elements `localNumber` defines which local edge or face number the boundary element has with respect to its parent element.

`GaussPoints` defines the number n of the Gauss points for the element. The value is calculated from $n = (\frac{2p_m+1}{2})^d$, where d is the element dimension and p_m the maximum polynomial degree. n is rounded up to the nearest integer. The value `GaussPoints` corresponds to the requirement that a polynomial of degree $2p_m$ is integrated accurately.

When modifying a solver to support higher-order elements, the maximum number of degrees of freedom associated with an element may be obtained from the mesh variable `MaxElementDOFs`. This variable is set by the read-in process of the mesh.

E.2.3 Higher-order basis and element mappings

Basis functions for higher-order elements are defined in the module `PElementBase`. The module contains also definitions for the functions ϕ_k and φ_k and Legendre polynomials. These definitions have been generated into explicit form with the symbolic program *Maple* [1] up to $p_{\max} = 20$. This means that no recursion is needed for the generation of the values of Legendre polynomials or other lower-level components based on them, if $p < p_{\max}$.

Generally higher-order basis functions take as their arguments the point in which to calculate the function value and indexing $i, m(i, j)$ or $m(i, j, k)$ depending on the function type. All edge functions take in addition to these parameters a special optional flag, namely `invertEdge`, which defines if the direction of the edge basis function needs to be inverted. In Elmer the parametrization of edges depends on the global node numbering. To describe this, let A and B be the global node numbers of an edge. The varying parameter of edge function then varies between $[-1, 1]$ as $A \rightarrow B$. Inversion is then used for enforcing the global continuity of edge basis functions which are not properly aligned. The edge rule is presented in figure E.2.3.




Figure E.1: The global direction of edge for global node indexes $A < B$

Most of the face functions take an optional orientation argument in order to specify which face functions are formed. This is formed according to the global numbers of face nodes. There are rules for triangular and square faces. Let A, B, C be the global nodes of a triangular face. Globally the face is aligned so that $A < B < C$. For square faces let $A = \min\{v_i\}$ where v_i are the global node indices of square face and B, C correspond to the nodes next to the node A on the face. The square face is aligned by the rule $A < B < C$ for these nodes. These rules are presented in figure E.2.3.

The tetrahedral element uses an exception to the above interface rules, as the edge and face functions of tetrahedral elements take the type of the tetrahedral element as their optional argument. This is due to fact that it is possible to reduce any tetrahedral element to one of the two reference tetrahedral elements for which all edges and faces are defined so that their local orientation matches global orientation. This means, that for tetrahedral elements, global continuity does not need to be enforced, if proper reduction to one of the two reference elements has been made.

Mappings from element nodal numbers to describe different p -element edges or faces are defined in the module `PElementMaps`. These mappings generally define which nodes of the element belong to a certain local edge or face of the element. Mappings to elements edges, faces and from faces to local




Figure E.2: The global direction of triangular and quadrilateral faces for global node indexes $A < B < C$; A has the lowest index among the indexes of the face.

edge numbers may be obtained from the routines `GetElementEdgeMap`, `GetElementFaceMap` and `GetElementFaceEdgeMap`. The mappings may also be accessed via methods `get T_e P_e Map`, where T_e is the element name and $P_e = \{\text{Edge}, \text{Face}\}$ is the part of the element for which the map is requested. The routine `getElementBoundaryMap` returns mappings for element boundaries depending on the element type.

For example, to get the global nodes for the brick face number 4, one would use the following `Fortran90` code

```
map(1:4) = getBrickFaceMap(4)
nodes(1:4) = Element % NodeIndexes(map)
```

E.3 Basis functions

Let $\lambda_1, \lambda_2, \lambda_3 \in \{\pm\xi, \pm\eta, \pm\zeta\}$ and additionally $\bigcap_i \lambda_i = \phi$.

E.3.1 Line




Figure E.3: Line element

Nodal basis

$$\begin{aligned} L_1 &= \frac{1 - \xi}{2} \\ L_2 &= \frac{1 + \xi}{2} \end{aligned}$$

Bubble basis

$$L_i^{(0)} = \phi_i(\xi), \quad i = 2, \dots, p$$

E.3.2 Quadrilateral




Figure E.4: Quadrilateral element

Nodal basis

$$\begin{aligned} N_1 &= \frac{1}{4}(1 - \xi)(1 - \eta) \\ N_2 &= \frac{1}{4}(1 + \xi)(1 - \eta) \\ N_3 &= \frac{1}{4}(1 + \xi)(1 + \eta) \\ N_4 &= \frac{1}{4}(1 - \xi)(1 + \eta) \end{aligned}$$

Edge basis

$$\begin{aligned} N_i^{(1,2)} &= \frac{1}{2}(1 - \eta)\phi_i(\xi), \quad i = 2, \dots, p \\ N_i^{(2,3)} &= \frac{1}{2}(1 + \xi)\phi_i(\eta), \quad i = 2, \dots, p \\ N_i^{(4,3)} &= \frac{1}{2}(1 + \eta)\phi_i(\xi), \quad i = 2, \dots, p \\ N_i^{(1,4)} &= \frac{1}{2}(1 - \xi)\phi_i(\eta), \quad i = 2, \dots, p \end{aligned}$$

Bubble basis

$$N_{m(i,j)}^{(0)} = \phi_i(\xi)\phi_j(\eta)$$

where $i, j \geq 2$, $i + j = 4, \dots, p$




Figure E.5: Triangle element

E.3.3 Triangle

Nodal basis

$$\begin{aligned} L_1 &= \frac{1}{2} \left(1 - \xi - \frac{1}{\sqrt{3}} \eta \right) \\ L_2 &= \frac{1}{2} \left(1 + \xi - \frac{1}{\sqrt{3}} \eta \right) \\ L_3 &= \frac{\eta}{\sqrt{3}} \end{aligned}$$

Edge basis

$$\begin{aligned} N_i^{(1,2)} &= L_1 L_2 \varphi_i(L_2 - L_1), \quad i = 2, \dots, p \\ N_i^{(2,3)} &= L_2 L_3 \varphi_i(L_3 - L_2), \quad i = 2, \dots, p \\ N_i^{(3,1)} &= L_3 L_1 \varphi_i(L_1 - L_3), \quad i = 2, \dots, p \end{aligned}$$

Bubble basis

$$N_{m(j,n)}^{(0)} = L_1 L_2 L_3 P_1(L_2 - L_1)^j P_1(2L_3 - 1)^n$$

where $j, n = 0, \dots, i-3, j+n = i-3, i = 3, \dots, p$




Figure E.6: Brick element

E.3.4 Brick

Nodal basis

$$\begin{aligned}
 N_1 &= \frac{1}{8}(1 - \xi)(1 - \eta)(1 - \zeta) \\
 N_2 &= \frac{1}{8}(1 + \xi)(1 - \eta)(1 - \zeta) \\
 N_3 &= \frac{1}{8}(1 + \xi)(1 + \eta)(1 - \zeta) \\
 N_4 &= \frac{1}{8}(1 - \xi)(1 + \eta)(1 - \zeta) \\
 N_5 &= \frac{1}{8}(1 - \xi)(1 - \eta)(1 + \zeta) \\
 N_6 &= \frac{1}{8}(1 + \xi)(1 - \eta)(1 + \zeta) \\
 N_7 &= \frac{1}{8}(1 + \xi)(1 + \eta)(1 + \zeta) \\
 N_8 &= \frac{1}{8}(1 - \xi)(1 + \eta)(1 + \zeta)
 \end{aligned}$$

Edge basis

$$\begin{aligned}
N_{i-1}^{1,2} &= \frac{1}{4}\phi_i(\xi)(1-\eta)(1-\zeta) \\
N_{i-1}^{2,3} &= \frac{1}{4}\phi_i(\eta)(1+\xi)(1-\zeta) \\
N_{i-1}^{4,3} &= \frac{1}{4}\phi_i(\xi)(1+\eta)(1-\zeta) \\
N_{i-1}^{1,4} &= \frac{1}{4}\phi_i(\eta)(1-\xi)(1-\zeta) \\
N_{i-1}^{1,5} &= \frac{1}{4}\phi_i(\zeta)(1-\xi)(1-\eta) \\
N_{i-1}^{2,6} &= \frac{1}{4}\phi_i(\zeta)(1+\xi)(1-\eta) \\
N_{i-1}^{3,7} &= \frac{1}{4}\phi_i(\zeta)(1+\xi)(1+\eta) \\
N_{i-1}^{4,8} &= \frac{1}{4}\phi_i(\zeta)(1-\xi)(1+\eta) \\
N_{i-1}^{5,6} &= \frac{1}{4}\phi_i(\xi)(1-\eta)(1+\zeta) \\
N_{i-1}^{6,7} &= \frac{1}{4}\phi_i(\eta)(1+\xi)(1+\zeta) \\
N_{i-1}^{8,7} &= \frac{1}{4}\phi_i(\xi)(1+\eta)(1+\zeta) \\
N_{i-1}^{5,8} &= \frac{1}{4}\phi_i(\eta)(1-\xi)(1+\zeta)
\end{aligned}$$

Face basis

$$\begin{aligned}
N_{m(i,j)}^{(1,2,5,6)} &= \frac{1}{2}(1-\eta)\phi_i(\xi)\phi_j(\zeta) \\
N_{m(i,j)}^{(1,2,4,3)} &= \frac{1}{2}(1-\zeta)\phi_i(\xi)\phi_j(\eta) \\
N_{m(i,j)}^{(1,4,5,8)} &= \frac{1}{2}(1-\xi)\phi_i(\eta)\phi_j(\zeta) \\
N_{m(i,j)}^{(4,3,8,7)} &= \frac{1}{2}(1+\eta)\phi_i(\xi)\phi_j(\zeta) \\
N_{m(i,j)}^{(5,6,8,7)} &= \frac{1}{2}(1+\zeta)\phi_i(\xi)\phi_j(\eta) \\
N_{m(i,j)}^{(2,3,6,7)} &= \frac{1}{2}(1+\xi)\phi_i(\eta)\phi_j(\zeta)
\end{aligned}$$

where $i, j = 2, 3, \dots, p-2$, $i+j = 4, 5, \dots, p$

Bubble basis

$$N_{m(i,j,k)}^{(0)} = \phi_i(\xi)\phi_j(\eta)\phi_k(\zeta)$$

where $i, j, k = 2, 3, \dots, p-4$, $i+j+k = 6, 7, \dots, p$




Figure E.7: Tetrahedral elements of types 1 and 2

E.3.5 Tetrahedron

Nodal basis

$$\begin{aligned}
 L_1 &= \frac{1}{2} \left(1 - \xi - \frac{1}{\sqrt{3}} \eta - \frac{1}{\sqrt{6}} \zeta \right) \\
 L_2 &= \frac{1}{2} \left(1 + \xi - \frac{1}{\sqrt{3}} \eta - \frac{1}{\sqrt{6}} \zeta \right) \\
 L_3 &= \frac{\sqrt{3}}{3} \left(\eta - \frac{1}{\sqrt{8}} \zeta \right) \\
 L_4 &= \sqrt{\frac{3}{8}} \zeta
 \end{aligned}$$

Edge basis

Type 1

$$\begin{aligned}
 N_{i-1}^{(1,2)} &= L_1 L_2 \varphi_i (L_2 - L_1), \quad i = 2, \dots, p \\
 N_{i-1}^{(1,3)} &= L_1 L_3 \varphi_i (L_3 - L_1), \quad i = 2, \dots, p \\
 N_{i-1}^{(1,4)} &= L_1 L_4 \varphi_i (L_4 - L_1), \quad i = 2, \dots, p \\
 N_{i-1}^{(2,3)} &= L_2 L_3 \varphi_i (L_3 - L_2), \quad i = 2, \dots, p \\
 N_{i-1}^{(2,4)} &= L_2 L_4 \varphi_i (L_4 - L_2), \quad i = 2, \dots, p \\
 N_{i-1}^{(3,4)} &= L_3 L_4 \varphi_i (L_4 - L_3), \quad i = 2, \dots, p
 \end{aligned}$$

Type 2

$$N_{i-1}^{(3,2)} = L_3 L_2 \varphi_i (L_2 - L_3), \quad i = 2, \dots, p$$

Edges (1, 2), (1, 3), (1, 4), (2, 4) and (3, 4) according to type 1.

Face basis

Type 1

$$\begin{aligned}
N_{m(i,j)}^{(1,2,3)} &= L_1 L_2 L_3 P_i (L_2 - L_1) P_j (2L_3 - 1) \\
N_{m(i,j)}^{(1,2,4)} &= L_1 L_2 L_4 P_i (L_2 - L_1) P_j (2L_4 - 1) \\
N_{m(i,j)}^{(1,3,4)} &= L_1 L_4 L_3 P_i (L_3 - L_1) P_j (2L_4 - 1) \\
N_{m(i,j)}^{(2,3,4)} &= L_2 L_3 L_4 P_i (L_3 - L_2) P_j (2L_4 - 1)
\end{aligned}$$

Type 2

$$\begin{aligned}
N_{m(i,j)}^{(1,3,2)} &= L_1 L_3 L_2 P_i (L_3 - L_1) P_j (2L_2 - 1) \\
N_{m(i,j)}^{(3,2,4)} &= L_3 L_2 L_4 P_i (L_2 - L_3) P_j (2L_4 - 1)
\end{aligned}$$

where $i, j = 0, 1, 2, \dots, p-3$, $i+j = 0, 1, \dots, p-3$. Faces (1, 2, 4) and (1, 3, 4) defined according to type 1.

Bubble basis

$$N_{m(i,j,k)}^{(0)} = L_1 L_2 L_3 L_4 P_i (L_2 - L_1) P_j (2L_3 - 1) P_k (2L_4 - 1)$$

where $i, j, k = 0, 1, \dots, p-4$, $i+j+k = 0, 1, \dots, p-4$


E.3.6 Pyramid

Figure E.8: Pyramidal element

Nodal basis

$$\begin{aligned}
T_0(c, t) &= \frac{(1 - \frac{t}{\sqrt{2}}) - c}{2(1 - \frac{t}{\sqrt{2}})} \\
T_1(c, t) &= \frac{(1 - \frac{t}{\sqrt{2}}) + c}{2(1 - \frac{t}{\sqrt{2}})}
\end{aligned}$$

$$\begin{aligned}
P_1 &= T_0(\xi, \zeta)T_0(\eta, \zeta)(1 - \frac{\zeta}{\sqrt{2}}) \\
P_2 &= T_1(\xi, \zeta)T_0(\eta, \zeta)(1 - \frac{\zeta}{\sqrt{2}}) \\
P_3 &= T_1(\xi, \zeta)T_1(\eta, \zeta)(1 - \frac{\zeta}{\sqrt{2}}) \\
P_4 &= T_0(\xi, \zeta)T_1(\eta, \zeta)(1 - \frac{\zeta}{\sqrt{2}}) \\
P_5 &= \frac{1}{\sqrt{2}}\zeta
\end{aligned}$$

Edge basis

$$\begin{aligned}
P_{i-1}^{(1,2)} &= P_1(\xi, \eta, \zeta)P_2(\xi, \eta, \zeta)\varphi_i(\xi) \\
P_{i-1}^{(2,3)} &= P_2(\xi, \eta, \zeta)P_3(\xi, \eta, \zeta)\varphi_i(\eta) \\
P_{i-1}^{(4,3)} &= P_4(\xi, \eta, \zeta)P_3(\xi, \eta, \zeta)\varphi_i(\xi) \\
P_{i-1}^{(1,4)} &= P_1(\xi, \eta, \zeta)P_4(\xi, \eta, \zeta)\varphi_i(\eta) \\
P_{i-1}^{(1,5)} &= P_1(\xi, \eta, \zeta)P_5(\xi, \eta, \zeta)\varphi_i(\frac{\xi}{2} + \frac{\eta}{2} + \frac{\zeta}{\sqrt{2}}) \\
P_{i-1}^{(2,5)} &= P_2(\xi, \eta, \zeta)P_5(\xi, \eta, \zeta)\varphi_i(-\frac{\xi}{2} + \frac{\eta}{2} + \frac{\zeta}{\sqrt{2}}) \\
P_{i-1}^{(3,5)} &= P_3(\xi, \eta, \zeta)P_5(\xi, \eta, \zeta)\varphi_i(-\frac{\xi}{2} - \frac{\eta}{2} + \frac{\zeta}{\sqrt{2}}) \\
P_{i-1}^{(4,5)} &= P_4(\xi, \eta, \zeta)P_5(\xi, \eta, \zeta)\varphi_i(\frac{\xi}{2} - \frac{\eta}{2} + \frac{\zeta}{\sqrt{2}})
\end{aligned}$$

Face basis**Square face**

$$P_{m(i,j)}^{(1,2,3,4)} = P_1(\xi, \eta, \zeta)P_3(\xi, \eta, \zeta)\varphi_i(\xi)\varphi_j(\eta)$$

where $i, j = 2, \dots, p-2, i+j = 4, \dots, p$.

Triangular faces

$$\begin{aligned}
P_{m(i,j)}^{(1,2,5)} &= P_1(\xi, \eta, \zeta)P_2(\xi, \eta, \zeta)P_5(\xi, \eta, \zeta)P_i(P_2(\xi, \eta, \zeta) - P_1(\xi, \eta, \zeta))P_j(2P_5(\xi, \eta, \zeta) - 1) \\
P_{m(i,j)}^{(2,3,5)} &= P_2(\xi, \eta, \zeta)P_3(\xi, \eta, \zeta)P_5(\xi, \eta, \zeta)P_i(P_3(\xi, \eta, \zeta) - P_2(\xi, \eta, \zeta))P_j(2P_5(\xi, \eta, \zeta) - 1) \\
P_{m(i,j)}^{(3,4,5)} &= P_3(\xi, \eta, \zeta)P_4(\xi, \eta, \zeta)P_5(\xi, \eta, \zeta)P_i(P_4(\xi, \eta, \zeta) - P_3(\xi, \eta, \zeta))P_j(2P_5(\xi, \eta, \zeta) - 1) \\
P_{m(i,j)}^{(4,1,5)} &= P_4(\xi, \eta, \zeta)P_1(\xi, \eta, \zeta)P_5(\xi, \eta, \zeta)P_i(P_1(\xi, \eta, \zeta) - P_4(\xi, \eta, \zeta))P_j(2P_5(\xi, \eta, \zeta) - 1)
\end{aligned}$$

where $i, j = 0, \dots, p-3, i+j = 0, \dots, p-3$ and P_i, P_j Legendre polynomials.

Bubble basis

$$P_{m(i,j,k)}^{(0)} = P_1(\xi, \eta, \zeta) P_3(\xi, \eta, \zeta) P_5(\xi, \eta, \zeta) P_i\left(\frac{\xi}{1 - \frac{\zeta}{\sqrt{2}}}\right) P_j\left(\frac{\eta}{1 - \frac{\zeta}{\sqrt{2}}}\right) P_k\left(\frac{\zeta}{\sqrt{2}}\right)$$

where $i, j, k = 0, \dots, p-4$, $i+j+k = 0, \dots, p-4$ and P_i, P_j, P_k Legendre polynomials


E.3.7 Wedge

Figure E.9: Wedge element

Nodal basis

$$L_1 = \frac{1}{2} \left(1 - \xi - \frac{1}{\sqrt{3}} \eta \right)$$

$$L_2 = \frac{1}{2} \left(1 + \xi - \frac{1}{\sqrt{3}} \eta \right)$$

$$L_3 = \frac{\sqrt{3}}{3} \eta$$

$$H_1 = \frac{1}{2} L_1 (1 - \zeta)$$

$$H_2 = \frac{1}{2} L_2 (1 - \zeta)$$

$$H_3 = \frac{1}{2} L_3 (1 - \zeta)$$

$$H_4 = \frac{1}{2} L_1 (1 + \zeta)$$

$$H_5 = \frac{1}{2} L_2 (1 + \zeta)$$

$$H_6 = \frac{1}{2} L_3 (1 + \zeta)$$

Edge basis

$$\begin{aligned}
H_{i-1}^{(1,2)} &= \frac{1}{2}L_1L_2\varphi_i(L_2 - L_1)(1 - \zeta) \\
H_{i-1}^{(2,3)} &= \frac{1}{2}L_2L_3\varphi_i(L_3 - L_2)(1 - \zeta) \\
H_{i-1}^{(3,1)} &= \frac{1}{2}L_3L_1\varphi_i(L_1 - L_3)(1 - \zeta) \\
H_{i-1}^{(4,5)} &= \frac{1}{2}L_4L_5\varphi_i(L_5 - L_4)(1 + \zeta) \\
H_{i-1}^{(5,6)} &= \frac{1}{2}L_5L_6\varphi_i(L_6 - L_5)(1 + \zeta) \\
H_{i-1}^{(6,4)} &= \frac{1}{2}L_6L_4\varphi_i(L_4 - L_6)(1 + \zeta) \\
H_{i-1}^{(1,4)} &= L_1\phi_i(\zeta) \\
H_{i-1}^{(2,5)} &= L_2\phi_i(\zeta) \\
H_{i-1}^{(3,6)} &= L_3\phi_i(\zeta)
\end{aligned}$$

Face basis**Triangular faces**

$$\begin{aligned}
H_{m(i,j)}^{(1,2,3)} &= \frac{1}{2}(1 - \zeta)P_i(L_2 - L_1)P_j(2L_3 - 1)L_1L_2L_3 \\
H_{m(i,j)}^{(4,5,6)} &= \frac{1}{2}(1 + \zeta)P_i(L_2 - L_1)P_j(2L_3 - 1)L_1L_2L_3
\end{aligned}$$

where $i, j = 0, 1, \dots, p-3$, $i+j = 0, 1, \dots, p-3$ and P_i, P_j Legendre polynomials.

Square faces

$$\begin{aligned}
H_{m(i,j)}^{(1,2,5,4)} &= \varphi_i(L_2 - L_1)\phi_j(\zeta)L_1L_2 \\
H_{m(i,j)}^{(2,3,6,5)} &= \varphi_i(L_3 - L_2)\phi_j(\zeta)L_2L_3 \\
H_{m(i,j)}^{(3,1,4,6)} &= \varphi_i(L_1 - L_3)\phi_j(\zeta)L_3L_1
\end{aligned}$$

where $i, j = 2, \dots, p-2$, $i+j = 4, \dots, p$.

Bubble basis

$$H_{m(i,j,k)}^{(0)} = \phi_k(\zeta)L_1L_2L_3P_i(L_2 - L_1)P_j(2L_3 - 1)$$

where $i, j = 0, \dots, p-5$, $k = 2, \dots, p-3$, $i+j+k = 2, \dots, p-3$.

Bibliography

- [1] Maplesoft home page. <http://www.maplesoft.com/>.
- [2] P.R.B. Devloo. On the definition of high order shape functions for finite elements. Available online: <http://www.fec.unicamp.br/~phil/downloads/shape.zip>.
- [3] B. Szabo and I. Babuska. *Finite Element Analysis*. John Wiley & Sons Ltd., 1991.
- [4] P. Šolín et al. *Higher-Order Finite Element Methods*. Chapman & Hall / CRC, 2004.

Appendix F

Face and edge elements

In Elmer the way to define a high-degree finite element discretization is based on the idea that a background mesh for representing the standard lowest-degree continuous finite element expansion is first provided and then additional degrees of freedom are introduced in order to enhance the approximation defined originally over the background mesh. The same idea has also been adapted to create other alternate finite element formulations. For example, it is possible to enhance the approximation on the background mesh by introducing additional degrees of freedom associated with either faces or edges of the elements and omitting then the original nodal degrees of freedom. This leads to a suitable set of unknowns for creating discretizations based on the face or edge element interpolation. If $L_2(\Omega)$ is used to denote the set of square-integrable scalar functions and $\Omega \subset \mathbb{R}^d$, we are thus led to bases for approximating vector fields in finite dimensional versions of the spaces

$$\mathbf{H}(\text{div}, \Omega) = \{ \mathbf{v} \in L_2(\Omega)^d \mid \text{div } \mathbf{v} \in L_2(\Omega) \}$$

or

$$\mathbf{H}(\text{curl}, \Omega) = \{ \mathbf{v} \in L_2(\Omega)^d \mid \text{curl } \mathbf{v} \in L_2(\Omega)^d \}.$$

A physical motivation for using these spaces is that fields with only either normal or tangential continuity on boundaries can then be approximated in a natural manner. They may also be used to approximate various mixed variational formulations. In this appendix we shall detail the construction of available bases for defining face and edge interpolation elements. Currently the selection of edge interpolation elements is more extensive.

F.1 The construction of face element interpolation

Here we focus on the construction of finite element spaces $\mathbf{X}_h(\Omega) \subset \mathbf{H}(\text{div}, \Omega)$. To this end, we use local spaces $\mathbf{X}(K)$ which are defined over elements K in the finite element mesh \mathcal{M}_h associated with the region Ω . Then $\mathbf{H}(\text{div}, \Omega)$ -conforming approximations are obtained as members of the space

$$\mathbf{X}_h = \{ \mathbf{v}_h \in \mathbf{H}(\text{div}, \Omega) \mid \mathbf{v}_h|_K \in \mathbf{X}(K) \forall K \in \mathcal{M}_h \}.$$

We note that the functions belonging to \mathbf{X}_h must have continuous normal components across the interior element faces of the mesh \mathcal{M}_h .

Any real element $K \subset \Omega$ is obtained as an image of a reference element \hat{K} under the element mapping \mathbf{f}_K so that $K = \mathbf{f}_K(\hat{K})$. The element space $\mathbf{X}(K)$ is defined via introducing a local space $\mathbf{X}(\hat{K})$ on the reference element and showing how its members are transformed in connection with the element mapping to get the basis for spanning $\mathbf{X}(K)$. To this end, we use the Piola transformation so that the connection between the two element spaces is given by

$$\mathbf{X}(K) = \{ \mathbf{v}_h \mid \mathbf{v}_h(\mathbf{x}) = \frac{1}{\det \mathbf{F}(\mathbf{f}_K^{-1}(\mathbf{x}))} \mathbf{F}(\mathbf{f}_K^{-1}(\mathbf{x})) (\hat{\mathbf{v}}_h \circ \mathbf{f}_K^{-1})(\mathbf{x}), \text{ with } \hat{\mathbf{v}}_h \in \mathbf{X}(\hat{K}) \}$$

where $\mathbf{F} = \nabla \mathbf{f}_K$ is the gradient of \mathbf{f}_K with respect to $\hat{\mathbf{x}} \in \hat{K}$. It then follows that the basis functions for $\mathbf{X}(K)$ are obtained by applying the Piola transformation to the basis of $\mathbf{X}(\hat{K})$.

Currently the Elmer solver contains the implementation of face element bases for triangular, quadrilateral and tetrahedral elements. In the case of triangular and quadrilateral elements the reference element is chosen to be either an equilateral triangle or a square, while tetrahedral elements use a regular tetrahedron as the reference element. These choices are convenient as using reference elements which have the faces of an equal size simplifies implementation. These reference elements have also been employed in connection with the high-order elements. The element mapping $\mathbf{f}_K : \hat{K} \rightarrow \mathbb{R}^3$ is always defined in terms of the lowest-order Lagrange interpolation basis functions \hat{L}_j associated with the reference element as

$$\mathbf{x} = \mathbf{f}_K(\hat{\mathbf{x}}) = \sum_j \hat{L}_j(\hat{\mathbf{x}}) \mathbf{x}_j,$$

where $\hat{\mathbf{x}}$ is a point on the reference element and the points \mathbf{x}_j give the vertices of the real element K .

F.1.1 The Raviart–Thomas space and its extension to three dimensions

In two dimensions ($d = 2$) the simplest face element is the triangular Raviart–Thomas element of order $k = 0$. The reference element space is then chosen to be

$$\mathbf{X}(\hat{K}) = \mathbf{RT}_0(\hat{K}) = [P_0(\hat{K})]^d \oplus P_0(\hat{K}) \begin{bmatrix} \hat{x}_1 \\ \hat{x}_2 \end{bmatrix} \quad (\text{F.1})$$

where $P_k(\hat{K})$ is the space of polynomials of at most degree k on \hat{K} . The dimension of $\mathbf{X}(\hat{K})$ is thus 3.

To define a basis for $\mathbf{RT}_0(\hat{K})$ each element face $\hat{\mathcal{F}}_{ij}$ connecting the reference element vertices $\hat{\mathbf{x}}_i$ and $\hat{\mathbf{x}}_j$ is associated with one element basis function $\hat{\psi}_{ij}$. We choose

$$\begin{aligned} \hat{\psi}_{12} &= -\frac{1}{2} \begin{bmatrix} 0 \\ 1 \end{bmatrix} + \frac{1}{2\sqrt{3}} \begin{bmatrix} \hat{x}_1 \\ \hat{x}_2 \end{bmatrix}, \\ \hat{\psi}_{23} &= \frac{1}{2\sqrt{3}} \begin{bmatrix} 1 \\ 0 \end{bmatrix} + \frac{1}{2\sqrt{3}} \begin{bmatrix} \hat{x}_1 \\ \hat{x}_2 \end{bmatrix}, \\ \hat{\psi}_{31} &= -\frac{1}{2\sqrt{3}} \begin{bmatrix} 1 \\ 0 \end{bmatrix} + \frac{1}{2\sqrt{3}} \begin{bmatrix} \hat{x}_1 \\ \hat{x}_2 \end{bmatrix}. \end{aligned} \quad (\text{F.2})$$

A finite element expansion

$$\hat{\mathbf{v}}_h(\hat{\mathbf{x}}) = \hat{d}_{12} \hat{\psi}_{12}(\hat{\mathbf{x}}) + \hat{d}_{23} \hat{\psi}_{23}(\hat{\mathbf{x}}) + \hat{d}_{31} \hat{\psi}_{31}(\hat{\mathbf{x}})$$

is referred to as the Raviart-Thomas interpolation function. Here the degrees of freedom \hat{d}_{ij} describe the flux across the element faces as

$$\hat{d}_{ij} = \int_{\hat{\mathcal{F}}_{ij}} \hat{\mathbf{v}}_h(\hat{\mathbf{x}}) \cdot \hat{\mathbf{n}}(\hat{\mathbf{x}}) d\hat{s},$$

where $\hat{\mathbf{n}}$ denotes the outward unit normal associated with the element face.

The tetrahedral element of the Nédélec face element family of the first kind generalizes the use of the space (F.1) to three dimensions ($d = 3$). If we introduce the Whitney form $\hat{\phi}_{ijk}$ by

$$\hat{\phi}_{ijk} = \hat{\lambda}_i \nabla \hat{\lambda}_j \times \nabla \hat{\lambda}_k - \hat{\lambda}_j \nabla \hat{\lambda}_i \times \nabla \hat{\lambda}_k + \hat{\lambda}_k \nabla \hat{\lambda}_i \times \nabla \hat{\lambda}_j,$$

the reference element basis functions $\hat{\psi}_{ijk}$ associated with the element faces $\hat{\mathcal{F}}_{ijk}$ may be defined as

$$\hat{\psi}_{213} = -2\hat{\phi}_{123}, \quad \hat{\psi}_{124} = 2\hat{\phi}_{124}, \quad \hat{\psi}_{234} = 2\hat{\phi}_{234}, \quad \hat{\psi}_{314} = -2\hat{\phi}_{134},$$

so that

$$\begin{aligned} \hat{\psi}_{213} &= \begin{bmatrix} 0 \\ -\sqrt{6}/12 \\ -1/\sqrt{3} \end{bmatrix} + \frac{\sqrt{2}}{4} \begin{bmatrix} \hat{x}_1 \\ \hat{x}_2 \\ \hat{x}_3 \end{bmatrix}, \quad \hat{\psi}_{124} = \begin{bmatrix} 0 \\ -\sqrt{6}/4 \\ 0 \end{bmatrix} + \frac{\sqrt{2}}{4} \begin{bmatrix} \hat{x}_1 \\ \hat{x}_2 \\ \hat{x}_3 \end{bmatrix}, \\ \hat{\psi}_{234} &= \begin{bmatrix} \sqrt{2}/4 \\ 0 \\ 0 \end{bmatrix} + \frac{\sqrt{2}}{4} \begin{bmatrix} \hat{x}_1 \\ \hat{x}_2 \\ \hat{x}_3 \end{bmatrix}, \quad \hat{\psi}_{314} = \begin{bmatrix} -\sqrt{2}/4 \\ 0 \\ 0 \end{bmatrix} + \frac{\sqrt{2}}{4} \begin{bmatrix} \hat{x}_1 \\ \hat{x}_2 \\ \hat{x}_3 \end{bmatrix}. \end{aligned} \quad (\text{F.3})$$

The degrees of freedom for the corresponding finite element expansion $\hat{v}_h(\hat{\mathbf{x}})$ expressed in terms of these basis functions are then characterized by

$$\hat{d}_{ijk} = \int_{\hat{\mathcal{F}}_{ijk}} \hat{v}_h(\hat{\mathbf{x}}) \cdot \hat{\mathbf{n}}(\hat{\mathbf{x}}) d\hat{S}.$$

It can be verified that under the Piola transformation the basis functions of the reference element space $\mathbf{X}(\hat{K}) = \mathbf{RT}_0(\hat{K})$ transform to a basis for $\mathbf{X}(K) = \mathbf{RT}_0(K)$. In addition, the degrees of freedom associated with the resulting basis functions ψ_{ij} or ψ_{ijk} are then characterized by

$$d_{ij} = \int_{\mathcal{F}_{ij}} \mathbf{v}_h(\mathbf{x}) \cdot \mathbf{n}(\mathbf{x}) dS \quad \text{or} \quad d_{ijk} = \int_{\mathcal{F}_{ijk}} \mathbf{v}_h(\mathbf{x}) \cdot \mathbf{n}(\mathbf{x}) dS.$$

In order to define a unique set of degrees of freedom for assembling global finite element equations, each element face in the mesh \mathcal{M}_h is associated with a uniquely defined unit normal \mathbf{m} , so that either $\mathbf{n}(\mathbf{x}) \cdot \mathbf{m}(\mathbf{x}) = 1$ or $\mathbf{n}(\mathbf{x}) \cdot \mathbf{m}(\mathbf{x}) = -1$ for all \mathbf{x} on the same element face. The degrees of freedom for assembling the global equations from elementwise contributions are then selected to be

$$D_{ij} = \int_{\mathcal{F}_{ij}} \mathbf{v}_h(\mathbf{x}) \cdot \mathbf{m}(\mathbf{x}) dS \quad \text{or} \quad D_{ijk} = \int_{\mathcal{F}_{ijk}} \mathbf{v}_h(\mathbf{x}) \cdot \mathbf{m}(\mathbf{x}) dS.$$

The restriction of the global finite element expansion \mathbf{v}_h on the element can then be written in terms of the local basis functions and the global degrees of freedom as

$$\mathbf{v}_h|_K = \sum d_{ij} \psi_{ij} = \sum D_{ij} (\mathbf{n} \cdot \mathbf{m})_{\mathbf{x} \in \mathcal{F}_{ij}} \psi_{ij}$$

or

$$\mathbf{v}_h|_K = \sum d_{ijk} \psi_{ijk} = \sum D_{ijk} (\mathbf{n} \cdot \mathbf{m})_{\mathbf{x} \in \mathcal{F}_{ijk}} \psi_{ijk}.$$

In practice this is implemented by applying sign reversions to the local basis functions when $\mathbf{n} \cdot \mathbf{m} = -1$. It should be noted that the normal continuity of the resulting face element expansion is then in-built.

F.1.2 The Brezzi-Douglas-Marini space and its extension to three dimensions

In the following we consider choosing the reference element space as $\mathbf{X}(\hat{K}) = [P_1(\hat{K})]^d$. In two dimensions this choice leads to the construction of the Brezzi-Douglas-Marini space of order 1 (BDM_1), while in three dimensions it brings us to the definition of the Nédélec face elements of the second kind.

It seems that ways to obtain computational bases for face element discretizations are far less standardized in comparison with implementing the basic element types. In Elmer we have followed the idea employed in [1] to obtain a basis for the reference element space. Each basis function may then be associated with a certain numerical integration point located on one of the element faces. The normal component of the basis function is normalized to yield the unity at the integration point to which the basis function is associated, while the normal component of this basis function is required to vanish at other integration points located on the element faces. It follows from this strategy that convenient integral descriptions of the degrees of freedom are also obtained.

For triangular elements the dimension of $\mathbf{X}(\hat{K})$ is 6 and hence two basis functions $\hat{\psi}_{ij}^k$, $k = 1, 2$, are associated with each face $\hat{\mathcal{F}}_{ij}$. A possible choice would be to use $\hat{\lambda}_i \mathbf{rot} \hat{\lambda}_j$ and $\hat{\lambda}_j \mathbf{rot} \hat{\lambda}_i$ as these basis functions, but the basis functions implemented into Elmer are chosen to be their linear combinations such that the normal components of the resulting alternate basis functions satisfy the version of the Kronecker delta property explained above. We note that if the element face is parametrized as $\hat{\mathcal{F}}_{ij} = \{ \hat{\mathbf{x}} = \hat{\mathbf{s}}(\xi) \mid -1 \leq \xi \leq 1 \}$, the two basis functions are associated with the Gaussian points $\xi = -1/\sqrt{3}$ and $\xi = 1/\sqrt{3}$. The degrees of freedom associated with the face element interpolation function $\hat{v}_h(\hat{\mathbf{x}})$ can then be characterized by

$$\hat{d}_{ij}^k = \int_{\hat{\mathcal{F}}_{ij}} \hat{v}_h(\hat{\mathbf{x}}) \cdot \hat{\mathbf{n}}(\hat{\mathbf{x}}) \hat{\phi}_k(\hat{\mathbf{x}}) d\hat{s},$$

where $\hat{\phi}_k(\hat{\mathbf{x}})$ are the linear Lagrange interpolation functions which have the Gaussian points of the element face as the interpolation points.

The construction of the basis for the reference element space in three dimensions is analogous to the two-dimensional case. The dimension of $\mathbf{X}(\hat{K})$ is then 12 and each face $\hat{\mathcal{F}}_{ijk}$ of the reference tetrahedron is associated with three basis functions. The basis functions $\hat{\psi}_{ijk}^l$, $l = 1, 2, 3$, associated with each face $\hat{\mathcal{F}}_{ijk}$ are determined as linear combinations of the alternate basis functions (see, for example, [4]) $\hat{\lambda}_k \nabla \hat{\lambda}_i \times \nabla \hat{\lambda}_j$, $\hat{\lambda}_j \nabla \hat{\lambda}_i \times \nabla \hat{\lambda}_k$ and $\hat{\lambda}_i \nabla \hat{\lambda}_j \times \nabla \hat{\lambda}_k$. Parametrizing the element face in the coordinates ξ and η as shown in Figure E.5, the three basis functions for the element face are associated with the numerical integration points $(-1/2, \sqrt{3}/6)$, $(1/2, \sqrt{3}/6)$, and $(0, 2/\sqrt{3})$. The degrees of freedom for the face element interpolation function $\hat{v}_h(\hat{\mathbf{x}})$ in terms of this basis may again be characterized as

$$d_{ijk}^l = \int_{\hat{\mathcal{F}}_{ijk}} \hat{v}_h(\hat{\mathbf{x}}) \cdot \hat{\mathbf{n}}(\hat{\mathbf{x}}) \hat{\phi}_l(\hat{\mathbf{x}}) d\hat{S},$$

where $\hat{\phi}_l(\hat{\mathbf{x}})$ are the linear Lagrange interpolation functions having the numerical integration points of the element face as the interpolation points.

The Piola transformation is again used to transform the basis functions of the reference element space $\mathbf{X}(\hat{K}) = [P_1(\hat{K})]^d$ to a basis for $\mathbf{X}(K) = [P_1(K)]^d$. Each element face in the mesh \mathcal{M}_h is also associated with a uniquely defined unit normal \mathbf{m} , so that global finite element equations can be assembled using the unique set of degrees of freedom as explained in connection with the Raviart-Thomas element. Additionally an in-built mechanism is also applied to permute the order of basis functions so that the continuity of the normal components is maintained when elementwise contributions are assembled to the global linear system using the standard assembly subroutines of the Elmer solver.

F.1.3 The quadrilateral Arnold-Boffi-Falk element

The quadrilateral Arnold-Boffi-Falk edge element [3] is designed to provide the optimal accuracy even when the physical elements are not obtained as affine images of the reference element. In this case the reference element space is taken to be

$$\mathbf{X}(\hat{K}) = \mathbf{ABF}_0(\hat{K}) = Q_{2,0}(\hat{K}) \times Q_{0,2}(\hat{K}) \quad (\text{F.4})$$

where $Q_{m,n}(\hat{K})$ denotes the space of polynomial functions on \hat{K} of degree at most m in \hat{x}_1 and at most n in \hat{x}_2 . The dimension of this space is 6. Four degrees of freedom are associated with the faces of the element and the remaining two are local.

F.1.4 Using integration over the reference element

As with standard finite elements, it is convenient to convert integrals over elements K to integrals over the reference element \hat{K} in order to perform the numerical integration. Here the effect of using the Piola transformation needs to be taken into account and hence the implementation of elementwise integration is somewhat different as compared with using the standard transformation laws for the basic finite elements.

In the following $\mathbf{u} : K \rightarrow \mathbb{R}^d$ and $\mathbf{v} : K \rightarrow \mathbb{R}^d$ are the Piola transformations of $\hat{\mathbf{u}} : \hat{K} \rightarrow \mathbb{R}^d$ and $\hat{\mathbf{v}} : \hat{K} \rightarrow \mathbb{R}^d$. A scalar field $\phi : K \rightarrow \mathbb{R}$ is obtained from a field $\hat{\phi} : \hat{K} \rightarrow \mathbb{R}$ via $\phi(\mathbf{x}) = (\hat{\phi} \circ \mathbf{f}_K^{-1})(\mathbf{x})$. In addition, \mathbf{grad} and \mathbf{div} are the gradient and divergence operators with respect to \mathbf{x} , while ∇ and \mathbf{Div} are used to denote the gradient and divergence operators with respect to $\hat{\mathbf{x}}$. The following transformation laws

can then be used to convert integrals over elements K to integrals over the reference element \hat{K} :

$$\begin{aligned}
 \int_{\partial K} \mathbf{u}(\mathbf{x}) \cdot \mathbf{n}(\mathbf{x}) \phi(\mathbf{x}) dS &= \int_{\partial \hat{K}} \hat{\mathbf{u}}(\hat{\mathbf{x}}) \cdot \hat{\mathbf{n}}(\hat{\mathbf{x}}) \hat{\phi}(\hat{\mathbf{x}}) d\hat{S}, \\
 \int_K \mathbf{u}(\mathbf{x}) \cdot \mathbf{v}(\mathbf{x}) d\Omega &= \int_{\hat{K}} \mathbf{F} \hat{\mathbf{u}}(\hat{\mathbf{x}}) \cdot \mathbf{F} \hat{\mathbf{v}}(\hat{\mathbf{x}}) \frac{1}{\det \mathbf{F}} d\hat{\Omega}, \\
 \int_K \operatorname{div} \mathbf{u}(\mathbf{x}) \phi(\mathbf{x}) d\Omega &= \int_{\hat{K}} \operatorname{Div} \hat{\mathbf{u}}(\hat{\mathbf{x}}) \hat{\phi}(\hat{\mathbf{x}}) d\hat{\Omega}, \\
 \int_K \mathbf{u}(\mathbf{x}) \cdot \mathbf{grad} \phi(\mathbf{x}) d\Omega &= \int_{\hat{K}} \hat{\mathbf{u}}(\hat{\mathbf{x}}) \cdot \nabla \hat{\phi}(\hat{\mathbf{x}}) d\hat{\Omega}.
 \end{aligned} \tag{F.5}$$

F.1.5 Examples

The classic Poisson equation

$$-\operatorname{div} \mathbf{grad} u = f \quad \text{on } \Omega$$

can be written as a mixed problem such that the problem consists of finding simultaneously u in $L_2(\Omega)$ and the flux vector $\mathbf{q} = \mathbf{grad} u$ in $\mathbf{H}(\operatorname{div}, \Omega)$. The files in the directories

- `../trunk/fem/tests/RaviartThomas2D`
- `../trunk/fem/tests/BDM2D`
- `../trunk/fem/tests/RaviartThomas3D`
- `../trunk/fem/tests/BDM3D`

contained in the source code repository of Elmer exemplify how the approximation of this mixed formulation can be implemented using the face elements which have been considered in this appendix.

F.2 The construction of edge element interpolation

The definition of finite element spaces $\mathbf{X}_h(\Omega) \subset \mathbf{H}(\operatorname{curl}, \Omega)$ is also given by using local spaces $\mathbf{X}(K)$ so that $\mathbf{H}(\operatorname{curl}, \Omega)$ -conforming approximations are obtained as members of the space

$$\mathbf{X}_h = \{ \mathbf{v}_h \in \mathbf{H}(\operatorname{curl}, \Omega) \mid \mathbf{v}_h|_K \in \mathbf{X}(K) \forall K \in \mathcal{M}_h \}.$$

We note that the functions belonging to \mathbf{X}_h must now have continuous tangential components across the interior element faces of the finite element mesh \mathcal{M}_h associated with the region Ω .

The element space $\mathbf{X}(K)$ is again constructed via introducing a local space $\mathbf{X}(\hat{K})$ on the reference element and showing how its members are transformed in connection with the element mapping to obtain the basis for spanning $\mathbf{X}(K)$. To this end, we use a transformation which is similar to the Piola transformation used in the construction of the face element interpolation. Here the connection between the two element spaces is given by

$$\mathbf{X}(K) = \{ \mathbf{v}_h \mid \mathbf{v}_h(\mathbf{x}) = [\mathbf{F}(\mathbf{f}_K^{-1}(\mathbf{x}))]^{-T} \hat{\mathbf{v}}_h(\mathbf{f}_K^{-1}(\mathbf{x})), \text{ with } \hat{\mathbf{v}}_h \in \mathbf{X}(\hat{K}) \}.$$

It then follows that the basis functions for $\mathbf{X}(K)$ are obtained by applying the $\mathbf{H}(\operatorname{curl})$ -version of the Piola transformation to the basis of $\mathbf{X}(\hat{K})$.

F.2.1 The lowest-order triangular and tetrahedral edge elements of Nédélec's first and second family

In the case of triangular and tetrahedral elements the Nédélec edge elements of the first kind are based on choosing the reference element space as

$$\mathbf{X}(\hat{K}) = [P_{k-1}(\hat{K})]^d + \mathbf{S}_k(\hat{K}) \quad (\text{F.6})$$

where

$$\mathbf{S}_k(\hat{K}) = \{ \mathbf{q} \in [\tilde{P}_k]^d \mid \mathbf{q} \cdot \mathbf{x} = 0 \} \quad (\text{F.7})$$

with \tilde{P}_k the set of homogeneous polynomials of degree k , i.e. the span of monomials of degree k . On the other hand, the Nédélec edge elements of the second kind are obtained by selecting

$$\mathbf{X}(\hat{K}) = [P_k(\hat{K})]^d. \quad (\text{F.8})$$

In the case of Nédélec's first family of degree $k = 1$ each reference element edge $\hat{\mathcal{E}}_{ij}$ connecting the reference element vertices $\hat{\mathbf{x}}_i$ and $\hat{\mathbf{x}}_j$ is associated with one basis function $\hat{\psi}_{ij}$. Given a finite element expansion $\hat{\mathbf{v}}_h(\hat{\mathbf{x}})$ in terms of these basis functions, the associated degrees of freedom \hat{d}_{ij} satisfy

$$\hat{d}_{ij} = \int_{\hat{\mathcal{E}}_{ij}} \hat{\mathbf{v}}_h(\hat{\mathbf{x}}) \cdot \hat{\mathbf{t}}(\hat{\mathbf{x}}) d\hat{s},$$

where $\hat{\mathbf{t}}$ denotes the unit tangent vector associated with the element edge.

The $\mathbf{H}(\text{curl}, \Omega)$ -version of the Piola transformation is applied to transform the basis functions of the reference element space $\mathbf{X}(\hat{K})$ to a basis for $\mathbf{X}(K)$. The degrees of freedom associated with the resulting basis functions ψ_{ij} are then characterized by

$$d_{ij} = \int_{\mathcal{E}_{ij}} \mathbf{v}_h(\mathbf{x}) \cdot \mathbf{t}(\mathbf{x}) ds,$$

with \mathbf{t} the unit tangent vector that is oriented similarly with $\hat{\mathbf{t}}$. In order to define a unique set of degrees of freedom for assembling global finite element equations, each element edge in the mesh \mathcal{M}_h is associated with a uniquely defined unit tangent $\boldsymbol{\tau}$, so that either $\mathbf{t}(\mathbf{x}) \cdot \boldsymbol{\tau}(\mathbf{x}) = 1$ or $\mathbf{t}(\mathbf{x}) \cdot \boldsymbol{\tau}(\mathbf{x}) = -1$ for all \mathbf{x} on the same element edge. The degrees of freedom for assembling the global equations from elementwise contributions are then selected to be

$$D_{ij} = \int_{\mathcal{E}_{ij}} \mathbf{v}_h(\mathbf{x}) \cdot \boldsymbol{\tau}(\mathbf{x}) ds.$$

The restriction of the global finite element expansion \mathbf{v}_h on the element can then be written in terms of the local basis functions and the global degrees of freedom as

$$\mathbf{v}_{h|K} = \sum d_{ij} \psi_{ij} = \sum D_{ij} (\mathbf{t} \cdot \boldsymbol{\tau})_{\mathbf{x} \in \mathcal{E}_{ij}} \psi_{ij}.$$

In practice this is implemented by applying sign reversions to the local basis functions when $\mathbf{t} \cdot \boldsymbol{\tau} = -1$. It should be noted that the tangential continuity of the resulting edge element expansion is then in-built.

The strategy to construct basis functions for the Nédélec edge elements of the second kind is analogous to that described in Section F.1.2 for the face elements. That is, each basis function is associated with a certain numerical integration point located on one of the element edges. The tangential component of the basis function is normalized to yield the unity at the integration point to which the basis function is associated, while the tangential component of this basis function is required to vanish at other integration points located on the element edges. Considering the lowest-degree case, in both two and three dimensions each element edge is associated with two basis functions (as the dimension of $\mathbf{X}(\hat{K})$ is 6 for the triangular element and 12

for the tetrahedral element). The degrees of freedom associated with the edge element interpolation function $\hat{v}_h(\hat{\mathbf{x}})$ can then be characterized by

$$\hat{d}_{ij}^k = \int_{\hat{\mathcal{E}}_{ij}} \hat{v}_h(\hat{\mathbf{x}}) \cdot \hat{\mathbf{t}}(\hat{\mathbf{x}}) \hat{\phi}_k(\hat{\mathbf{x}}) d\hat{s},$$

with $\hat{\phi}_k(\hat{\mathbf{x}})$ the linear Lagrange interpolation functions which have the Gaussian points of the element edge as the interpolation points.

The basis of the element space $\mathbf{X}(K)$ is again obtained via applying the adopted version of the Piola transformation. The global degrees of freedom are also defined in the same way as in the case of Nédélec's first family. In addition, an in-built mechanism is now applied to permute the order of basis functions so that the continuity of the tangential components is maintained when elementwise contributions are assembled to the global linear system using the standard assembly subroutines of the Elmer solver.

F.2.2 The optimal edge element family of the lowest degree

The simplest edge elements which have other shapes than a triangle or a tetrahedron may not provide optimal accuracy when the physical elements are not obtained as affine images of the reference element. In this section, we describe a family of edge elements designed to provide optimal accuracy even for non-affine element shapes. The triangular and tetrahedral elements of this family correspond to the Nédélec edge element family of the first kind; cf. the previous section. The other edge element types including quadrilateral, pyramidal, prismatic and hexahedral elements are described in the following.

The curl-conforming version of the quadrilateral Arnold-Boffi-Falk element

In this case the reference element space is taken to be

$$\mathbf{X}(\hat{K}) = Q_{0,2}(\hat{K}) \times Q_{2,0}(\hat{K}). \quad (\text{F.9})$$

The dimension of this space is 6. Four degrees of freedom are associated with the edges of the element and the remaining two are local.

A 10-DOF pyramidal element

The construction of the pyramidal element is based on [2]. Using the reference element of the high-order family, the reference element space is given by

$$\mathbf{X}(\hat{K}) = [P_0(\hat{K})]^3 \oplus \text{span}\{\hat{\varphi}_j \mid j = 4, \dots, 10\} \quad (\text{F.10})$$

where

$$\begin{aligned} \hat{\varphi}_4 &= \frac{1}{(1 - \hat{x}_3/a)^2} \begin{bmatrix} \hat{x}_2(1 - \hat{x}_3/a) \\ \hat{x}_1(1 - \hat{x}_3/a) \\ \hat{x}_1\hat{x}_2/a \end{bmatrix}, \quad \hat{\varphi}_5 = \begin{bmatrix} 1 - \hat{x}_3/a \\ 0 \\ \hat{x}_1/a \end{bmatrix}, \quad \hat{\varphi}_6 = \begin{bmatrix} 0 \\ 1 - \hat{x}_3/a \\ \hat{x}_2/a \end{bmatrix}, \\ \hat{\varphi}_7 &= \frac{\hat{x}_2}{1 - \hat{x}_3/a} \begin{bmatrix} 1 - \hat{x}_3/a \\ 0 \\ \hat{x}_1/a \end{bmatrix}, \quad \hat{\varphi}_8 = \frac{\hat{x}_1}{1 - \hat{x}_3/a} \begin{bmatrix} 0 \\ 1 - \hat{x}_3/a \\ \hat{x}_2/a \end{bmatrix}, \\ \hat{\varphi}_9 &= \frac{\hat{x}_2^2}{(1 - \hat{x}_3/a)^2} \begin{bmatrix} 1 - \hat{x}_3/a \\ 0 \\ \hat{x}_1/a \end{bmatrix}, \quad \hat{\varphi}_{10} = \frac{\hat{x}_1^2}{(1 - \hat{x}_3/a)^2} \begin{bmatrix} 0 \\ 1 - \hat{x}_3/a \\ \hat{x}_2/a \end{bmatrix}, \end{aligned} \quad (\text{F.11})$$

with $a = \sqrt{2}$.

A 15-DOF prismatic element

The construction of the prismatic element is also given in [2]. If the prismatic reference element is represented as $\hat{K} = \hat{C} \times \{\hat{x}_3 \mid -1 \leq \hat{x}_3 \leq 1\}$, with \hat{C} a triangle on a plane $\hat{x}_3 = \text{const}$, the reference element space is defined by

$$\mathbf{X}(\hat{K}) = [\mathbf{Nd}_1(\hat{C}) \otimes P_2(\{\hat{x}_3 \mid -1 \leq \hat{x}_3 \leq 1\})] \times P_2(\hat{C}) \quad (\text{F.12})$$

where $\mathbf{Nd}_1(\hat{C})$ denotes the Nédélec edge element space of the first kind and $P_2(D)$ is the space of the second-order polynomials defined over the domain D .

A 27-DOF hexahedral element

The optimal hexahedral edge element has been developed in [5]. In this case the reference element space is taken to be

$$\mathbf{X}(\hat{K}) = Q_{0,2,2}(\hat{K}) \times Q_{2,0,2}(\hat{K}) \times Q_{2,2,0}(\hat{K}). \quad (\text{F.13})$$

F.2.3 Second-order edge elements

Second-order edge finite elements are also available. They are members of the conventional Nédélec's first family. The background mesh can then consist of either the lowest-order or second-order nodal elements. For example, a tetrahedral background mesh may consist of elements of type 504 or 510.

In the case of triangles and tetrahedra the reference element space is given by (F.6) with $k = 2$ and thus has the dimension 8 or 20, respectively. In the case of quadrilateral elements, the reference element space is

$$\mathbf{X}(\hat{K}) = Q_{1,2}(\hat{K}) \times Q_{2,1}(\hat{K}), \quad (\text{F.14})$$

having the dimension 12.

If the prismatic reference element is represented as $\hat{K} = \hat{C} \times \hat{I}$, with \hat{C} a triangle on a plane $\hat{x}_3 = \text{const}$ and $\hat{I} = \{\hat{x}_3 \mid -1 \leq \hat{x}_3 \leq 1\}$, the reference element space for the second-order prism is defined by

$$\mathbf{X}(\hat{K}) = [\mathbf{Nd}_2(\hat{C}) \otimes P_2(\hat{I})] \times [P_2(\hat{C}) \otimes P_1(\hat{I})] \quad (\text{F.15})$$

where $\mathbf{Nd}_2(\hat{C})$ denotes the second-order Nédélec edge element space of the first kind and $P_k(\hat{I})$ is the space of the k th order polynomials defined over the domain \hat{I} . This space has the dimension 36. On the other hand, the reference element space for bricks is

$$\mathbf{X}(\hat{K}) = Q_{1,2,2}(\hat{K}) \times Q_{2,1,2}(\hat{K}) \times Q_{2,2,1}(\hat{K}), \quad (\text{F.16})$$

having the dimension 54. Finally, the construction of the pyramidal element of degree 2 follows [2] by omitting some of the basis functions associated with the quadrilateral face to obtain the conformity with the elements of the Nédélec's first family. The resulting element has 31 DOFs.

F.2.4 Using integration over the reference element

The effect of the applied version of the Piola transformation needs to be taken into account when integrals over elements K are converted to integrals over the reference element \hat{K} . Therefore the implementation of elementwise integration is again somewhat different as compared with using the standard transformation laws for the basic finite elements.

In the following $\mathbf{u} : K \rightarrow \mathbb{R}^d$ and $\mathbf{v} : K \rightarrow \mathbb{R}^d$ are the transformations of $\hat{\mathbf{u}} : \hat{K} \rightarrow \mathbb{R}^d$ and $\hat{\mathbf{v}} : \hat{K} \rightarrow \mathbb{R}^d$ under the $\mathbf{H}(\text{curl})$ -version of the Piola transformation. A scalar field $\phi : K \rightarrow \mathbb{R}$ is obtained from a field $\hat{\phi} : \hat{K} \rightarrow \mathbb{R}$ via $\phi(\mathbf{x}) = (\hat{\phi} \circ \mathbf{f}_K^{-1})(\mathbf{x})$. In addition, the curl operator with respect to \mathbf{x} is denoted by curl , while \mathbf{Curl} is used to denote the curl operator with respect to $\hat{\mathbf{x}}$. The following transformation laws can then be used to convert integrals over elements K to integrals over the reference element \hat{K} :

$$\begin{aligned} \int_{\partial K} \mathbf{u}(\mathbf{x}) \cdot \mathbf{t}(\mathbf{x}) \phi(\mathbf{x}) dS &= \int_{\partial \hat{K}} \hat{\mathbf{u}}(\hat{\mathbf{x}}) \cdot \hat{\mathbf{t}}(\hat{\mathbf{x}}) \hat{\phi}(\hat{\mathbf{x}}) d\hat{S}, \\ \int_K \mathbf{u}(\mathbf{x}) \cdot \mathbf{v}(\mathbf{x}) d\Omega &= \int_{\hat{K}} [\mathbf{F}(\hat{\mathbf{x}})]^{-T} \hat{\mathbf{u}}(\hat{\mathbf{x}}) \cdot [\mathbf{F}(\hat{\mathbf{x}})]^{-T} \hat{\mathbf{v}}(\hat{\mathbf{x}}) \det \mathbf{F}(\hat{\mathbf{x}}) d\hat{\Omega}. \end{aligned} \quad (\text{F.17})$$

In addition, the referential description of the field $\mathbf{curl} \mathbf{u}(\mathbf{x})$ is given by

$$(\mathbf{curl} \mathbf{u})(\mathbf{f}_K(\hat{\mathbf{x}})) = \frac{1}{\det \mathbf{F}(\hat{\mathbf{x}})} \mathbf{F}(\hat{\mathbf{x}}) \mathbf{Curl} \hat{\mathbf{u}}(\hat{\mathbf{x}}). \quad (\text{F.18})$$

In two dimensions the only nontrivial component of the curl field transforms in particular as

$$(\mathbf{curl} \mathbf{u} \cdot \mathbf{e}_3)(\mathbf{f}_K(\hat{\mathbf{x}})) = \frac{1}{\det \mathbf{F}(\hat{\mathbf{x}})} \mathbf{Curl} \hat{\mathbf{u}}(\hat{\mathbf{x}}) \cdot \hat{\mathbf{e}}_3. \quad (\text{F.19})$$

In view of (F.18), we thus have

$$\int_K \mathbf{curl} \mathbf{u}(\mathbf{x}) \cdot \mathbf{curl} \mathbf{v}(\mathbf{x}) d\Omega = \int_{\hat{K}} \mathbf{F}(\hat{\mathbf{x}}) \mathbf{Curl} \hat{\mathbf{u}}(\hat{\mathbf{x}}) \cdot \mathbf{F}(\hat{\mathbf{x}}) \mathbf{Curl} \hat{\mathbf{v}}(\hat{\mathbf{x}}) \frac{1}{\det \mathbf{F}(\hat{\mathbf{x}})} d\hat{\Omega}. \quad (\text{F.20})$$

F.2.5 Setting Dirichlet boundary conditions

In connection with edge element interpolation Dirichlet constraints can be specified to prescribe the tangential components of the unknown field on the boundary. In the case of the simplest edge elements the values of the prescribed degrees of freedom can be generated by computing integrals of the type

$$D = \int_{\mathcal{E}_{ij}} \mathbf{u}_h(\mathbf{x}) \cdot \boldsymbol{\tau}(\mathbf{x}) \phi(\mathbf{x}) ds$$

over element edges \mathcal{E}_{ij} . In the case of elements of the second degree or elements from the optimal family some degrees of freedom may be associated with element faces and a more general approach must be followed. In that case the values of the degrees of freedom associated with the element faces are obtained by seeking the best approximation of the data in L_2 .

Specific syntax must be used in the solver input file to define boundary conditions for the edge element approximation. The following commands can be used in the `Boundary Condition` section to define the boundary values of the solver variable `U` which corresponds to a continuous field \mathbf{u} and is approximated with the edge element basis functions:

`U {e} j Real`

This keyword is used to define the vector \mathbf{g} so that its tangential trace $\mathbf{g} \times \mathbf{n}$ is approximated by $\mathbf{u}_h \times \mathbf{n}$, with \mathbf{u}_h the finite element interpolating function. The value of this keyword defines the components g_j , $j \in \{1, 2, 3\}$, with respect to the global Cartesian coordinate system.

`U {e} Real`

If the edge element has only degrees of freedom which are associated with edges, this keyword can be used to define a scalar field \hat{g}_τ so that the values of the prescribed degrees of freedom are computed as

$$D = \int_{\mathcal{E}_{ij}} [\mathbf{g}(\mathbf{x}) \cdot \boldsymbol{\tau}(\mathbf{x}) + \hat{g}_\tau] \phi(\mathbf{x}) ds.$$

This keyword should not be used if additional degrees of freedom associated with the element faces are present (that is, this feature should not be used for elements of the second degree or elements from the optimal family).

Bibliography

- [1] Ervin V. J. Computational bases for RT_k and BDM_k on triangles. *Computers & Mathematics with Applications*, 64:2765–2774, 2012.
- [2] Bergot M. and Duruflé M. High-order optimal edge elements for pyramids, prisms and hexahedra. *Journal of Computational Physics*, 232:189–213, 2013.

- [3] Arnold D. N., Boffi D., and Falk R. S. Quadrilateral $H(\text{div})$ finite elements. *SIAM Journal on Numerical Analysis*, 42:2429–2451, 2005.
- [4] Arnold D. N., Falk R. S., and Winther R. Geometric decompositions and local bases for spaces of finite element differential forms. *Computer Methods in Applied Mechanics and Engineering*, 198:1660–1672, 2009.
- [5] Falk R. S., Gatto P., and Monk P. Hexahedral $H(\text{div})$ and $H(\text{curl})$ finite elements. *ESAIM: Mathematical Modelling and Numerical Analysis*, 45:115–143, 2011.

Appendix G

Advanced finite element definitions

Advanced finite element formulations may require that several fields are approximated simultaneously using different types of interpolation. For example curl-conforming (edge finite element) approximation may need to be used in combination with the standard Lagrange interpolation of another field. This requires advanced flexibility in creating suitable degrees of freedom for finite elements.

Elmer offers general support for creating advanced finite element formulations when a background mesh for representing the standard lowest-degree continuous finite element expansion is available. Additional degrees of freedom may then be added by associating them with the edges, faces or interiors of the elements of the background mesh. This enhancement strategy is used internally for example in the creation of the second-order curl-conforming finite elements. In general, the writer of a solver must nevertheless decide how the additional degrees of freedom are utilized in the finite element approximation.

G.1 Specification of the degrees of freedom

Each solver specification in the `sif` file may contain an additional command of the type

```
Element = "n:N ... "
```

where the string between the quotation marks controls what degrees of freedom (DOFs) are allocated. The specification string can contain the following parts:

- `"n:N"` to create $N \geq 0$ DOFs associated with each node of the background mesh.
- `"e:E"` to create $E \geq 0$ DOFs associated with each edge of the background mesh.
- `"f:F"` to create $F \geq 0$ DOFs associated with each face contained in the background mesh.
- As a more general alternate for the previous definition, `"-quad_face b:FQ -tri_face b:FT"` associates F_Q DOFs with every quadrilateral face and F_T DOFs with every triangular face of the background mesh.
- `"b:B"` to create $B \geq 0$ DOFs associated with each element interior
- As a more general alternate for the previous definition, use

```
"-tetra b:BT -pyramid b:Bp -prism b:BW -brick b:BB"
```

controls the number of the local DOFs associated with the element interiors in a manner that depends on the element shape (tetrahedron, pyramid, prism or brick).

Index

- dofs, 58
- global, 58
- nooutput, 58

- Adaptive Error Limit, 68
- Adaptive Error Measure, 37
- Adaptive Keep Smallest, 37
- Adaptive Max Change, 68
- Adaptive Max H, 68
- Adaptive Mesh Refinement, 68
- Adaptive Min H, 68
- Adaptive Min Timestep, 37
- Adaptive Remesh, 68
- Adaptive Save Mesh, 68
- Adaptive Time Error, 37
- Adaptive Timestepping, 36
- After Linsolve, 30
- algebraic multigrid, 21
- Anti Periodic BC VarName, 47
- Anti Rotational Projector, 49
- Anti Sliding Projector, 49
- Apply Contact BCs, 50, 53
- Apply Limiter, 45
- Apply Mortar BCs, 49, 53

- Backward Differences Formulae, 34
- bandwidth optimization, 26
- BDF, 34
- BDF Order, 35
- Before Linsolve, 29, 45
- BiCGstab(ℓ), 24
- BiCGstab, biconjugate gradient stabilized, 24
- BiCGstabl polynomial degree, 25
- Binary output, 120
- Biorthogonal Basis, 50
- Biorthogonal Dual LCoeff, 50
- Biorthogonal Dual Master, 50
- Body, 8
- Body Force, 8, 44, 45, 55, 60
- BoomerAMG, 78
- BoomerAMG Coarsen Type, 78
- BoomerAMG Cycle Type, 79
- BoomerAMG Interpolation Type, 79
- Boomeramg Max Levels, 79
- BoomerAMG Num Functions, 79
- BoomerAMG Num Sweeps, 79

- BoomerAMG Relax Type, 78
- BoomerAMG Smooth Type, 79
- Bossak Alpha, 36
- Boundary Condition, 9, 44–46, 49, 50, 54
- Boundary Condition bc_id, 64

- Calculate Acceleration, 59
- Calculate Energy Norm, 56
- Calculate Loads, 55
- Calculate Velocity, 59
- Calculate Weights, 56
- CG, conjugate gradient, 24
- CGS, conjugate gradient squared, 24
- Check Keywords, 12
- compilation instructions, 82
- Component, 9
- component name, 57
- Compressed Row Storage, 28
- Constants, 8
- Constrained Modes Varname, 54
- constraint modes, 54
- Constraint Modes Analysis, 54
- Constraint Modes Fluxes, 54
- Constraint Modes Fluxes Symmetric, 54
- Constraint Modes Lumped, 54
- Contact Active Condition, 50
- Contact Active Set Minimum, 51
- Contact Depth Offset, 50
- Contact Depth Offset Initial, 50
- Contact Passive Condition, 51
- Contact Type, 51
- Contact Velocity(dim), 51
- Convergence Tolerance, 32
- Coordinate Mapping(dim), 61
- Coordinate Scaling, 61
- Coordinate Transformation , 61
- Coordinate Transformation Radius, 61
- Corrector method, 38
- Crank-Nicolson method, 34
- Cuthill-McKee algorithm, 26
- Cylindrical Projector, 49

- DG Indirect Connections, 65
- DG Reduced Basis, 65
- DG Reduced Basis Mapping, 65
- DG Reduced Basis Master Bodies, 65

- DG Reduced Basis Slave Bodies, 65
- Direct methods, 19
- direct solver, 24
- Dirichlet constraints, 43
- Discontinuous BC, 64
- Discontinuous Boundary, 64
- Discontinuous Boundary Greedy, 64
- Discontinuous Bulk Greedy, 64
- Discontinuous Galerkin, 65
- Discontinuous Galerkin method, 75
- Discontinuous Target Bodies, 64
- domain decomposition, 71, 76
- Dynamic Friction Coefficient, 51

- Edge Basis, 49
- Eigen Analysis, 41
- Eigen System Complex, 41
- Eigen System Compute Residuals, 41
- Eigen System Convergence Tolerance, 41
- Eigen System Damped, 41
- Eigen System Lanczos Vectors, 41
- Eigen System Max Iterations, 41
- Eigen System Select, 41
- Eigen System Shift, 41
- Eigen System Use Identity, 41
- Eigen System Values, 41
- Eigenvalue problems, 39
- element definition command, 130
- Eliminate From Master, 53
- Eliminate Linear Constraints, 53
- Eliminate Slave, 53
- ElmerPost input file, 124
- ElmerSolver mesh file, 117
- Equation, 8
- error indicator, 66
- Exec Condition, 57
- Exec Intervals, 57
- Exec Solver, 57
- Export Lagrange Multiplier, 53
- Exported and derived variables, 58
- Exported Variable i, 59
- Exported Variable i Mask, 59
- Extruded Coordinate Index, 63
- Extruded Max Coordinate, 63
- Extruded Mesh Density, 63
- Extruded Mesh Levels, 63
- Extruded Mesh Ratio, 63
- Extruded Min Coordinate, 63

- Friction Contact, 51

- GCR, generalized conjugate residual, 24
- geometric multigrid, 21
- Geometric Stiffness, 41
- GMRES, generalized minimal residual, 24

- Header, 8

- ILU preconditioners, 21
- ILU, incomplete LU-decomposition, 25
- Initial Condition, 9
- Initialization of dependent variables, 17
- Initialize Dirichlet Condition, 17
- iterative solver, 24

- keyword syntax, 12
- Krylov methods, 24
- Krylov subspace methods, 20

- Lagrange Multiplier Name, 53
- LAPACK, 19
- Level Projector, 49
- Level Projector Extruded Edges Strong, 50
- Level Projector Nodes Strong, 49
- Level Projector Plane Edges Strong, 50
- limiter, 45
- Limiter Load Sign Negative, 45
- Limiter Load Tolerance, 45
- Limiter Value Tolerance, 45
- linear algebra, 19
- Linear System Abort Not Converged, 25
- Linear System Convergence Tolerance, 25
- Linear System Direct Method, 24
- Linear System GCR Restart, 25
- Linear System GMRES Restart, 25
- Linear System ILUT Tolerance, 25
- Linear System Iterative Method, 24
- Linear System Max Iterations, 25
- Linear System Precondition Recompute, 26
- Linear System Preconditioning, 25, 77, 78
- Linear System Residual Mode, 25
- Linear System Residual Output, 26
- Linear System Solver, 24
- Linear System Symmetric, 45
- Linear System Timing, 60
- Linear System Timing Cumulative, 60
- Linear System Use Hypre, 77
- linear systems, 19
- lower limit, 45
- Lumped Mass Matrix, 36

- MATC, 12, 13
 - parameterized keyword commands, 14
- Material, 8
- mesh database, 72
- Mesh DB, 9
- Mesh Grading Power, 62
- Mesh Keep, 62
- Mesh Keep Grading, 62
- Mesh Levels, 62
- mesh refinement, 66

- MG Cluster Alpha, 28
- MG Cluster Size, 28
- MG Convergence Tolerance, 26
- MG Direct Interpolate, 28
- MG Direct Interpolate Limit, 28
- MG Eliminate Dirichlet, 27
- MG Eliminate Dirichlet Limit, 27
- MG Equal Split, 26
- MG ILUT Tolerance, 27
- MG Levels, 26
- MG Lowest Linear Solver Limit, 27
- MG Max Iterations, 26
- MG Mesh Name, 26
- MG Positive Connection Limit, 28
- MG Post Smoothing Iterations, 27
- MG Pre Smoothing Iterations, 27
- MG Preconditioning, 27
- MG Projection Limit, 28
- MG Recompute Projector, 27
- MG Smoother, 27
- MG SOR Relax, 27
- MG Strong Connection Limit, 28
- MG Strong Connection Minimum, 28
- Mortar BC, 49
- Mortar BC Scaling, 49
- multigrid solver, 24
- Mumps, 20

- nodal loads, 55
- Nonlinear Calculate Acceleration, 59
- Nonlinear Calculate Velocity, 59
- Nonlinear System Convergence Absolute, 32
- Nonlinear System Convergence Measure, 31
- Nonlinear System Convergence Tolerance, 32
- Nonlinear System Convergence Without Constraints, 53
- Nonlinear System Max Iterations, 32
- Nonlinear System Newton After Iterations, 32
- Nonlinear System Newton After Tolerance, 32
- Nonlinear System Norm Degree, 32
- Nonlinear System Norm Dofs, 32
- Nonlinear System Relaxation Factor, 32
- Nonlinear Update Exported Variables, 59
- nonlinearity, 31
- Normal-Tangential VarName, 51

- Optimize Bandwidth, 26
- Output Intervals, 36

- Parasails, 77
- Pardiso, 20
- Periodic BC, 46
- Periodic BC Explicit, 46
- Periodic BC Rotate(3), 46
- Periodic BC Scale(3), 47
- Periodic BC Translate(3), 46
- Periodic BC Use Lagrange Coefficients, 47
- Periodic BC VarName, 47
- Periodic conditions, 46
- Plane Projector, 49
- preconditioning, 21
- Predictor method, 38
- Predictor-Corrector Control, 38
- Predictor-Corrector Control Beta 1, 38
- Predictor-Corrector Control Beta 2, 38
- Predictor-Corrector Control Tolerance, 38
- Predictor-Corrector Save Error, 38
- Predictor-Corrector Scheme Order, 38
- Preserve Baseline, 63
- Projector Skip Edges, 50
- Projector Skip Nodes, 50

- Radial Projector, 49
- Reference Norm, 60
- Reference Norm Tolerance, 60
- Relaxation Factor, 32
- Restart Before Initial Conditions, 17
- Restart file, 16
- Restart Position, 16
- Rotational Projector, 49

- Save Projector, 50
- Set Dirichlet BCs by BC Numbering, 44
- sif file, 8
- Simulation, 8, 43, 61–64
- Skip First Timestep, 38
- Slide Contact, 51
- Sliding Projector, 49
- Solver, 9, 45, 49, 50, 54–56, 59, 60, 65
- Solver activation, 57
- solver input file, 8
 - section names, 8
- Solver Timing, 60
- Solver Timing Cumulative, 60
- SOLVER.KEYWORDS, 12
- Stability Analysis, 41
- Start Time, 57
- Static Friction Coefficient, 51
- Steady State Convergence Measure, 32
- Steady State Convergence Tolerance, 33
- Steady State Max Iterations, 33
- Steady State Min Iterations, 33
- Steady State Norm Degree, 32
- Steady State Norm Dofs, 32
- Steady State Relaxation Factor, 33
- Stick Active Condition, 51
- Stick Contact, 51
- Stick Contact Coefficient, 51
- Stick Passive Condition, 51
- Stop Time, 57

SuperLU, 20

Target Boundaries(n), 44

Target Coordinates(n,DIM), 44

Target Nodes(n), 44

Tie Contact, 51

Time Derivative Order, 35

time-dependent, 34

timeintegration, 34

Timestep Intervals, 36

Timestep Size, 36

Timestep Sizes, 36

Timestepping Method, 35

U {e}, 152

U {e} j, 152

Umfpack, 19

Update Exported Variables, 59

upper limit, 45

Use Transpose values, 53

Variable name, 57

Varname, 44

Varname {e}, 44

Varname {e} i, 44

VarName Calculate Acceleration, 59

VarName Calculate Velocity, 59

Varname Condition, 44, 45

Varname i, 44

Varname Load , 55

Varname Lower Limit, 45

Varname Passive, 60

VarName Transient, 59

Varname Upper Limit, 45